

DScope: Detecting Real-World Data Corruption Hang Bugs in Cloud Server Systems

Ting Dai¹, Jingzhu He¹, Xiaohui (Helen) Gu¹, Shan Lu², Peipei Wang¹

¹*NC State University*

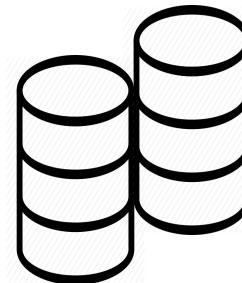
²*University of Chicago*



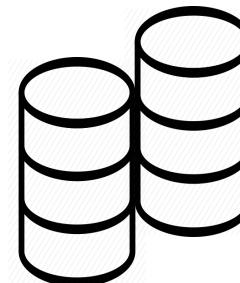
Real-World Data Corruption Problem



British Airway service was down for **hours** with financial penalty of **£ 100 million**.



Primary data center

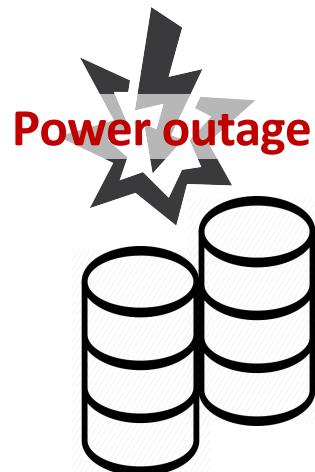


Backup data center

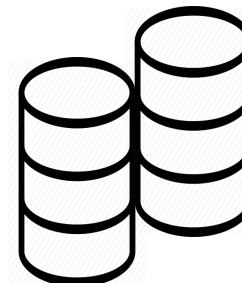
Real-World Data Corruption Problem



British Airway service was down for **hours** with financial penalty of **£ 100 million**.



Primary data center

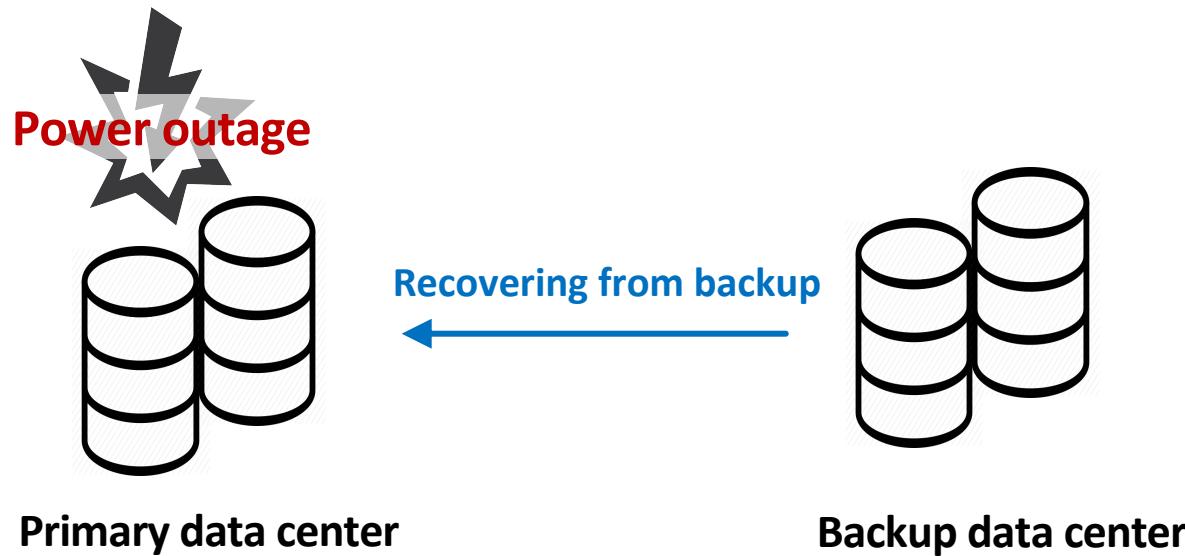


Backup data center

Real-World Data Corruption Problem



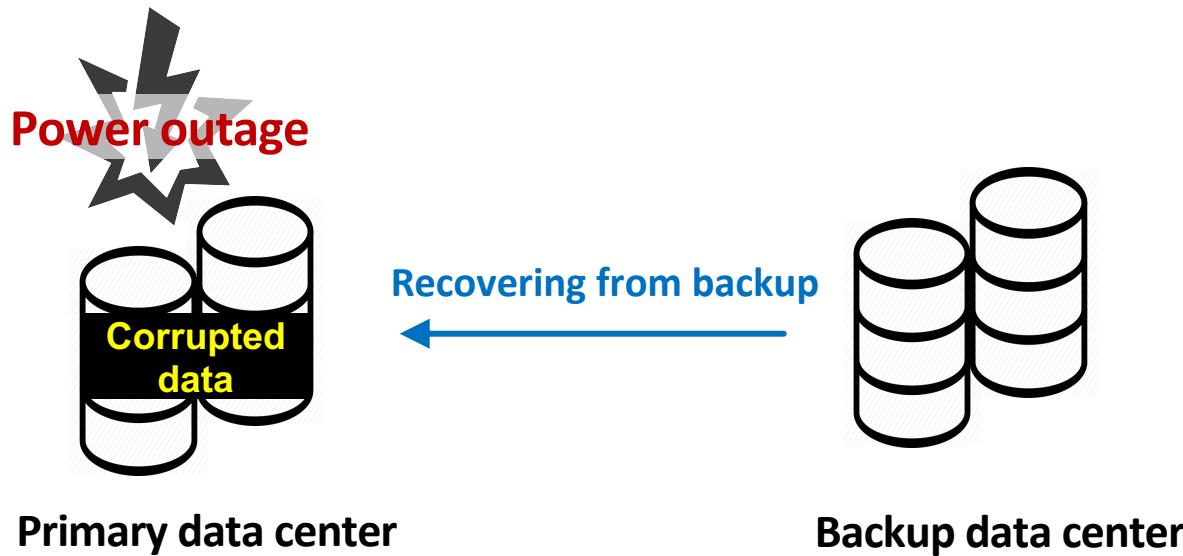
British Airway service was down for **hours** with financial penalty of **£ 100 million**.



Real-World Data Corruption Problem



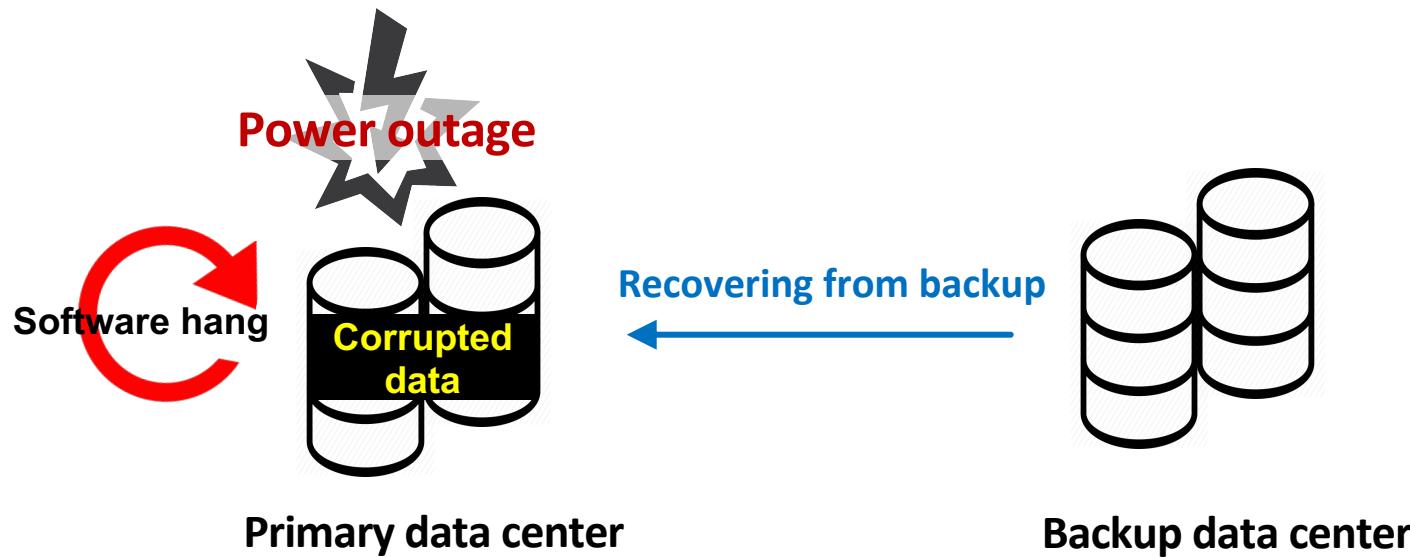
British Airway service was down for **hours** with financial penalty of **£ 100 million**.



Real-World Data Corruption Problem



British Airway service was down for **hours** with financial penalty of **£ 100 million**.



A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len);  
        ...  
        ...  
189         len -= ret;  
190     }  
191 }
```

Overview of DSscope

A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
InputStream  
        ...  
        ...  
189         len -= ret;  
190     }  
191 }
```

Overview of DSscope

A Data Corruption Hang Bug Example

Overview of DSscope

Hadoop-8614

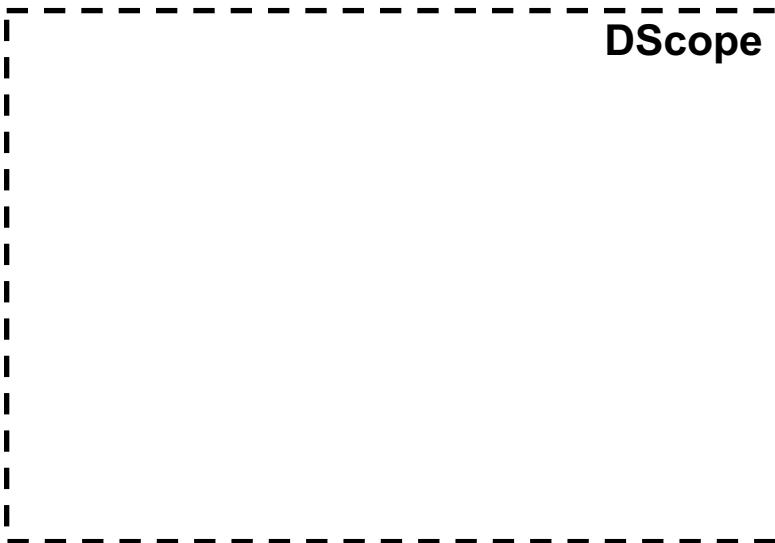
```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
186         ... InputStream  
187         ...  
188         len -= ret; The loop stride (ret) is  
189     } always 0 when in is  
190 } corrupted.  
191 }
```

A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
186         ... InputStream  
187         ...  
188         len -= ret; The loop stride (ret)  
189     } is always 0 when in is  
190 } corrupted.  
191 }
```

Overview of DSscope

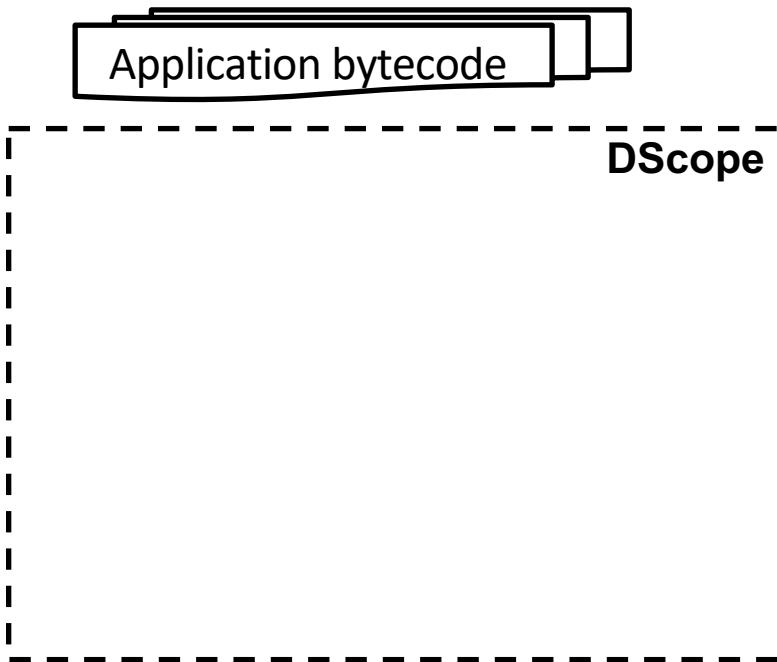


A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
186         ... InputStream  
187         ...  
188         len -= ret; The loop stride (ret)  
189     } is always 0 when in is  
190 } corrupted.  
191 }
```

Overview of DScope

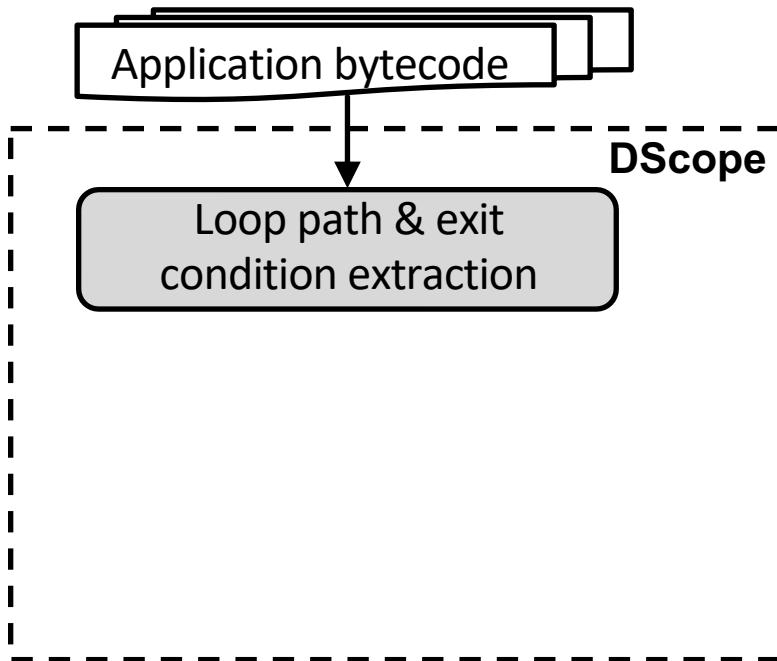


A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
InputStream  
        ...  
        ...  
189         len -= ret; The loop stride (ret) is  
always 0 when in is  
corrupted.  
190     }  
191 }
```

Overview of DScope

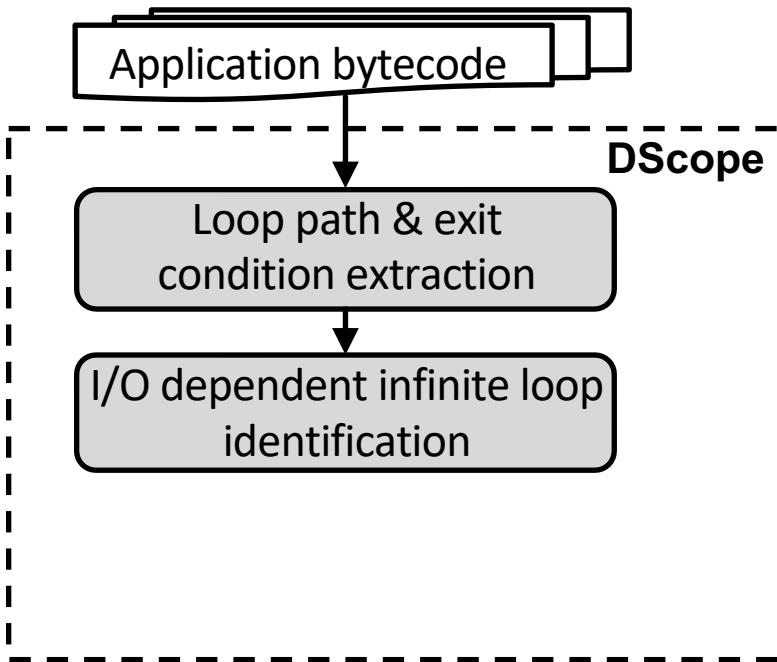


A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
InputStream  
        ...  
        ...  
189         len -= ret; The loop stride (ret) is  
always 0 when in is  
corrupted.  
190     }  
191 }
```

Overview of DScope

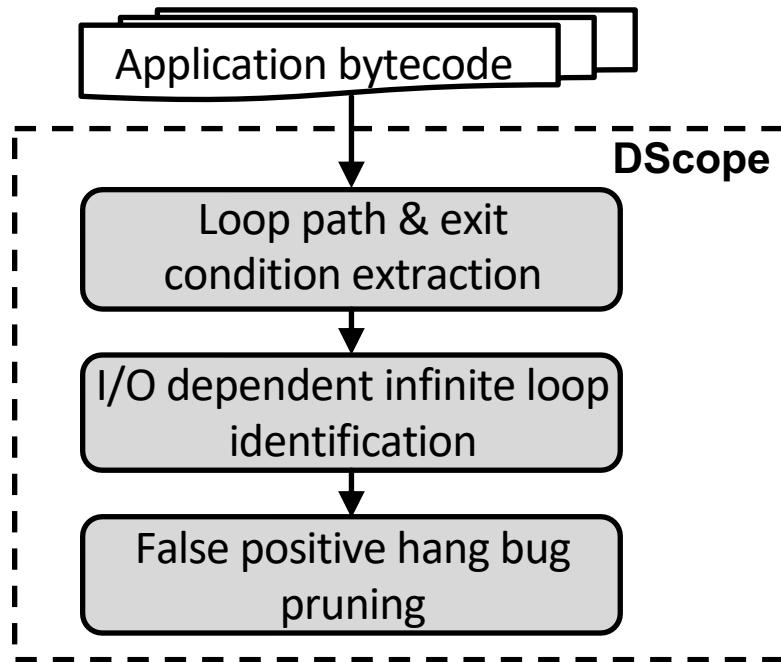


A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
InputStream  
        ...  
        ...  
189         len -= ret; The loop stride (ret) is  
always 0 when in is  
corrupted.  
190     }  
191 }
```

Overview of DScope

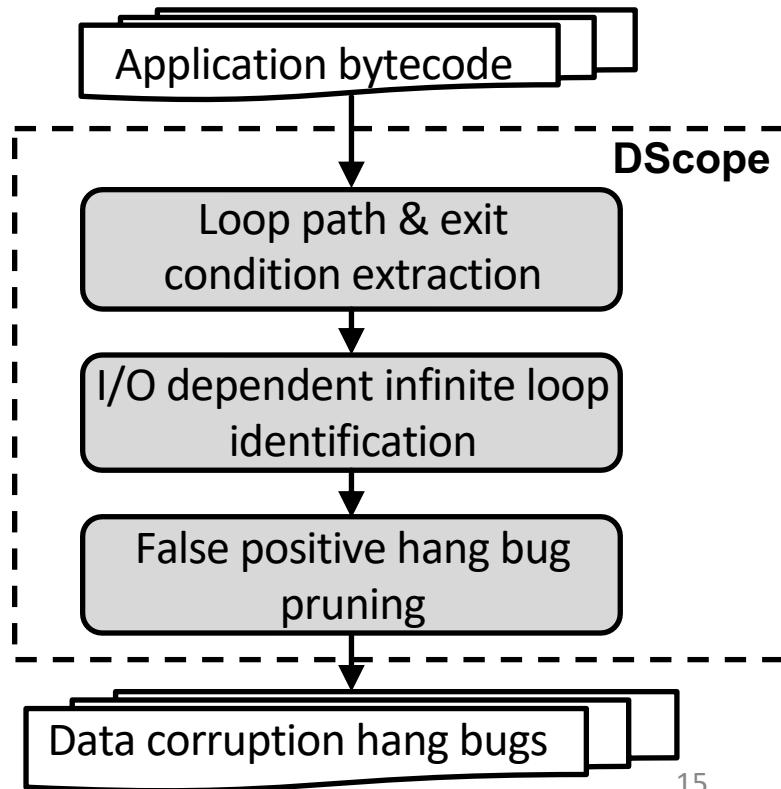


A Data Corruption Hang Bug Example

Hadoop-8614

```
183 public static void skipFully(  
    InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len); Corrupted  
InputStream  
        ...  
        ...  
189         len -= ret; The loop stride (ret) is  
always 0 when in is  
corrupted.  
190     }  
191 }
```

Overview of DSscope



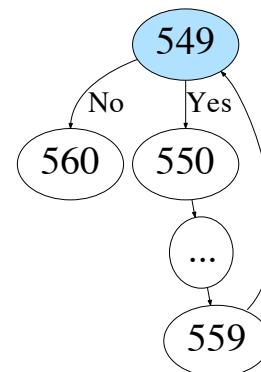
Loop Path & Exit Condition Extraction

- Simple Loops

Loop Path & Exit Condition Extraction

- Simple Loops

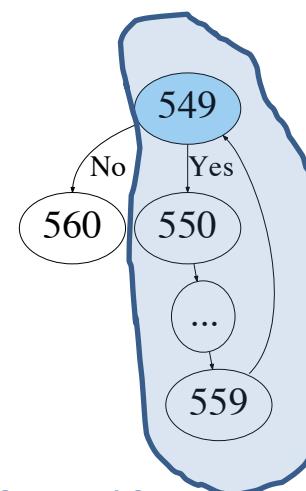
```
549  for ( int j = 0; j < length; j++) {  
550      String rack = racks[j] ;  
...  
559 }  
560
```



Loop Path & Exit Condition Extraction

- Simple Loops

```
549  for ( int j = 0; j < length; j++) {  
550      String rack = racks[j] ;  
...  
559 }  
560
```



Loop path: 549 → 550 → ... → 559 → 560 → 549

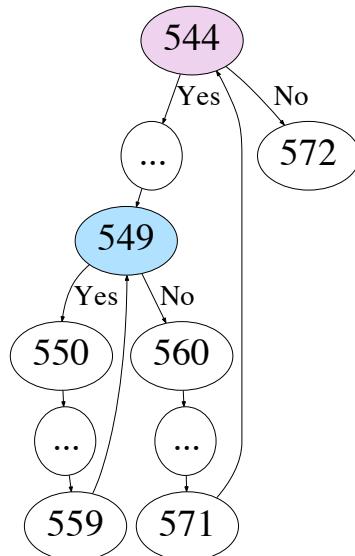
Exit condition: $j \geq length$

Loop Path & Exit Condition Extraction

- Nested Loops

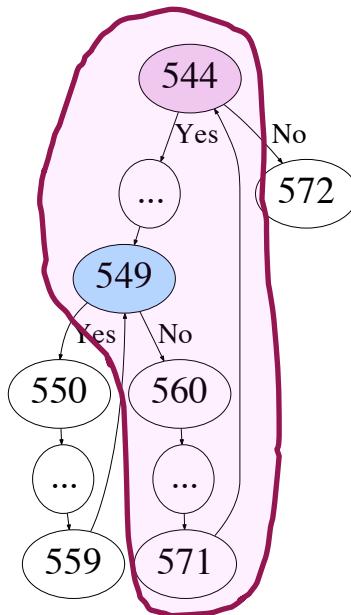
Loop Path & Exit Condition Extraction

- Nested Loops



Loop Path & Exit Condition Extraction

- Nested Loops

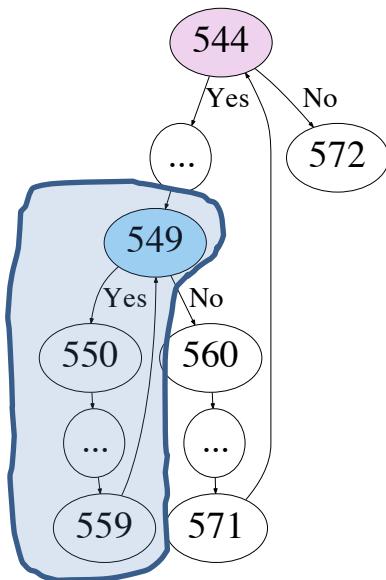


Loop paths:

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Loop Path & Exit Condition Extraction

- Nested Loops



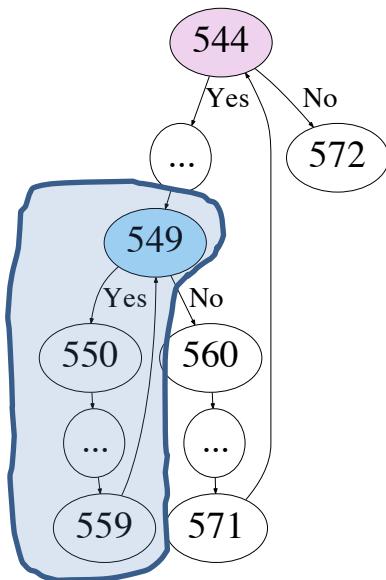
Loop paths:

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Inner: 549 → 550 → ... → 559 → 549

Loop Path & Exit Condition Extraction

- Nested Loops



Loop paths:

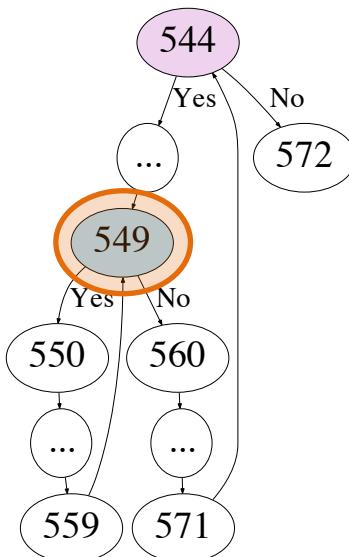
Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Inner: 549 → 550 → ... → 559 → 549

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Loop Path & Exit Condition Extraction

- Nested Loops



Loop paths:

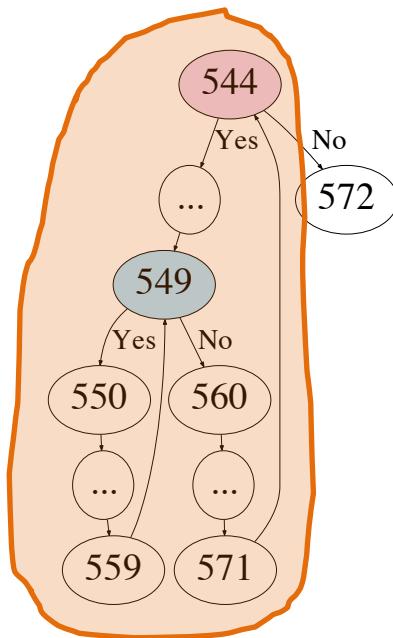
Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Inner: 549 → 550 → ... → 559 → 549

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Loop Path & Exit Condition Extraction

- Nested Loops



Loop paths:

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

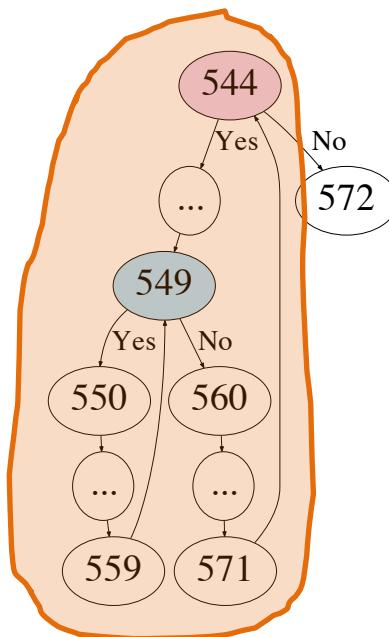
Inner: 549 → 550 → ... → 559 → 549

Outer: 544 → ...

560 → ... → 571 → 544

Loop Path & Exit Condition Extraction

- Nested Loops



Loop paths:

Outer: 544 → ... → 549 → 560 → ... → 571 → 544

Inner: 549 → 550 → ... → 559 → 549

Outer: 544 → ...

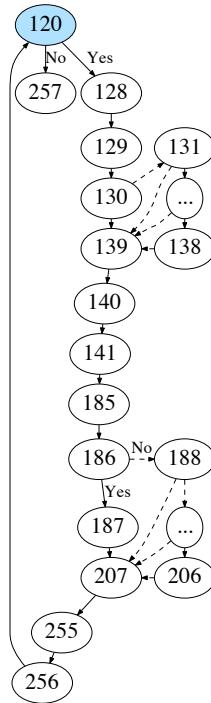
560 → ... → 571 → 544

DScope then extracts the exit conditions for each loop path.

Loop Path & Exit Condition Extraction

- Loops with exception handling

```
120 while (!dataFile.isEOF()) {  
    ...  
129     try {  
130         key = decorateKey(...dataFile);  
        ...  
139     } catch (Throwable th) {  
140         //ignore exception  
141     }  
    ...  
185     try {  
186         if (key == null)  
187             throw new IOError(...);  
        ...  
207     } catch (Throwable th) {  
208         //ignore exception  
    } }
```



Loop Path & Exit Condition Extraction

- Loops with exception handling

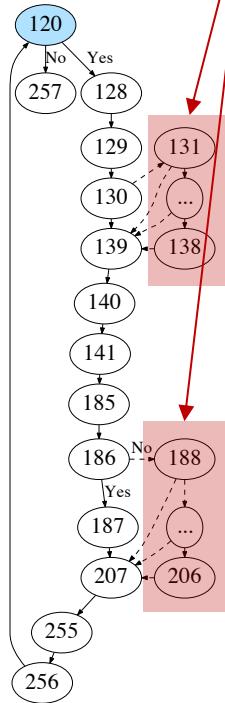
```
120 while (!dataFile.isEOF()) {  
    ...  
129     try {  
130         key = decorateKey(...dataFile);  
131         ...  
139     } catch (Throwable th) {  
140         //ignore exception  
141     }  
142     ...  
185     try {  
186         if (key == null)  
187             throw new IOException(...);  
188         ...  
207     } catch (Throwable th) {  
208         //ignore exception  
209     }  
210 }
```

throw exception

Corrupted dataFile

throw exception

Infeasible path



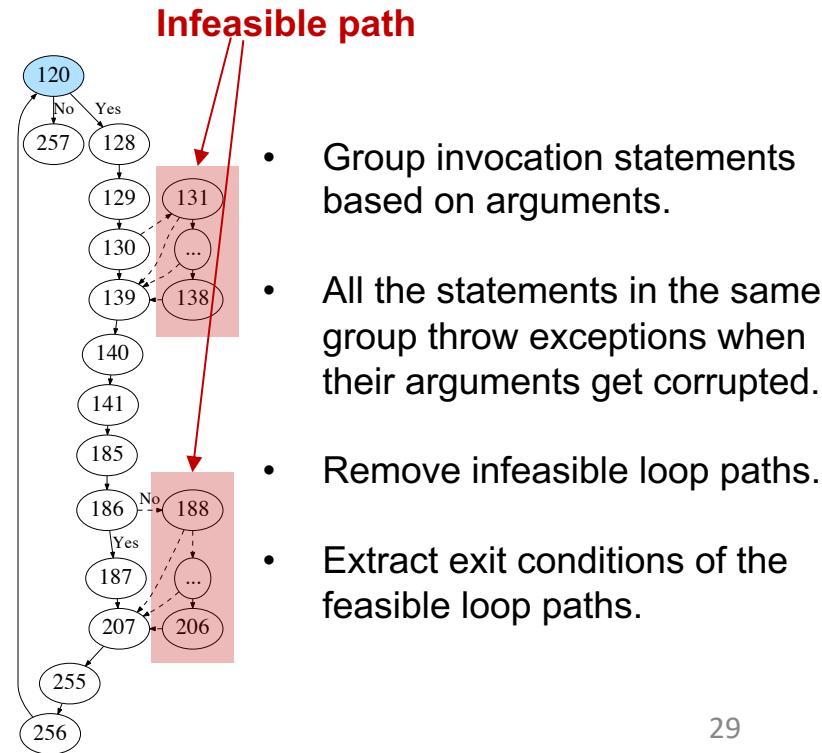
Loop Path & Exit Condition Extraction

- Loops with exception handling

```

120 while (!dataFile.isEOF()) {
...
129 try {
    key = decorateKey(...dataFile);
Corrupted dataFile
...
139 } catch (Throwable th) {
140     //ignore exception
141 }
...
185 try {
186     if (key == null)
187         throw new IOException(...);
throw exception
...
207 } catch (Throwable th) {
208     //ignore exception
}
}

```



I/O Dependent Infinite Loop Identification

- Exit conditions **directly** depend on I/O operations

//Soot IR

```
198 $i1 = r0.<InputStream: read()>(r2) // $i1 is an I/O related variable  
199 if $i1 == -1 goto line #203           // ``$i1 == -1'' is the exit condition  
...  
202 goto line #198
```

I/O Dependent Infinite Loop Identification

- Exit conditions **indirectly** depend on I/O operations

//Soot IR

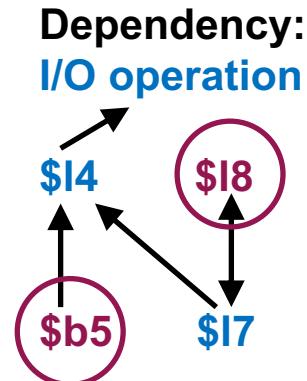
```
3 if |I8 >= |I0| goto line #12 //``|I8 >= |I0|'' is the exit condition  
...  
5 $I2 = |I0 - |I8|  
6 $I4 = $r2.<InputStream: skip>($I2) //|$I4 is an I/O related variable  
7 $b5 = $I4 cmp 0L  
8 if |$b5 == 0| goto line #12 //``|$b5 == 0|'' is the exit condition  
9 $I7 = $I8 + $I4  
10 i8 = $I7  
11 goto line #3
```

I/O Dependent Infinite Loop Identification

- Exit conditions **indirectly** depend on I/O operations

//Soot IR

```
3 if $I8 >= $I0 goto line #12 //``$I8 >= $I0'' is the exit condition  
...  
5 $I2 = $I0 - $I8  
6 $I4 = $r2.<InputStream: skip>($I2) //``$I4 is an I/O related variable  
7 $b5 = $I4 cmp 0L  
8 if $b5 == 0 goto line #12 //``$b5 == 0'' is the exit condition  
9 $I7 = $I8 + $I4  
10 $I8 = $I7  
11 goto line #3
```



I/O Dependent Infinite Loop Identification

- Exit conditions depend on **complex** I/O related variables
 - DScope performs an integrated analysis by linking variable information from IR code, Java source code, and Java bytecode.
 - User annotated I/O variables.

False Positive Filtering

Hadoop v2.5.0 WritableUtils.java

```
307 public static long readVLong(DataInput stream)...{  
308     byte firstByte = stream.readByte();  
309     int len = decodeVIntSize(firstByte);  
310     ...  
314     for (int idx = 0; idx < len-1; idx++) {  
315         ...  
316     } }
```

False Positive Filtering

Hadoop v2.5.0 WritableUtils.java

```
307 public static long readVLong(DataInput stream)...{  
308     byte firstByte = stream.readByte();  
309     int len = decodeVIntSize(firstByte);  
310     ...  
314     for (int idx = 0; idx < len-1; idx++) {  
315         ...  
316     } }  
    ...  
    len is I/O dependent
```

False Positive Filtering

Hadoop v2.5.0 WritableUtils.java

```
307 public static long readVLong(DataInput stream)...{  
308     byte firstByte = stream.readByte();  
309     int len = decodeVIntSize(firstByte);  
310     ...  
314     for (int idx = 0; idx < len-1; idx++) {  
315         ...  
316     } }
```

len is I/O dependent

It's a FP because the loop stride is always 1 and the upper bound (`len-1`) is fixed.

False Positive Filtering

Hadoop v2.5.0 WritableUtils.java

```
307 public static long readVLong(DataInput stream)...{  
308     byte firstByte = stream.readByte();  
309     int len = decodeVIntSize(firstByte);  
310     ...  
314     for (int idx = 0; idx < len-1; idx++) {  
315         ...  
316     }  
317 }
```

len is I/O dependent

It's a FP because the loop stride is always 1 and the upper bound (len-1) is fixed.

- **False positive condition:**
 - The loop stride is always **positive** when the loop index has a fixed **upper** bound;
 - The loop stride is always **negative** when the loop index has a fixed **lower** bound.

Loop Stride and Bound Inference

Loop Stride and Bound Inference

- Stride and bounds are denoted by
 - Numeric primitives

```
for (int idx = 0; idx < len-1; idx++) {
```

```
    ...
```

```
}
```

Bound (len-1)

Stride (1)

Loop Stride and Bound Inference

- **Stride and bounds are denoted by**
 - **APIs in 60 commonly used Java classes**
Forward index Reverse index Check bounds
Reset index Update bounds

Loop Stride and Bound Inference

- Stride and bounds are denoted by
 - APIs in 60 commonly used Java classes
 - Forward index Reverse index Check bounds
 - Reset index Update bounds

```
RandomAccessReader dataFile;
```

```
while (!dataFile.isEOF()) {
```

Bound checking

```
...
```

```
    dataSize = dataFile.readLong();
```

Stride forwarding

```
}
```

Evaluation

System	Description	# of bugs
Cassandra	Distributed database management system	2
Compress	Libraries for I/O ops on compressed file	2
Hadoop Common	Hadoop utilities and libraries	10
Mapreduce	Hadoop big data processing framework	5
HDFS	Hadoop distributed file system	4
Yarn	Hadoop resource management platform	4
Hive	Data warehouse	12
Kafka	Distributed streaming platform	1
Lucene	Indexing and search server	2

- Implemented a prototype of DScope using Soot;
- State-of-the-art static bug detectors:
 - Findbugs
 - Infer

Bug Detection Results

System		DScope		Findbugs	Infer
		TP	FP	TP	TP
Cassandra	v2.0.8	2	1	0	1
Compress	v1.0	2	2	0	-
Hadoop Common	v0.23.0	4	6	0	0
	v2.5.0	6	6	0	0
Mapreduce	v0.23.0	3	0	0	0
	v2.5.0	2	0	0	0
HDFS	v0.23.0	1	1	0	0
	v2.5.0	3	5	1	-
Yarn	v0.23.0	2	2	1	0
	v2.5.0	2	5	0	0
Hive	v1.0.0	7	6	0	-
	v2.3.2	5	1	0	0
Kafka	v0.10.0.0	1	1	0	0
Lucene	V2.1.0	2	1	0	0
Total		42	37	2	1

Bug Detection Results

System		DScope		Findbugs	Infer
		TP	FP	TP	TP
Cassandra	v2.0.8	2	1	0	1
Compress	v1.0	2	2	0	-
Hadoop Common	v0.23.0	4	6	0	0
	v2.5.0	6	6	0	0
Mapreduce	v0.23.0	3	0	0	0
	v2.5.0	2	0	0	0
HDFS	v0.23.0	1	1	0	0
	v2.5.0	3	5	1	-
Yarn	v0.23.0	2	2	1	0
	v2.5.0	2	5	0	0
Hive	v1.0.0	7	6	0	-
	v2.3.2	5	1	0	0
Kafka	v0.10.0.0	1	1	0	0
Lucene	V2.1.0	2	1	0	0
Total		42	37	2	1

Data Corruption Hang Bug Types

- Type 1: Error codes returned by I/O operations directly affect loop strides.
- Type 2: Corrupted data content indirectly affects loop strides.
- Type 3: Improper exception handling directly affects loop strides.
- Type 4: Improper exception handling indirectly affects loop strides.

Data Corruption Hang Bug Types

- Type 1: Error codes returned by I/O operations directly affect loop strides.

Hadoop-8614

```
183 public static void skipFully(InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len);  
            ...  
            ...  
189         len -= ret;  
    } }
```

Data Corruption Hang Bug Types

- Type 1: Error codes returned by I/O operations directly affect loop strides.

Hadoop-8614

```
183 public static void skipFully(InputStream in, long len) ... {  
184     while (len > 0) {  
185         long ret = in.skip(len);  Corrupted InputStream  
...  
...  
189         len -= ret;  
    } }
```



The loop stride (ret) is always 0 when in is corrupted.

Data Corruption Hang Bug Types

- Type 2: Corrupted data content **indirectly affects loop strides.**

HDFS-13514

```
194 BUFFER_SIZE = conf.getInt();
```

```
78  private void readLocalFile(Path path, ...) ... {  
    ...  
84      byte[] data = new byte[BUFFER_SIZE];  
85      long size = 0;  
86      while (size >= 0) {  
87          size = in.read(data);  
    } }
```

Data Corruption Hang Bug Types

- Type 2: Corrupted data content indirectly affects loop strides.

HDFS-13514

```
194 BUFFER_SIZE = conf.getInt(); Corrupted configuration file
```

```
78  private void readLocalFile(Path path, ...) ... {  
     ...  
84      byte[] data = new byte[BUFFER_SIZE];  
85      long size = 0;  
86      while (size >= 0) {  
87          size = in.read(data);  
     } }
```

Data Corruption Hang Bug Types

- Type 2: Corrupted data content indirectly affects loop strides.

HDFS-13514

194 BUFFER_SIZE = conf.getInt(); **Corrupted configuration file**

78 private void ~~readLocalFile~~(Path path, ...) ... {

...

84 byte[] data = new byte[BUFFER_SIZE];

85 long size = 0;

86 while (size >= 0) {

87 size = in.read(data);

} }

0

empty array

Data Corruption Hang Bug Types

- Type 2: Corrupted data content indirectly affects loop strides.

HDFS-13514

194 BUFFER_SIZE = conf.getInt(); **Corrupted configuration file**

78 private void ~~readLocalFile~~(Path path, ...) ... {

...

84 byte[] data = new byte[BUFFER_SIZE];

85 long size = 0;

86 while (size >= 0) {

87 size = in.read(data);

} }

0

empty array

The loop stride (size) is always 0 when conducting read op on an empty array.

False Negative Example

The loop index, stride or bounds are **only** related to specific application I/O functions.

False Negative Example

The loop index, stride or bounds are **only** related to specific application I/O functions.

HDFS-5438

```
1668 while (!fileComplete) {  
1669     fileComplete = dfsClient.namenode.complete(src,  
                                         dfsClient.clientName, last);  
...  
1689 }
```

False Negative Example

The loop index, stride or bounds are **only** related to specific application I/O functions.

HDFS-5438

```
1668 while (!fileComplete) {  
1669     fileComplete = dfsClient.namenode.complete(src,  
                                         dfsClient.clientName, last); Corrupted block  
...  
1689 }
```

False Negative Example

The loop index, stride or bounds are **only** related to specific application I/O functions.

HDFS-5438

```
1668 while (!fileComplete) {  
1669     fileComplete = dfsClient.namenode.complete(src,  
                           dfsClient.clientName, last); Corrupted block  
...  
1689 }
```

Application
function

False Positive Example

Hadoop v2.5 BlockReaderLocal.java

```
472 private int readWithBounceBuffer(  
        ByteBuffer buf...) ...{  
481     do {  
...  
502         bb = drainDataBuf(buf);  
512     } while (buf.remaining() > 0);  
...  
514 }
```

Check bounds

```
277 private int drainDataBuf(  
        ByteBuffer buf) {  
...  
286     buf.put(dataBuf);  
...  
291 }
```

Forward index

False Positive Example

Hadoop v2.5 BlockReaderLocal.java

```
472 private int readWithBounceBuffer(  
        ByteBuffer buf...) ...{  
481     do {  
...  
502         bb = drainDataBuf(buf);  
512     } while (buf.remaining() > 0);  
...  
514 }
```

Check bounds

```
277 private int drainDataBuf(  
        ByteBuffer buf) {  
...  
286     buf.put(dataBuf);  
...  
291 }
```

Forward index

- The **forwarding-index** Java APIs and the **checking-bounds** Java APIs are located in **different** application function.

Conclusion

- DScope is a new data corruption hang bug detection tool for cloud server systems.
 - Combines candidate bug discovery and false positive filtering.
 - Evaluated over 9 cloud server systems and detects **42** true data corruption hang bugs including **29** new bugs.

Acknowledgements

- DScope is supported in part by NSF CNS1513942 grant and NSF CNS1149445 grant.

Thank you