

Automatically Detecting Risky Scripts in Infrastructure Code

Ting Dai
IBM Research
ting.dai@ibm.com

Grzegorz Koper
IBM GTS
grzegorz.koper@pl.ibm.com

Alexei Karve
IBM Research
karve@us.ibm.com

Sai Zeng
IBM Research
saizeng@us.ibm.com

ABSTRACT

Infrastructure code supports embedded scripting languages such as Shell and PowerShell to manage the infrastructure resources and conduct life-cycle operations. Risky patterns in the embedded scripts have widespread of negative impacts across the whole infrastructure, causing disastrous consequences. In this paper, we propose an analysis framework, which can automatically extract and compose the embedded scripts from infrastructure code before detecting their risky code patterns with correlated severity levels and negative impacts. We implement SecureCode based on the proposed framework to check infrastructure code supported by Ansible, i.e., Ansible playbooks. We integrate SecureCode with the DevOp pipeline deployed in IBM cloud and test SecureCode on 45 IBM Services community repositories. Our evaluation shows that SecureCode can efficiently and effectively identify 3419 true issues with 116 false positives in minutes. Among the 3419 true issues, 1691 have high severity levels.

CCS CONCEPTS

• **Computer systems organization** → **Reliability; Availability**; • **Software and its engineering** → **Software testing and debugging**.

KEYWORDS

Infrastructure-as-Code, Ansible, Shell, PowerShell, static analysis, availability, reliability, performance, security

ACM Reference Format:

Ting Dai, Alexei Karve, Grzegorz Koper, and Sai Zeng. 2020. Automatically Detecting Risky Scripts in Infrastructure Code. In *ACM Symposium on Cloud Computing (SoCC '20), October 19–21, 2020, Virtual Event, USA*. ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3419111.3421303>

1 INTRODUCTION

Cloud computing revolutionizes the way enterprises manage their infrastructure. Cloud providers today host mission critical workloads for thousands of clients in large-scale and distributed data centers. To manage infrastructure spanning across hundreds of thousands of servers, automation is extensively used with improved productivity for infrastructure life-cycle operations such as provisioning and deployment, change management, security compliance management, and event and incident management. Such automation practices are often referred to as Infrastructure-as-Code (IaC) while the codified automation procedures are referred to as infrastructure code. In recent years, many IaC tools and techniques have been developed and adopted to facilitate the practice, such as Ansible [12], Chef [14], Puppet [16], and Terraform [19]. Ansible, a Redhat branded tool, is one of the mainstream automation tools for infrastructure provisioning and configuration management. It is widely used by enterprise clients.

Infrastructure code, such as Ansible playbook and Chef cookbook, codifies mission critical operations. Not only does it make sure infrastructure is managed for availability, reliability, and security compliance, but also it guarantees that the running applications on the infrastructure are healthy with desirable performance. Modern IaC tools have their own programming paradigm to write infrastructure code, e.g., YAML for Ansible playbooks and Ruby for Chef cookbooks. They also support embedded scripting languages such as Shell and PowerShell to either manage the infrastructure

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SoCC '20, October 19–21, 2020, Virtual Event, USA

© 2020 Association for Computing Machinery.

ACM ISBN 978-1-4503-8137-6/20/10...\$15.00

<https://doi.org/10.1145/3419111.3421303>

resources like operating systems, storage and networks or interact with applications to execute the automation procedures. Risky patterns in infrastructure scripts introduce bugs and expose vulnerabilities which lead to widespread of negative impacts on availability, reliability, performance, and security across the entire infrastructure composed of hundreds of thousands of servers, causing disastrous consequences. Example risky scripts such as `Remove-Partition-DriveLetter 'C'` which removes the disk could result in business disruption when they are invoked in the production environment. Coding patterns with infinite loops can drain the computation resource, severely degrading the performance of production workloads. An unauthorized user privilege escalation or unprotected storage of credentials make the infrastructure vulnerable to cyber-attacks.

Unfortunately, different from traditional software programs like Java and C/C++ which are equipped with mature bug/vulnerability detection, testing, and verification tools [23, 24, 27, 31–33, 43], existing techniques and practices for infrastructure code and embedded scripts are rudimentary [25]. It gets more challenging when the modern code development shifts to community-based approach, where a team of contributors have mixed skills, experiences, and responsibilities. Particularly, most contributors are system administrators who lack the same level of understanding and debugging support compared with software developers [44]. Studies [2] have shown that nearly 75% of system downtime is caused by human errors. Many service outage incidents are caused by mistakes made by system administrators [1]. For example, in 2017, the Amazon S3 service became unavailable due to a removal command invoked by a system administrator who inadvertently removed a large set of servers, affecting other AWS services in the US-EAST-1 region that rely on S3 for storage, such as S3 console, EC2 new instance launches, and AWS Lambda [8]. This service disruption lasted for five hours [4] with financial loss of \$150 million [5].

State-of-the-practice script checking tools, such as ShellCheck [29] and PSScriptAnalyzer [35] have brought to light the opportunity of conducting static checking over the infrastructure scripts. Those generic script-analyzers report issues in the scripts by checking their formats and syntaxes. However, without correlating the identified issues with their risky behaviors, users (e.g., system administrators) have no understanding about how the risks manifest in the production environment, what the potential business consequences of those risks have, and how severe those negative consequences are. Bridging the gap between generic script-analyzers and business consequence is one of the motivations of this work.

In this paper, we propose an automated analysis tool to identify the risky scripts in the infrastructure code, which

cause availability, reliability, and performance issues or impose security vulnerabilities. Our analysis tool is static without running the target infrastructure code; thus it achieves high code coverage than the dynamic detection approach which can be interrupted by unexpected runtime errors. Moreover, our tool is platform agnostic which requires no system-specific or platform-specific knowledge, such as OS version and kernel version. To realize efficient analysis, we design a scalable analysis tool, including structured representation tree generation, script detection and composition to identify, transform and compose scripts from the original infrastructure code. The composed scripts are then checked by the open source script analyzers. The analysis results are assigned with severity levels considering business configurations and consequences. The locations of the identified issues in the composed scripts are traced back to the locations of the original source code along with the reported issues.

Our work makes the following contributions:

- We design a generic analysis framework to identify risky scripts in infrastructure code. Such product is what the current business market lacks, which has a high internal demand in IBM.
- We generate our risky code knowledge base with severity levels and business impact categories after empirically studying all 409 rules from ShellCheck and PSScriptAnalyzer.
- We implement a real-world solution based on the proposed framework, i.e., SecureCode, to check infrastructure code supported by Ansible, one of the industry mainstream automation tools.
- We integrate SecureCode with the DevOp pipeline deployed in IBM cloud and test it on 45 IBM Services community repositories. Our evaluation shows that SecureCode can identify 3419 availability, reliability, performance and security issues residing in the infrastructure code, within which 1691 have high severity levels.

The rest of the paper is organized as follows. Section 2 describes the design of our analysis framework. Section 3 presents SecureCode's implementation using our analysis framework to detect risky code in Ansible playbooks. Section 4 shows the experimental evaluation. Section 5 discusses the future work. Section 6 compares our work with related work. Finally, the paper concludes in Section 7.

2 SYSTEM DESIGN

This section discusses our design goal, followed by the overview of our analysis platform. We then describe the individual analysis modules.

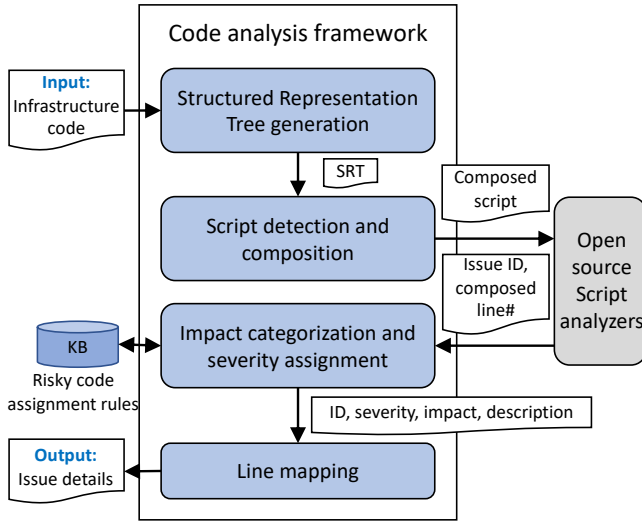


Figure 1: System architecture and program execution flow.

2.1 Design Goal

Our goal is to build a generic, efficient and configurable analysis system.

Generic. Our analysis system is proposed to be generic, which is applicable to most infrastructure code. We provide programming interfaces for users to detect and compose risky scripts in different infrastructure code.

Efficient. Because of the static property, our system does not have runtime delays, e.g., task A can not start until task B finishes. Moreover, without being interrupted by unexpected runtime errors, we can achieve high code coverage.

Configurable. We support user-defined configurations to better output identified issues with respect to business clients' compliance requirements.

2.2 System Overview

To reach the aforementioned design goal, our analysis platform comes with the architecture and execution flows shown in Figure 1. Our system takes an infrastructure code repository as input. It first generates structured representation trees (SRTs) for the repository. By traversing each SRT, it composes script contents by removing template renderings after identifying that the infrastructure code invokes script executions. After passing the composed scripts to script-analyzers, it receives the checking results with issue IDs and the line numbers of the composed scripts where the issues reside. To satisfy configurable business compliance requirements, it assigns the severity level and potential impact for each identified issue based on pre-defined assignment rules retrieved from the knowledge base, and maps those issues

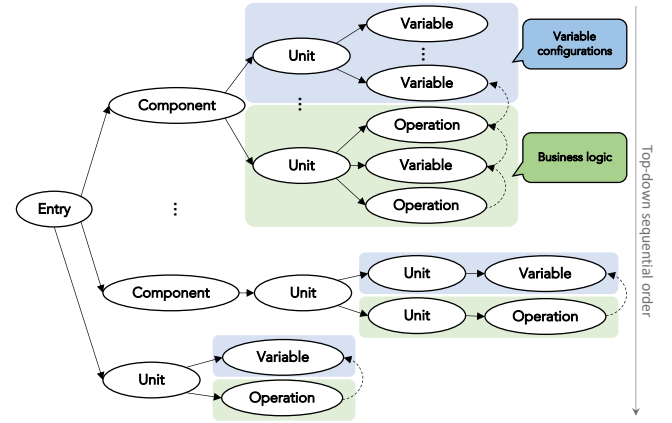


Figure 2: The structured representation tree. The “→” represents the containment relationship while the “- ->” represents the dependency flow.

to the corresponding infrastructure source-code lines before returning them back to the user.

2.3 Structured Representation Tree Generation

What is an SRT. An SRT is a tree data-structure to represent the syntax abstraction of the infrastructure code. It helps to extract and organize the key information in the infrastructure code and exclude details in order to ease the process of accurate script extraction. As shown by Figure 2, starting from the entry node, each SRT contains multiple components. A component consists of multiple units. A unit indicates variable configurations or represents business logic operations. A unit can also contain other units. Operations and variables are dependent on each other—an operation uses the statically configured variables, while a variable can be assigned with the value of an operation's runtime output.

Figure 2 shows that units, variables and operations are necessary in an SRT. But components are not. It is especially common when infrastructure code only contains a few automation steps and the system administrator puts all steps in a single automation file.

Top-down sequential order. All the operations in an SRT are top-down sequentially ordered. So are the units and components. For example, in Figure 2, the variable configuration unit is executed before business logic unit. When constructing an SRT, the ordering dependency between different execution steps in the same single automation file can be easily extracted—the topmost executes earliest. However, for the implicit ordering dependency, their execution sequence is defined inside of the IaC platforms rather than in the infrastructure code. The implicit ordering usually happens when referencing a statically configured variable var in




Figure 3: The happens-before relationship in infrastructure code. The “ \rightarrow ” represents the containment relationship while the “ \Rightarrow ” represents the happens-before relationship.

an operation op. We sculpture the ordering of their (i.e., var’s and op’s) corresponding units by following the “reference after define” policy.

SRT forests. An infrastructure code repository contains multiple entry nodes, representing multiple automation procedures, denoted by a forest of SRTs. The ordering dependency of those SRTs is not always available or required. For example, an administrator can trigger an install procedure before a delete procedure or vice versa. However, the automation procedures’ ordering does not affect the generation of their SRTs.

Merged SRTs. In other cases where two automation procedures have a happens-before relationship explicitly shown in the infrastructure code, we merge the two SRTs into one by introducing a new entry node and creating two containment flows from the new entry node to the two procedures. The first-executed procedure is on the top part of the merged SRT while the latter-executed one is on the bottom part. For example, as shown in Figure 3, *Procedure_A* happens before *Procedure_B*, we merge them into one SRT, where the top-down sequential ordering guarantees the happens-before relationship.

When the happens-before relationships involve more than two procedures and all procedures are partially ordered, we first categorize them in different sets. Each set contains part of the procedures which can be ordered in a determined way. We then merge the procedures from the same set into an SRT. For example, we have three procedures, P_1 , P_2 , and P_3 , which follow the happens before relationships as $P_1 \xrightarrow{HB} P_2$ and $P_1 \xrightarrow{HB} P_3$. We merge P_1 and P_2 into one SRT, and merge a duplicated P_1 and P_3 into another SRT. This also applies to the case when all procedures are totally ordered, e.g., $P_1 \xrightarrow{HB} P_2 \xrightarrow{HB} P_3$. We just need to merge all of them into the same SRT.

SRT construction. As shown in Algorithm 1, our SRT generation follows the Breadth-First Search (BFS) procedure. We first use the `getContainedObj()` function to retrieve all

Algorithm 1: SRT construction

input : infrastructure code repository r
output : a set of SRTs S

```

1 begin
2    $Q \leftarrow \text{QUEUE}(), V \leftarrow \emptyset$   $\triangleright$  initialize the queue and visit set
3    $V \leftarrow V \cup \{r\}$   $\triangleright$  add  $r$  in  $V$ 
4    $P \leftarrow \text{GETCONTAINEDOBJ}(r)$   $\triangleright$  get contained objects
5   for  $p \in P$  do
6      $\text{ENQUEUE}(Q, p)$ 
7   while  $Q \neq \emptyset$  do
8      $c \leftarrow \text{DEQUEUE}(Q)$ 
9      $\text{node}_c \leftarrow \text{PARSEOBJ}(c)$   $\triangleright$  create SRT node
10    if  $\text{parent}[c] = \emptyset$  then  $\triangleright$  not contained
11       $S \leftarrow S \cup \{\text{node}_c\}$   $\triangleright$  add root node in  $S$ 
12     $N \leftarrow \text{GETCONTAINEDOBJ}(c)$   $\triangleright$  get contained objects
13    for  $n \in N$  do
14      if  $n \notin V$  then
15         $\text{ENQUEUE}(Q, n)$ 
16         $\text{parent}[n] = c$ 
17         $\text{children}[c] \leftarrow \text{children}[c] \cup \{n\}$   $\triangleright$  create
18          containment relationship
19     $V \leftarrow V \cup \{c\}$ 
20   $O \leftarrow \text{ORDER}(S)$   $\triangleright$  get ordered lists for all objects
21  for  $R \subset O$  do
22     $e, \text{node}_e \leftarrow \text{EMPTYNODE}()$   $\triangleright$  create empty object, node
23    for  $\text{node}_r \in R$  do
24       $S \leftarrow S \setminus \{\text{node}_r\}$   $\triangleright$  remove node from  $S$ 
25       $\text{children}[e] \leftarrow \text{children}[e] \cup \{r\}$   $\triangleright$  create
26        containment relationship
27     $S \leftarrow S \cup \{\text{node}_e\}$   $\triangleright$  add new root node in  $S$ 
28  return  $S$ 

```

the automation procedures in the root directory r (line #4) and add them in the queue (line #6). We then create a set of SRTs for all the procedures in a while loop (line #7-18). Specifically, in each iteration, we poll the head c from the queue (line #8) and use the `parseObj()` function to create an SRT node node_c to abstract the syntax of object c (line #9). If object c is not contained by others, which means it is an independent automation procedure, component or unit, we consider node_c as the SRT root node and add it in the tree set S (line #10-11). Next, we use the `getContainedObj()` function to extract all the contained objects in c (line #12). For each contained object n , if not visited, we add it in the queue (line #15) and create a containment relationship from object n to object c (line #16-17). In the end of the iteration, we add object c in the visit set (line #18). Our SRT generation algorithm scans the whole repository layer by layer until all automation objects (e.g., files, folders) are abstracted in an SRT node (i.e., entry, procedure, component, unit, variable and operation).

An SRT node contains a list of attribute fields, including `type`, `_file_` and `_line_`. If the node's type is *operation*, this node contains the content of the operation in the form of module-command pairs, e.g., "shell: ls -al". If the node's type is *variable*, it contains the variable name with the configured value, e.g., "user: root". The `_file_` attribute stores the source code file path while the `_line_` attribute stores the first line number of an operation or a variable. In an infrastructure code repository, an SRT node is uniquely identified by the $\langle_file_,_line_ \rangle$ tuple.

To handle happens-before relationships among a forest of SRTs in S , we first extract the ordered lists¹ O for all of them (line #19). Each list R contains a sequential order among a set of SRTs, i.e., $R \subseteq S$, $R = \{srt_1, srt_2, \dots, srt_n\}$, $srt_1 \rightarrow srt_2 \rightarrow \dots \rightarrow srt_n$. For each list R , we use the `emptynode()` function to create a new empty object e with its corresponding SRT node $node_e$ (line #21) and add $node_e$ in the tree set S (line #26). For each SRT node $node_r$ in the ordered list R , we first remove it from the tree set S (line #23) and then create a containment relationship between $node_e$ and $node_r$ (line #24).

We should note that the interface functions, i.e., `parseObj()` and `getContainedObj()`, highlighted in Algorithm 1, should be implemented specifically to support the target infrastructure code's programming paradigm. Furthermore, the `getContainedObj()` function must guarantee that its returned values/objects are sequentially ordered—the leftmost executes earliest². In such way, the sequential order is preserved among all the leaves in an SRT from left to right. The detailed implementation of the two functions will be discussed in Section 3. The other functions, including `queue()`, `enqueue()`, `dequeue()`, `order()`, `emptynode()` are commonly used library functions. We do not explicitly describe their implementation details due to page limit.

2.4 Script Detection and Composition

Script invocation in infrastructure code. Modern IaC platforms provide their own programming paradigm, syntax and libraries to support script invocations. For example, Ansible supports Shell/PowerShell script invocations in YAML playbooks using the following Ansible modules: `command`, `shell`, `script`, `raw`, `install`, `before_install`, `win_command`, and `win_shell`. Chef supports Shell/PowerShell script invocations in Ruby cookbooks using the following Chef resources: `execute`, `script`, `powershell_script`, `bash`,

`csh`, and `ksh`. Puppet supports Shell/PowerShell script invocations in DSL files using the Puppet `exec` resource. Terraform supports Shell/PowerShell script invocations in TF files using the `local-exec` and `remote-exec` provisioners.

Script detection. To detect whether the infrastructure code invokes scripts, we traverse the generated SRTs, obtain their leaf nodes, and check whether the operation leaf nodes use the aforementioned script-related IaC libraries, e.g., Ansible modules, Chef resources, Puppet resources, and Terraform provisioners. As shown in Algorithm 2, our script detection module provides the `hasScript()` interface function (line #8) to check the existence of scripts in SRTs.

After detecting that the operation nodes involve the script invocation, simply dumping raw scripts from the operation nodes using the `extractScript()` interface function (line #9) is not enough, due to the template rendering in modern IaC platforms. Those raw scripts can be embedded with templated variables in Jinja2, Django, JSP, PHP, etc. To achieve accurate and modularized script checking, we need to have interfaces to compose the scripts in correct scripting language syntax and format, before passing them to the script-analyzers.

Variable map. In infrastructure code, variables are either configured in configuration files or defined by run-time jobs' execution results. The statically configured variables are in an SRT's variable leaf nodes associated with default values, while the dynamically assigned variables are in an SRT's operation leaf nodes associated with specific operations. Our system provides the `extractVar()` function to extract variable-value pairs from all variable leaf nodes and append them in the variable map M . The variable map is updated in-time—it is updated at the beginning of each loop iteration (line #6)—to guarantee the define-reference order in two aspects:

First, the variable map contains the configured variables' values, which can be retrieved and referenced in the later script composition. The sequential order in SRT generation makes sure that variable nodes are on the left side of operation nodes in the same component. Thus, the variable nodes have higher priority to be processed earlier.

Second, a dynamically defined variable in an operation node can only be referenced by the later operations. For example, there are three operation nodes op_1 , op_2 , and op_3 and leftmost executes earliest. If a variable var_2 is defined in op_2 , the in-time updated variable map makes sure that var_2 can be referenced in op_3 but not in op_1 , when we compose the scripts in those operations.

We should note that IaC tools handle lexically scope variables by overriding the old value with the new value globally, which is supported by our variable map.

¹The ordered lists are provided by user-specific inputs or pre-requisites, defined in the README file.

²It is the implementation version of SRT's top-down sequential order.

Algorithm 2: script detection and composition

```

input : a set of SRTs  $S$ 
output: a set of composed scripts  $C$ 

1 begin
2   for  $node_r \in S$  do                                ▶ SRT root node
3      $L \leftarrow \emptyset, M \leftarrow \emptyset$ 
4      $L \leftarrow \text{GETLEAVES}(node_r, L)$ 
5     for  $l \in L$  do
6        $M \leftarrow M \cup \text{EXTRACTVAR}(l)$                 ▶ var map
7       if  $l.type = 'operation'$  then                    ▶ operation leaf node
8         if  $\text{HASSCRIPT}(l)$  then
9            $c \leftarrow \text{EXTRACTSCRIPT}(l)$ 
10           $c \leftarrow \text{COMPOSESCRIPT}(c, M)$ 
11           $C \leftarrow C \cup \text{REFORMATTEMPLATED}(c)$ 
12   return  $C$ 

13 function  $\text{GETLEAVES}(node_r, L)$ 
14   if  $children[r] = \emptyset$  then                        ▶ leaf node
15     return  $\{node_r\}$ 
16   for  $n \in children[r]$  do
17      $L \leftarrow L \cup \text{GETLEAVES}(node_n, L)$  ▶ union children's leaves
18   return  $L$ 

19 function  $\text{COMPOSESCRIPT}(c, M)$ 
20    $c' \leftarrow c$ 
21    $V \leftarrow \text{EXTRACT}(c)$                             ▶ extract the referenced variables
22   for  $var \in V$  do
23     if  $\text{CONTAINSKEY}(M, var)$  then
24        $val \leftarrow \text{GET}(M, var)$ 
25        $c \leftarrow \text{REPLACE}(c, var, val)$ 
26   if  $c \neq c'$  then
27     return  $\text{COMPOSESCRIPT}(c, M)$                     ▶ compose recursively
28   return  $c$ 

```

Script composition. As shown in Algorithm 2, the script composition (line #19-28) is a process of replacing a templated variable var in raw scripts with its value val extracted from the variable map M . To handle cases with nested variable references, our $\text{composeScript}()$ function recursively conducts the replacement (line #27) until all the referenced variables are replaced by their defined values or until the script cannot be further composed, i.e., the variable map M does not contain the value for a referenced variable.

For example, we have a raw shell script as `rm -rf {{dir}}`, where the `{{dir}}` variable is in Jinja2 format. The variable map contains the key-value pairs as `dir={{path}}`, `path='/'`. We call the $\text{composeScript}()$ function to extract the `dir` variable (line #21), query the map (line #23), and retrieve its value as `path` (line #24). We update the script by replacing `{{dir}}` with `{{path}}` and get `rm -rf {{path}}` (line #25). A recursive function call is invoked to check whether the script can be further composed (line #27). After

a second-round of extraction, query, and replacement, the script is updated as `rm -rf /`. This is the final composed script, since there is no referenced variable embedded in it.

Templated variable reformat. The composed scripts may still contain templated variables which are undefined in the infrastructure code, i.e., the variable map M does not contain the value for a referenced variable. Passing those syntactically problematic scripts to script-analyzers results in inaccuracy and false positives. To remove the templating language from the composed scripts while still preserving the variable reference, we use the $\text{reformatTemplated}()$ interface function (line #11) to reformat those templated variables in correct script syntax.

We should note that the $\text{reformatTemplated}()$ interface function along with others, i.e., $\text{extractVar}()$, $\text{hasScript}()$, $\text{extractScript}()$, and $\text{extract}()$, highlighted in Algorithm 2, should be implemented specifically for the target infrastructure code due to the corresponding programming paradigm. The implementation details of these functions will be discussed in Section 3. Other functions, such as $\text{containsKey}()$, $\text{get}()$, and $\text{replace}()$ are commonly used library functions. Their implementation is omitted in the paper.

2.5 Business Impact Categorization and Severity Assignment

Our analysis system leverages the open-source state-of-the-practice script-analyzers such as ShellCheck and PSScriptAnalyzer to conduct risky code checking on the composed scripts. However, the issues reported by those script-analyzers are generic. They label each issue with the type of “style”, “info”, “warning”, or “error” without clearly showcasing the potential business impacts of those issues and how severe those impacts are, when the reported issues’ risky behaviors manifest in the production environment. To satisfy business compliance requirements, our analysis system regulates the identified issues with impact categorization and severity assignment.

We conduct an empirical study on the existing script-analyzers’ rulesets, including 345 rules from ShellCheck and 64 rules from PSScriptAnalyzer. We study the code patterns checked by each rule to identify what risky behaviors and consequent negative impacts the problematic code can have. We categorize the issues based on the potential impacts as “none” issues, “security” vulnerabilities, “availability” issues, “performance” issues or “reliability” issues. For each category of risky code, we assign severity levels accordingly.

None risky behavior. For the problems labeled as “style” by the script-analyzers, they usually do not have risky behaviors. We mark their impacts as “none” and assign their severity levels as “low”. For example, ShellCheck considers that the shell command `echo $(cat foo.txt)` has a style

issue, wherein `echo` is useless. This command does not have any risky behavior³ but a simplified version is preferred, i.e., `cat foo.txt`.

Security. For the issues which contain security vulnerabilities, we assign their severity levels as “high”, based on IBM business security requirements [6, 18]. For example, in the shell command `find . -name '*.txt' -exec sh -c 'echo "{}" \;`, the filename is passed in by an injected shell string (i.e., “{}”). Any shell meta-characters in the filename can be interpreted as part of the script, which allows arbitrary code execution exploitation.

Availability. For the issues which can cause the system to be unavailable, leading to potential service outages [11], we assign their severity levels as “high”. For example, the shell command `rm -rf /` deletes the whole system directory, which severely impacts system availability, causing catastrophic consequences.

Performance. For the issues which degrade system performance, we assign their severity levels as “medium”. For example, the shell command `if ["$(find . | grep 'IMG[0-9]')]"` iterates the entire directory and reads all matching lines into memory before making a decision rather than stopping at the first matching line. It prolongs the infrastructure code execution, which can decrease system performance. When a performance issue involves abnormally high CPU or memory usage, such as an infinite loop, it can significantly impact the whole system, making the system become partially or entirely unavailable. However, without runtime information and specific user input, we currently cannot decide whether the risky code has performance issues or performance-induced availability issues.

Reliability. For the issues which make the scripts’ output become unreliable, further affecting infrastructure life-cycle operations, we measure them from a context specific perspective. Specifically, we extract the infrastructure operation criticality from operation names, automation file names, and comments. We classify the operations as “critical” if they are related to reboot, restart, update, backup, database, service, etc. We classify the operations as “moderate” if they are related to file read/write, word count, logs, etc. The other operations are considered as “trivial”, such as printing messages on a terminal. For the unreliable scripts which are related to the infrastructure critical, moderate, or trivial operations, we consider their severity levels as “high”, “medium”, or “low”, respectively. For example, the shell command `ps ax | grep python` is unreliable. It not only searches python processes but also tries to match against other fields, such as if the user’s name was *pythonguru*. When the command’s output

is used in a subsequent reboot operation, all the processes created by the *pythonguru* user get restarted, affecting all the running workloads. Thus, this unreliable script’s severity level should be assigned as high. However, when this shell command’s output is not used or simply printed on a terminal, this unreliable script’s severity level should be assigned as low.

We generate a configurable knowledge-base to store the impact and severity information for each rule. Each rule has an ID in the format of analyzer-abbreviation and digit numbers. For the ShellCheck rules, their IDs are in the range of SC1000-SC9999, while for the PSScriptAnalyzer rules, their IDs are in the range of PS1000-PS9999. Our knowledge-base also contains short and detailed rule descriptions associated with each rule ID. A short description is a message containing both pattern and impact information to briefly summarize the risky code. A detailed discussion with examples and specific scenarios are explained in a description file. We should note that, the knowledge-base is configurable—users can easily turn on/off a rule or adjust a rule’s severity level. In our evaluation, we include all the rules from the knowledge-base with the severity levels assigned by our empirical study results.

2.6 Output Format with Line Mapping

Our analysis system provides a user-friendly output for each detected issue, including ① a rule ID, ② an issue type inherited from the script-analyzers, ③ a severity level, ④ the potential business impact, ⑤ a short description for the matched rule, ⑥ a file link for the rule’s detailed descriptions, ⑦ the content of the original risky script with optional template rendering, ⑧ the content of the composed risky script, and ⑨ the source code file path with ⑩ the line number where the risky script reside. Example output is discussed in Section 4.1.

The above first six fields can be retrieved from the knowledge-base. The original and composed risky code contents can be derived from our script composition module. The source code file path can be obtained from the `_file_` attribute in the operation SRT node where the script is invoked. The risky script line number can be obtained from the following formula “ $L = L_0 + i - 1$ ”, where L_0 is the line number of the embedded scripts stored in the `_line_` attribute of the SRT node; i is the line number of the risky code reported by script-analyzers in the composed script.

3 IMPLEMENTATION

We come up with a real-world solution, i.e., SecureCode, using our analysis platform to identify the risky scripts in Ansible infrastructure code, i.e., Ansible playbooks. To support the programming paradigm in Ansible playbooks, which are

³It might cause a millisecond overhead with negligible performance degradation.

written in YAML with Jinja2 template rendering, we propose our own template parser as well as reusing parsing functions in the Ansible-lint [13] tool to implement SecureCode, especially for implementing the interface functions in Section 2, including the `parseObj()` and `getContainedObj()` functions in Algorithm 1, the `extractVar()`, `hasScript()`, `extractScript()`, `reformatTemplated()`, and `extra-ct()` functions in Algorithm 2.

3.1 SRT Construction Implementation

In Ansible, a *playbook* represents an automation procedure. A *playbook* contains different *roles* as the SRT entry node contains multiple components. A *role* contains *vars*, *defaults*, and *tasks* directories as a component node contains different unit nodes. Variables are configured in the *vars* and *defaults* directories while the business logic is represented by the operations defined in the *tasks* directory.

Our `parseObj()` function's implementation takes the structured directory of Ansible playbooks into consideration, and creates the SRT entry nodes, component nodes and unit nodes accordingly. To create SRT variable nodes and operation nodes, we parse YAML files via the `parse_yaml_linenumbers()` function from Ansible-lint, and extract operation invocation statements and variable configuration statements.

We implement the `getContainedObj(c)` function to extract the containment relationships among Ansible playbooks, roles, variables, and operations in two ways. If the automation object *c* is a folder (e.g., the root directory, the *roles* directory and the *tasks* directory), the contained objects are all the files and subdirectories in this folder. Particularly, when the automation object *c* is the *roles* directory, the `getContainedObj(c)` function returns *defaults*, *vars*, and *tasks* in order, to make sure that the variable nodes are on the left side of operations nodes in an SRT. If the automation object *c* is a file (e.g., a *main.yml* file in the *tasks* directory), we extract all the contained objects by this file via the `play_children()` function from Ansible-lint. For example, if the *main.yml* file includes the *before.yml* file, then we extract *before.yml* as the contained object by *main.yml*.

3.2 Script Detection and Composition Implementation

In Ansible, every operation is conducted by specific Ansible modules or user-specified plugins. An Ansible operation is represented by a *module-parameter* pair or a *plugin-parameter* pair. For example, in the operation of "shell: cat {{fp}}.txt", shell is an Ansible module to execute shell commands while cat {{fp}}.txt is the command content to be executed.

We implement the `hasScript()` function to detect scripts by checking whether the module-parameter pairs in SRT

operation nodes are script related. The script-related Ansible modules include the `command`, `shell`, `script`, `raw`, `install`, `before_install`, `win_command`, and `win_shell` modules. Our `extractScript()` function then retrieves the parameter content from the script-related module-parameter pairs as the raw scripts. Currently, SecureCode does not check user customized script plugins during script detection and extraction, based on two reasons. First, the script-related modules provided by Ansible are abundant for users to execute Shell or PowerShell commands in infrastructure code. Second, user-specified shell plugins are not guaranteed to be compatible with state-of-the-practice script analyzers.

In Ansible, a statically configured variable is represented by a *variable-value* pair while the dynamically defined variable is represented by a *variable-operation* pair. We implement the `extractVar()` function to extract both static and dynamic variables. Specifically, the static variable-value pairs extracted from the variable nodes can be directly appended into the variable map *M*. As for the dynamic variable-operation pairs extracted from the operation nodes, we retrieve the variable's value based on the module type in the operation. If the operation contains script-related modules, such as `shell`, `win_shell`, we retrieve the command content in this operation as the variable's value. Otherwise, we simply assign a default value to this variable. In the current implementation, SecureCode does not understand all the Ansible modules, thus it cannot transform those script-unrelated operations into scripting languages.

Ansible uses Jinja2 template rendering to enable dynamic expression and access to variables. Those templated variables can be embedded in the raw scripts. We implement the `extract()` function with a template parser to extract the Jinja2 variables from the raw scripts during script composition. Our template parser contains a set of regular expressions for different matching and extraction purposes.

We first check whether the Jinja2 variables contain complex structures, i.e., filters and concatenations. We convert them into multiple simple ones via a set of well-designed regexes. For example, we convert `{{(disk_image.path | dirname) + '/' + disk_name}}` into `{{disk_image}}/{{disk_name}}`. Those simplified Jinja2 variables only contain letters, numbers, and the underscore symbol, surrounded by double braces, i.e., `{{}}`.

Then, the `extract()` function can get a list of template-free variables by removing the surrounding double braces in the simplified Jinja2 variables, e.g., `disk_images`, `disk_name`.

We implement the `reformatTemplated()` function to check whether the composed script still contains templated variables and reformat them in scripting language. Specifically, we use the `"{{[0-9A-Za-z_\\.]*}}"` regex to search the composed script and replace the double braces with the dollar symbol. For example, if the `{{file}}` variable is not defined

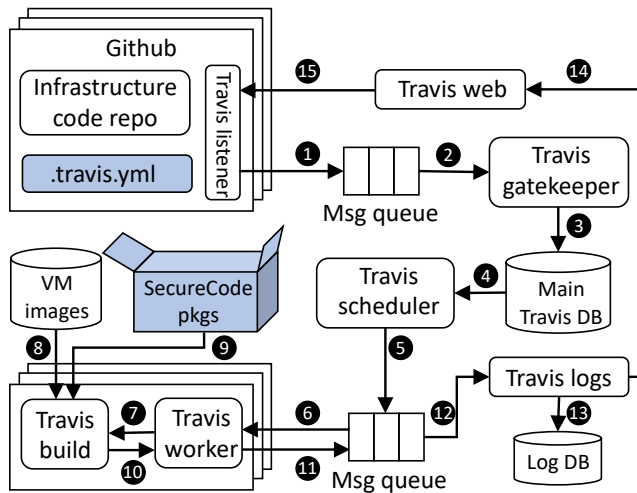


Figure 4: Travis flow. The `.travis.yml` file contains a Travis job configuration, including SecureCode’s pre-installation, installation and execution steps.

in the infrastructure code but still referenced in the shell script `cat {{file}}`, the `reformatTemplated()` function generates a syntactically correct script variable `"$file"` to replace `{{file}}` and updates the script as `cat "$file"`.

4 EVALUATION

We implement SecureCode in Python language, using parsing functions in Ansible-lint v4.2.0, Shell script analyzer ShellCheck v0.7.0, and PowerShell script analyzer PSScript-Analyzer v1.18.3. We push SecureCode in IBM’s enterprise github, integrate it with DevOp pipeline deployed in IBM Services using Travis CI [20], shown by Figure 4. We test SecureCode in 45 IBM Services community github repositories in the Call-for-Automation (C4A) contest hosted by the IBM Continuous Engineering team.

For each github repository, we create a Travis job to conduct SecureCode’s checking on the automation files with the extension of `.yaml`, `.yml`, `.sh`, and `.ps1`, in the master branch. A Travis job is defined by the `.travis.yml` configuration file, which contains three steps: pre-installation, installation, and a running command to invoke SecureCode. In our experiments, all Travis jobs are running on a Ubuntu v16.04.5 VM with kernel v4.19.52, provisioned on IBM Cloud.

4.1 Output and Case Study

As shown by Figure 4, every target Github repository is associated with a webhook, i.e., Travis listener. With each new Github commit, SecureCode is invoked to check the repository code in a VM (Step 1 to 9). After SecureCode finishes checking, the checking status returns to the corresponding Github commit (Step 10 to 15). If SecureCode identifies an

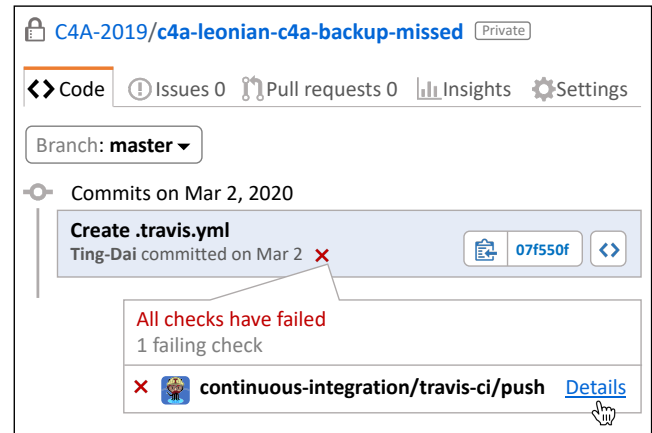


Figure 5: SecureCode’s checking status for the Github commits. To view SecureCode’s output, click on “Details”.

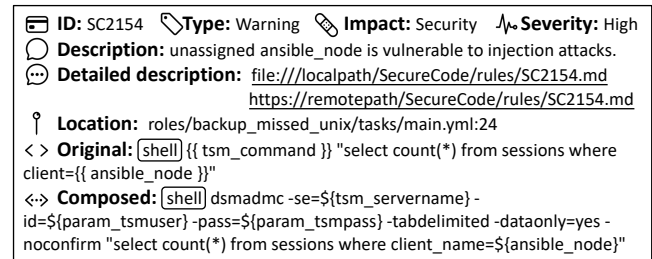


Figure 6: A snippet of SecureCode’s detection report presented in Travis web GUI.

issue, the commit is associated with a red cross mark (otherwise a green check mark), shown by Figure 5. The developer can click on “details” to see SecureCode’s detection reports.

Figure 6 shows the output of one of the detected issues by SecureCode. The risky code happens at line #24 in the `main.yml` file, where a shell script is invoked. The original shell script contains Jinja2 template rendering variables, i.e., `{{ tsm_command }}` and `{{ ansible_node }}`. SecureCode conducts script composition by replacing `{{ tsm_command }}` with its configured value extracted from the variable map, and then reformats the composed script by replacing double braces with the dollar symbol for those undefined Jinja2 variables, e.g., `{{ ansible_node }}`. The risky pattern of undefined variables matches ShellCheck’s Rule #2154 with the type of warning. SecureCode reports this issue containing security vulnerabilities with severity level as high based on our business compliances in the risky code knowledge-base. The statically unconfigured `ansible_node` variable allows a user to pass any value from a command line when he/she executes the corresponding Ansible playbooks. Without runtime sanity checking, an unexperienced user may falsely

pass wrong values to the infrastructure code. Even worse, a malicious user can inject untrusted inputs to the infrastructure, compromising the system and stealing confidential informations. In this example, `ansible_node` is used in a SQL command, which is vulnerable to SQL injection attacks. SecureCode also provides detailed descriptions about this issue with local file address and remote wiki page.




Figure 7: The statistics of identified issues with corresponding rules in each impact category with different severity levels. The y-axis represents a rule ID while the x-axis is the number of matched issues for a rule.

4.2 Detection Accuracy and Statistics

SecureCode identifies 3535 issues in total from the 45 repositories which contain 1492 automation files. We manually investigate all issues and filter out the false positives, which are 1) either falsely matched the code patterns by the corresponding rules in ShellCheck or PSScriptAnalyzer, 2) or misinterpreted by SecureCode’s script composition module. SecureCode’s detection is accurate with 116 out of 3535 issues as false positives.

The statistics of the 3419 true issues are shown in Figure 7. The number of matched issues for each rule varies drastically. Some rules such as SC2140, SC2086, SC2154, and SC1017 have the highest matched numbers, because their risky code patterns are related to common mistakes users make. We also have observed that one repository can have hundreds lines of code with the same risky pattern. It implies that users especially system administrators have strong coding personalities leading to the same mistakes [7]. We should notice that, there are fewer PowerShell scripts than Shell scripts in our benchmarks; thus the PSScriptAnalyzer rules match fewer issues than the ShellCheck rules, in general.

There are 1204 security issues, 485 reliability issues, and 2 availability issues having the high severity level. Particularly, the Rule SC2154 and SC2086 match most of them. The Rule SC2154 is that unassigned variable is vulnerable to injection attacks, with example discussed in Section 4.1. The Rule SC2086 is that unquoted variable/command causes globbing and word splitting, resulting in unreliable outputs. When those unreliable outputs are used in system critical operations, such as reboot, stop and install, they put the whole system in an unstable state with the potential service downtime or even outage.

SecureCode identifies 247 and 51 medium severity level issues which are related system reliability and performance, respectively. Among all the 298 medium severe issues, the Rule SC2086 match most of them when the corresponding scripts are conducting non-critical operations consuming moderate computation resources, e.g., logging in a loop. It is challenging for SecureCode to differentiate whether a loop can become infinite without specific inputs in a runtime environment. Thus, for all the loop-related issues, we currently mark their severity levels as medium.

SecureCode identifies 1430 issues with low severity level. They 1) either have none risky behaviors, such as Rule SC2034 (i.e., unused variable) and Rule PS1012 (i.e., lines end with whitespace), 2) or are related to trivial operations such as simply printing messages on a terminal.

Baseline comparison. We conduct a comparison experiment by running vanilla Ansible-lint, ShellCheck and PSScriptAnalyzer on the same 45 github repositories as SecureCode. To satisfy language capability in order to get valid

checking results, we run Ansible-lint on all the automation files with the `.yaml` and `.yml` extensions, ShellCheck on all the `.sh` automation files, and PSScriptAnalyzer on all the `.ps1` automation files. Among all the 3535 issues in the 45 repositories detected by SecureCode, 1718 can be detected by ShellCheck and 20 can be detected by PSScriptAnalyzer. The remaining 1797 issues are caused by scripts invoked inside `.yaml` and `.yml` files, which can only be detected by SecureCode. We observe that Ansible-lint can barely detect any script-related issues when they are invoked in different Ansible modules. This is because Ansible-lint is specifically designed to check the formats and best practices of YAML language. As for ShellCheck and PSScriptAnalyzer, they can detect all the issues in `.sh` and `.ps1` files as SecureCode, which is expected. This is because, unlike Shell and PowerShell scripts invoked in the `.yaml` and `.yml` files, the automation procedures listed in the `.sh` and `.ps1` files do not require SecureCode's script detection and composition. SecureCode leverages ShellCheck and PSScriptAnalyzer to conduct script checking, thus the 1718 issues in the `.sh` files can be detected by both ShellCheck and SecureCode while the 20 issues in the `.ps1` files can be detected by both PSScriptAnalyzer and SecureCode.

4.3 Preparation and Detection Time

Preparation time. Before conducting detection on the target github repositories, SecureCode has a preparation stage, including pre-installation and installation. Figure 8 shows SecureCode's preparation time in each Travis job. SecureCode's average preparation time is 136 seconds, including 85 seconds of pre-installation and 51 seconds of installation.

SecureCode's preparation time does not include all the latencies from when a new Travis job is created to when the job starts to execute (Step ① to ⑦ in Figure 4), including queuing delays, scheduling delays, database read/write delays, etc. It also does not contain the environment building delay, which refers to the time of launching a VM with default environment setups [21], such as installing python and git (Step ⑧ in Figure 4). Travis is running in parallel, the aforementioned latencies (or delays) heavily depend on the workloads and are easily affected by network connections. We focus on evaluate the execution time of SecureCode not the latency of the Travis system. Thus, we do not include those delays in our evaluation.

Pre-installation time. SecureCode's average pre-installation time is 85 seconds, with 27 out of 45 Travis jobs pre-installed in 106.50 ± 5.06 seconds and 18 out of 45 jobs pre-installed in 51.54 ± 2.35 seconds. Among all the pre-installation steps, installing PSScripAnalyzer on PowerShell Core consumes the majority (i.e. 54%) of the time on average. The consumption time of this step is also highly fluctuant with the standard

deviation of 27.01, which may caused by the compatibility of using Microsoft products on the Linux system.

Installation time. Triggering SecureCode's installation is easy by operating the `pip install git+https` command. The pip package manager then automatically generates the setup scripts, i.e., `setup.py` from SecureCode's `setup.cfg` file which contains the default setup configurations and commands to build a distribution. It takes the Travis VM 50.63 ± 5.37 seconds on average to install SecureCode, shown in Figure 8.

Detection time. Figure 9 shows SecureCode's detection time and lines of code (LOC) in each repository. In most cases, the detection time is proportional to the LOC with about 1 second spent on analyzing 80 LOC. However, in some repositories, such as *Repo₃*, *Repo₉*, *Repo₁₅*, *Repo₁₈*, *Repo₂₅*, and *Repo₃₇*, the ratio of detection time to LOC is much lower than others. This is because SecureCode's detection time is not only related to LOC, but also related to the amount of script invocations and the complexity of script composition in each repository. We observe that in *Repo₃*, *Repo₁₅*, and *Repo₁₈*, there are fewer script invocations even they have a large number of LOC. Meanwhile, in *Repo₉*, *Repo₂₅*, and *Repo₃₇*, they invoke simple scripts which do not require SecureCode's comprehensive composition step. On the contrary, for repositories containing complicated script invocations, such as *Repo₁₂* and *Repo₄₄*, SecureCode spends more time on script composition; thus the ratio of detection time to LOC in these repositories is higher than others.

4.4 Coding Skill Variation

We use both 1) the number of detected issues and 2) the number of detected issues per LOC (i.e., *ipl* ratio) to evaluate the code quality of each tested repository. As shown in Figure 10, the issue numbers and *ipl* ratios tend to fluctuate utterly, ranging from 0 to 1414, and 0 to 0.45, respectively. The average and median issue numbers are 79 and 6; while the average and median *ipl* ratios are 0.04 and 0.01. Repositories have more (severe) issues or higher *ipl* ratios tend to have lower quality of code. For example, *Repo₁* and *Repo₃₄* are considered to have lower code quality with a large number of issues (545 and 1414, respectively, $\gg 79 > 6$) and higher *ipl* ratios (0.36 and 0.29, respectively, $\gg 0.04 > 0.01$). We also observe that most of issues in these repositories have high or medium severity levels. Another examples are *Repo₂* and *Repo₉*. They have lower code quality with large numbers of issues (691 and 285, respectively, $\gg 79 > 6$), among which most have high severity levels, even with average *ipl* ratios, i.e., 0.04. On the contrary, repositories have less (severe) issues and lower *ipl* ratios tend to have higher quality of code, e.g., *Repo₁₁*, *Repo₁₄*, *Repo₁₆*.




Figure 8: SecureCode's pre-installation and installation time.




Figure 9: SecureCode's detection time and LOC in each repository.




Figure 10: LOC, the number of detected issues, and their ratios in each repository. The ratios are in an ascending order.

The variation of code qualities are caused by different development experience, coding practice, and scripting language expertise. Especially for the C4A contest, not all participants are fully trained, nor are they full-time dedicated on the project for the contest. SecureCode analyzes the submitted infrastructure code and checks whether they contain risky patterns on script invocations. However, judging the submitted code based on SecureCode's detection results is far from enough. Moreover, we never depreciate those repositories with large number of detected issues. They may propose new features and solutions to the Ansible infrastructure, which can be widely used and adopted in future after some refinement and refactoring.

4.5 User Experience

⁴Manual checking means checking the Ansible-embedded Shell/Powershell scripts manually without any tool.

⁵Users can compose their own Shell scripts from Ansible playbooks and then run those composed scripts using ShellCheck.

We have designed quantified metrics to access the user experience from IBM Continuous Engineering team by comparing SecureCode with manual checking ⁴, ShellCheck ⁵ and PSScriptAnalyzer ⁶ in terms of throughput improvement (i.e., LOCs reviewed per person per day) and efficiency improvement (i.e., number of issues to be identified). The pre-production pilot of SecureCode comes back with the feedback that SecureCode achieves an estimated 5x throughput improvement compared with manual check, 2x to 5x throughput improvement compared with ShellCheck, and 2x to 3x throughput improvement compared with PSScriptAnalyzer. SecureCode can identify issues which were overlooked by developers. Specifically, SecureCode can identify 5x more issues than manual check, 2x to 3x more issues than ShellCheck, and 2x to 3x more issues than PSScriptAnalyzer. This user experience matches our baseline comparison results, which showing that SecureCode removes the bottleneck of code governance/review and shortens the time to market.

⁶Users can compose their own Powershell scripts from Ansible playbooks and then run those composed scripts using PSScriptAnalyzer.

5 FUTURE WORK

In this work, we take a first step towards the direction by identifying risky scripts in infrastructure code. The fundamental idea is to bridge the gap between generic state-of-the-practice script-analyzers and business compliances when detecting risky patterns with correlated potential negative consequences in the business infrastructure environment.

SecureCode is a roadmap in this journey. Additional features to be implemented are on our schedule, including supporting more scripting languages, e.g., perl, python, and ruby, and supporting checking infrastructure code in other IaC tools, e.g., Chef cookbooks and Puppet manifests. To support other scripting languages, we need to integrate corresponding script analyzers into our system and one-time empirically study on the out-of-the-box rulesets, e.g., Bandit [37] for Python. To support other IaC tools, specific implementations on the interface functions in Algorithm 1 and Algorithm 2 are required. With the help of open-source libraries, e.g., puppet-strings [17], the implementation work is not a heavy load.

Our knowledge-base has a good coverage on common syntax and intermediate-level semantic bugs. In our pre-production testing, we do not encounter any false negatives. However, SecureCode cannot detect certain types of bugs, e.g., dirty copy-on-write, and fork bomb. Thus, extending the rulesets of SecureCode is also our future work.

6 RELATED WORK

Infrastructure code analysis. Previous work has been extensively done to identify quality concerns in infrastructure code [13, 15, 22, 40, 42]. Those quality analysis tools mainly check style issues, e.g., complex expressions, long statements, and unnamed tasks. However, these quality concerns mostly do not cause risky behaviors manifested in the production environment.

Many infrastructure code analysis approaches have been proposed to identify security, availability and reliability issues. SLIC [38] detects seven security smells in Puppet manifests, including admin by default, hard-coded secret, suspicious comments, etc. FSMoVe [41] identifies ordering violations and missing notifiers in Puppet programs, which can cause the infrastructure become unavailable [3] or unreliable [39]. Model-based testing [26, 30] is proposed to detect non-idempotent and non-convergent automations in Chef recipes and Puppet manifests. Infrastructure code which cannot idempotently converges to a desired state causes reliability issues.

In comparison, our work checks the risky patterns (including availability, reliability, security and performance issues) in Shell and PowerShell scripts embedded in infrastructure

code. We believe our work is complementary to the existing infrastructure code analysis approaches.

Shell and PowerShell script analysis. Recent work has been done to identify bugs and security vulnerabilities in Shell and PowerShell scripts. ABash [34] statically checks common expansion bugs in bash scripts using taint tracking and abstract interpretation. Shill [36] enforces the scripts' access control to system resources via declarative security policies and capability-based sandboxes. Machine Learning based detectors [9, 10, 28] are proposed to detect malicious PowerShell commands, using NLP, CNNs, and LSTM.

Different from existing Shell and PowerShell script checking approaches, which either are specific to certain types of bugs and vulnerabilities [34, 36] or require a large number of training data [9, 10, 28], our work leverages static pattern-based script-checking tools, i.e., ShellCheck [29] and PSScriptAnalyzer [35], to analyze infrastructure embedded Shell and PowerShell scripts and correlates them with business configurations and consequences.

7 CONCLUSION

Risky scripts in infrastructure code have widespread of negative business impacts, leading to disastrous consequences. To identify the risky scripts which cause availability, reliability, security and performance issues, we design an analysis framework to abstract the structured representations from infrastructure code before extracting and composing its embedded Shell and PowerShell scripts. Our analysis framework then leverages generic state-of-the-practice script analyzers, ShellCheck and PSScriptAnalyzer, to conduct script checking, and a configurable knowledge base to categorize business impacts and assign severity levels for the identified issues. We implement SecureCode based on our analysis framework to check Ansible playbooks. We integrate SecureCode with the DevOp pipeline deployed in IBM cloud and test SecureCode on 45 IBM Services community repositories. Our evaluation shows that SecureCode can identify 3419 true issues with 116 false positives, with the average pre-installation and installation time of 136 seconds, and average detection time of 1 second per 80 LOC. Among the identified 3419 true positives, 1691 have high severity levels which make the system become unavailable, unreliable, or vulnerable to security attacks. We also observe that the ratio of issue numbers to LOC varies drastically in each repository, indicating the community's variegated development experience, coding practice, and scripting language expertise.

REFERENCES

- [1] 2012. *Critical Cloud Computing: A CIIP perspective on cloud computing services*. <https://resilience.enisa.europa.eu/cloud-security-and-resilience/publications/critical-cloud-computing>

- [2] 2013. *Human error most likely cause of datacentre downtime, finds study*. <https://www.computerweekly.com/news/2240179651/Human-error-most-likely-cause-of-datacentre-downtime-finds-study>
- [3] 2014. *DNS Outage Post Mortem*. <https://github.blog/2014-01-18-dns-outage-post-mortem/>
- [4] 2017. *Amazon And The \$150 Million Typo*. <https://www.npr.org/sections/thetwo-way/2017/03/03/518322734/amazon-and-the-150-million-typo>
- [5] 2017. *AWS's S3 outage was so bad Amazon couldn't get into its own dashboard to warn the world*. https://www.theregister.co.uk/2017/03/01/aws_s3_outage/
- [6] 2017. *Cyber Security Is A Business Risk, Not Just An IT Problem*. <https://www.forbes.com/sites/edelmantechnology/2017/10/11/cyber-security-is-a-business-risk-not-just-an-it-problem>
- [7] 2017. *The Main Difference Between Junior Programmers And Senior Programmers*. <https://www.forbes.com/sites/quora/2017/12/05/the-main-difference-between-junior-programmers-and-senior-programmers/#724b335e67f3>
- [8] 2017. *Summary of the Amazon S3 Service Disruption in the Northern Virginia (US-EAST-1) Region*. <https://aws.amazon.com/message/41926/>
- [9] 2018. *Malicious PowerShell Detection via Machine Learning*. <https://www.fireeye.com/blog/threat-research/2018/07/malicious-powershell-detection-via-machine-learning.html>
- [10] 2019. *Deep learning rises: New methods for detecting malicious PowerShell*. <https://www.microsoft.com/security/blog/2019/09/03/deep-learning-rises-new-methods-for-detecting-malicious-powershell/>
- [11] 2019. *A shell script that deleted a database, and how ShellCheck could have helped*. <https://www.vidarholen.net/contents/blog/?p=746>
- [12] 2020. *Ansible is Simple IT Automation*. <https://www.ansible.com/>
- [13] 2020. *Ansible-lint: Best practices checker for Ansible*. <https://github.com/ansible/ansible-lint/>
- [14] 2020. *Chef: Deploy new code faster and more frequently*. <https://www.chef.io/>
- [15] 2020. *Puppet-lint: Check that your Puppet manifests conform to the style guide*. <https://github.com/rodjek/puppet-lint>
- [16] 2020. *Puppet: Powerful infrastructure automation and delivery*. <https://puppet.com/>
- [17] 2020. *puppet-strings*. <https://rubygems.org/gems/puppet-strings>
- [18] 2020. *Secure infrastructure from IBM lets you make sure your data stays safe — wherever it goes*. <https://www.ibm.com/it-infrastructure/solutions/security>
- [19] 2020. *Terraform: Use Infrastructure as Code to provision and manage any cloud, infrastructure, or service*. <https://www.terraform.io/>
- [20] 2020. *Travis CI*. <https://travis-ci.org/>
- [21] 2020. *The Xenial Build Environment*. <https://docs.travis-ci.com/user/reference/xenial/>
- [22] 2020. *YAML Lint: A linter for YAML files*. <https://github.com/adrienverge/yamllint>
- [23] Ting Dai, Daniel Joseph Dean, Peipei Wang, Xiaohui Gu, and Shan Lu. 2019. Hytrace: A Hybrid Approach to Performance Bug Diagnosis in Production Cloud Infrastructures. *IEEE Trans. Parallel Distrib. Syst.* 30, 1 (2019), 107–118.
- [24] Ting Dai, Jingzhu He, Xiaohui Gu, Shan Lu, and Peipei Wang. 2018. DScope: Detecting Real-World Data Corruption Hang Bugs in Cloud Server Systems. In *Proceedings of the ACM Symposium on Cloud Computing (SoCC'18)*.
- [25] Michael Greenberg and Austin J. Blatt. 2019. Executable Formal Semantics for the POSIX Shell. *Proc. ACM Program. Lang.* 4, POPL, Article 43 (2019), 30 pages.
- [26] Oliver Hanappi, Waldemar Hummer, and Schahram Dustdar. 2016. Asserting Reliable Convergence for Configuration Management Scripts. In *Proceedings of the 2016 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA'16)*.
- [27] Jingzhu He, Ting Dai, and Xiaohui Gu. 2018. TScope: Automatic Timeout Bug Identification for Server Systems. In *2018 IEEE International Conference on Autonomic Computing, ICAC 2018, Trento, Italy, September 3-7, 2018*. 1–10.
- [28] Danny Hendler, Shay Kels, and Amir Rubin. 2018. Detecting Malicious PowerShell Commands Using Deep Neural Networks. In *Proceedings of the 2018 on Asia Conference on Computer and Communications Security (ASIACCS'18)*.
- [29] Vidar Holen. 2020. *ShellCheck – shell script analysis tool, infrastructure, or service*. <https://www.shellcheck.net/>
- [30] Waldemar Hummer, Florian Rosenberg, Fábio Oliveira, and Tamar Eilam. 2013. Testing Idempotence for Infrastructure as Code. In *ACM/IFIP International Middleware Conference (Middleware'13)*.
- [31] Guangpu Li, Shan Lu, Madanlal Musuvathi, Suman Nath, and Rohan Padhye. 2019. Efficient Scalable Thread-Safety-Violation Detection: Finding Thousands of Concurrency Bugs during Testing. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles (SOSP'19)*.
- [32] Jiaxin Li, Yuxi Chen, Haopeng Liu, Shan Lu, Yiming Zhang, Haryadi S. Gunawi, Xiaohui Gu, Xicheng Lu, and Dongsheng Li. 2018. Pcatch: Automatically Detecting Performance Cascading Bugs in Cloud Systems. In *Proceedings of the Thirteenth EuroSys Conference (EuroSys'18)*.
- [33] V. Benjamin Livshits and Monica S. Lam. 2005. Finding Security Vulnerabilities in Java Applications with Static Analysis. In *Proceedings of the 14th Conference on USENIX Security Symposium*. USENIX Association, USA.
- [34] Karl Mazurak. 2007. ABash: Finding bugs in bash scripts. In *ACM SIGPLAN Workshop on Programming Languages and Analysis for Security (PLAS'07)*.
- [35] Microsoft. 2020. *PSScriptAnalyzer*. <https://github.com/PowerShell/PSScriptAnalyzer>
- [36] Scott Moore, Christos Dimoulas, Dan King, and Stephen Chong. 2014. SHILL: A Secure Shell Scripting Language. In *11th USENIX Symposium on Operating Systems Design and Implementation (OSDI'14)*.
- [37] PyCQA. 2020. *Bandit is a tool designed to find common security issues in Python code*. <https://github.com/PyCQA/bandit>
- [38] Akond Rahman, Chris Parnin, and Laurie Williams. 2019. The seven sins: security smells in infrastructure as code scripts. In *Proceedings of the 41st International Conference on Software Engineering (ICSE'19)*.
- [39] Rian Shambaugh, Aaron Weiss, and Arjun Guha. 2016. Rehearsal: A Configuration Verification Tool for Puppet. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'16)*.
- [40] Tushar Sharma, Marios Fragkoulis, and Diomidis Spinellis. 2016. Does Your Configuration Code Smell?. In *Proceedings of the 13th International Conference on Mining Software Repositories (MSR'16)*.
- [41] Thodoris Sotiropoulos, Dimitris Mitropoulos, and Diomidis Spinellis. 2020. Practical Fault Detection in Puppet Programs. In *Proceedings of the 42nd International Conference on Software Engineering (ICSE'20)*.
- [42] E. van der Bent, J. Hage, J. Visser, and G. Gousios. 2018. How good is your puppet? An empirically defined and validated quality model for puppet. In *2018 IEEE 25th International Conference on Software Analysis, Evolution and Reengineering (SANER'18)*.
- [43] Tianyin Xu, Xinxin Jin, Peng Huang, Yuanyuan Zhou, Shan Lu, Long Jin, and Shankar Pasupathy. 2016. Early Detection of Configuration Errors to Reduce Failure Damage. In *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI'16)*.
- [44] Tianyin Xu and Yuanyuan Zhou. 2015. Systems Approaches to Tackling Configuration Errors: A Survey. *ACM Comput. Surv.* 47, 4, Article 70 (July 2015), 41 pages.

DScope: Detecting Real-World Data Corruption Hang Bugs in Cloud Server Systems

Ting Dai
North Carolina State University
Raleigh, NC, USA
tdai@ncsu.edu

Jingzhu He
North Carolina State University
Raleigh, NC, USA
jhe16@ncsu.edu

Xiaohui Gu
North Carolina State University
Raleigh, NC, USA
xgu@ncsu.edu

Shan Lu
University of Chicago
Chicago, Illinois, USA
shanlu@cs.uchicago.edu

Peipei Wang
North Carolina State University
Raleigh, NC, USA
pwang7@ncsu.edu

ABSTRACT

Cloud server systems such as Hadoop and Cassandra have enabled many real-world data-intensive applications running inside computing clouds. However, those systems present many data-corruption and performance problems which are notoriously difficult to debug due to the lack of diagnosis information. In this paper, we present DScope, a tool that statically detects data-corruption related software hang bugs in cloud server systems. DScope statically analyzes I/O operations and loops in a software package, and identifies loops whose exit conditions can be affected by I/O operations through returned data, returned error code, or I/O exception handling. After identifying those loops which are prone to hang problems under data corruption, DScope conducts loop bound and loop stride analysis to prune out false positives. We have implemented DScope and evaluated it using 9 common cloud server systems. Our results show that DScope can detect 42 real software hang bugs including 29 newly discovered software hang bugs. In contrast, existing bug detection tools miss detecting most of those bugs.

CCS CONCEPTS

• **Computer systems organization** → **Cloud computing; Reliability**; • **Software and its engineering** → **Automated static analysis; Software performance**;

KEYWORDS

static analysis, data corruption, performance bug detection

ACM Reference Format:

Ting Dai, Jingzhu He, Xiaohui Gu, Shan Lu, and Peipei Wang. 2018. DScope: Detecting Real-World Data Corruption Hang Bugs in Cloud Server Systems. In *Proceedings of SoCC '18: ACM Symposium on Cloud Computing, Carlsbad, CA, USA, October 11–13, 2018 (SoCC'18)*, 13 pages.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SoCC'18, October 11–13, 2018, Carlsbad, CA, USA

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-6011-1/18/10...\$15.00

<https://doi.org/10.1145/3267809.3267844>

```
// LeaseManager.java #HDFS-4882(v0.23.0)
393 private synchronized void checkLeases() {
    ...
395 for( ; sortedLeases.size() > 0; ) {
    ...
411 try { //p is a file's lease path
412     if(fsnamesystem.internalReleaseLease(
413         oldest, p, ...)) {
        ... //remove p from sortedLeases
416     }
    ...
420 } catch (IOException e) {
        ... //remove p from sortedLeases
423     }
    ...
429 }
430 }
```

Figure 1: A real-world data corruption hang bug from HDFS. A corrupted file f associated with the lease path p makes the `internalReleaseLease` function fail for recovering the lease for f . When this failure happens, p is not removed from `sortedLeases` (skip updating loop index), `LeaseManager` keeps recovering lease for the file f endlessly.

<https://doi.org/10.1145/3267809.3267844>

1 INTRODUCTION

Cloud server systems such as Hadoop and Cassandra [3, 4] have enabled many real-world data-intensive applications ranging from security attack detection to business intelligence. However, due to their inherent complexity, those cloud server systems present many performance challenges. Particularly, previous studies [23, 25] have shown that many tricky performance bugs in cloud server systems are caused by unexpected data corruptions which are more likely to be overlooked by the developer. For example, in May 2017, a data corruption bug triggered in a data center failover operation brought down the British Airway service for hours [2].

Performance bugs¹ are notoriously difficult to debug because they typically produce little useful debugging information. The problem exacerbates in cloud server systems since the developer typically does not have the access to the original input data that triggered the performance bug or the large scale infrastructure

¹We use performance bugs to broadly refer to all non-functional bugs, which could cause slowdown or system unavailability.

to replay the failed production run. Although previous work has extensively studied data corruptions (e.g., [11, 15, 23, 28, 44]) and performance bugs (e.g., [17, 26, 31, 40]), little research has been done to study the intersection between the two, that is, the performance problems caused by data corruptions. Particularly, our work focuses on detecting software hang bugs that are triggered by data corruptions in cloud server systems. Software hang bugs make the system become unavailable to either part of or all of the users, which is one of the most severe performance problems production systems try to avoid [17–19, 29].

1.1 A Motivating Example

To better understand how real-world data corruption hang bugs happen, we use a known HDFS-4882² bug as one example shown by Figure 1. This hang bug happens when the improper handling of a corrupted file *f* causes the loop to skip updating its loop index. In HDFS, when a client's leases get expired, the lease recovery is triggered by the LeaseManager on the NameNode. The LeaseManager sends a lease recovery request for each lease in the sortedLeases set to the FSNamesystem via a RPC call (line #412-413). If a lease is successfully recovered (e.g., released or renewed) (line #412-416), or an IOException happens during the lease recovery (line #420-423), the lease path *p* is removed from the sortedLeases. The LeaseManager keeps recovering and removing the leases until the sortedLeases set is empty (line #395). However, the FSNamesystem only considers the case where the last block is corrupted if data corruption happens in a file. Thus, when the second-to-last block of a file *f* (i.e., an INode) is corrupted but the processing state of the last block of *f* is complete, the FSNamesystem improperly handles this case and returns false by mistake (line #412). This bug occurs when the HDFS client finished writing the second-to-last block, starts to write the last block and part of the DataNodes experience a shut-down failure. To resume the process, the NameNode marks the second-to-last block as committed, unblocks the HDFS client from writing the last block, and marks the last block as complete [1]. As a result, the lease path *p* is not removed from the sortedLeases, and the LeaseManager keeps invoking the lease recovery for the same lease endlessly.

1.2 Our Contribution

This paper presents DScope, an automated corruption-hang bug detection tool for server systems commonly used in computing clouds. DScope is a static analysis tool — it can detect data-corruption related hang bugs without running the target system and it requires no system-specific knowledge. To achieve both high coverage and low false positives, DScope first uses static control flow and data flow analysis to identify loops whose exit conditions may be affected by external data (i.e., I/O operations), and then conducts loop bound and loop stride analysis to filter out loops which are guaranteed not to have hang problems. To support such analysis, DScope models Java data-related APIs which are commonly used in cloud server systems.

This paper makes the following contributions:

- We present a hang-bug detection scheme that identifies potential infinite loops caused by data corruptions in cloud server systems.
- We describe a false-positive pruning technique that identifies always-exit loops through loop stride and bound analysis. Different from generic loop analysis, our analysis focuses on a wide variety of Java I/O APIs widely used in cloud systems, and helps greatly improve the accuracy of our data-corruption hang-bug detection.
- We categorize real-world data-corruption hang bugs into four common types based on DScope detection results. This categorization will help future work on avoiding, detecting, and preventing data-corruption hang bugs.

We have implemented DScope and evaluated it using 9 commonly used cloud server systems (e.g., Cassandra, HDFS, Mapreduce, Hive, etc). DScope reports 42 true data corruption hang bugs, with 29 of them are newly discovered bugs. We also applied two state of the art static bug detectors, Findbugs [7] and Infer [6], to the same set of systems. They detect very few corruption hang bugs (2 for Findbugs and 1 for Infer), indicating the need for a dedicated corruption hang bug detector like DScope.

The rest of the paper is organized as follows. §2 describes the design of the DScope system. §3 presents the types of the data corruption hang bugs. §4 shows the experimental evaluation. §5 discusses the future work. §6 compares our work with related work. Finally, the paper concludes in §7.

2 SYSTEM DESIGN

This section first provides an overview of DScope (§2.1). It then presents the detailed designs of how to discover corruption-hang bug candidates (§2.2) and how to prune false positives (§2.3).

2.1 Approach Overview

DScope focuses on detecting software hang bugs caused by potential data corruptions in cloud server systems. Our bug detection scheme consists of two major steps: 1) discovering all candidate data corruption hang bugs that aims at maximizing detection coverage; and 2) filtering out false positive detections by identifying code patterns that assure the program will not hang under any circumstance.

Since many software hang problems are caused by infinite loop bugs [19, 29], our work focuses on detecting possible infinite loops caused by data corruptions in cloud server systems written in Java. To detect those loop bugs, DScope leverages the Soot compiler framework [8] to compile application bytecode into intermediate representation (IR) code (i.e., Soot Jimple) and perform static analysis over the IR code in three steps: 1) loop path extraction, 2) I/O dependent loop identification, and 3) loop stride and bound analysis.

Specifically, DScope first extracts different execution paths which start from the loop header and end at the loop header by traversing the control flow graph (CFG) of all loops. Next, DScope derives the exit conditions of each loop path and checks whether those exit conditions depend on any I/O operations. The rationale is that if the loop exit condition depends on an I/O operation, a data corruption (e.g., hardware failure [11, 12, 27, 30, 36, 39], software

²We use "system name-bug #" to denote different bugs.

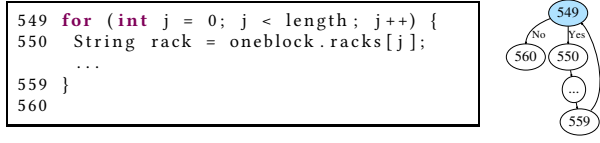


Figure 2: The example of a simple loop with the source code block and the corresponding CFG in the CombineFileInputFormat class in Hadoop v0.23.0.

fault [44]) can cause the loop exit condition to be never met and thus an infinite loop software hang bug.

After discovering candidate data corruption hang bugs, DScope performs false positive pattern filtering to improve the bug detection precision. The false positive filtering is based on the loop index, loop stride (i.e., the delta value applied to the loop index in each iteration) and loop bound analysis on every loop path. For example, if the loop stride is always positive when the loop bound is an upper bound (or if the loop stride is always negative when the loop bound is a lower bound), and if the loop bound is unchanged during each loop iteration and the loop exit conditions involve bound checking, we say that the detected hang bug is a false positive because the loop will always exit without causing a software hang.

2.2 Identify Bug Candidates

DScope discovers candidate corruption-hang bugs by 1) traversing the CFGs of all loops in application functions to derive their loop paths and 2) checking whether the exit conditions of those loop paths are I/O dependent.

Loop path extraction. For a simple loop, the execution path within one loop iteration, called a *loop path*, consists of all the statements that start from the loop header and end at the loop header. We can extract the loop path easily by traversing the CFG of the loop. For example, Figure 2 shows the source code and its CFG of a simple loop³. DScope generates a loop path {549, 550, ..., 559, 549}.

For a nested loop, the loop path consists of concatenations of the execution paths of both inner loops and outer loops. DScope extracts all loop paths using three steps. First, for the execution path, denoted as P_{outer} whose tail is the outer loop header, we add the path into a path set called S_{path} . Second, for the execution path P_{outer} whose tail is a loop body statement, we infer this statement must be the header of an inner loop and extracts the inner loop execution path denoted as P_{inner} from P_{outer} . Third, for each loop path in S_{path} , DScope first clones it and replaces any statement s_i with P_{inner} if s_i is an inner loop header and the current loop path does not contain P_{inner} . This new concatenated loop path is then added to the path set S_{path} . DScope repeats the third step until there is no more new loop path generated. Figure 3 shows an example of nested loops. First, DScope extracts one loop path {544, ..., 549, 550, ..., 571, 544}. Second, DScope extracts {549, 550, ..., 559, 549} as an inner loop path. Third, DScope replaces 549 with {549, 550, ..., 559, 549} on {544, ..., 549, 560, ..., 571, 544}, to




Figure 3: The example of nested loops with the source code block and the corresponding CFG in the CombineFileInputFormat class in Hadoop v0.23.0.

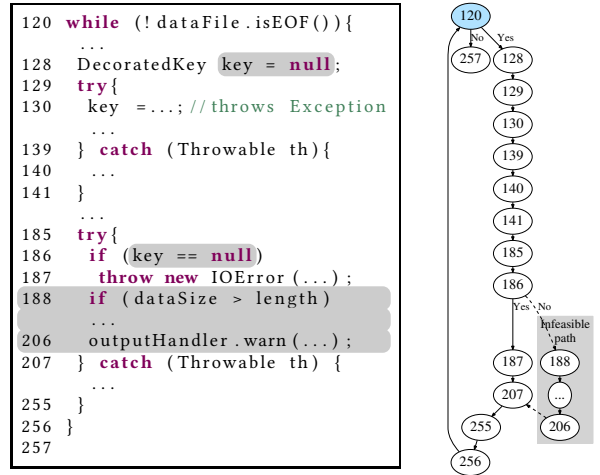


Figure 4: The example of a loop containing exception handling constructs with the source code block and the corresponding CFG in the Scrubber class in Cassandra v2.0.8.

create a new loop path {544, ..., 549, 550, ..., 559, 549, 560, ..., 571, 544}.

The third group of complicated loops involve exceptions. For those loops, some sub-paths become infeasible due to the exception handling, which should not be considered in our loop exit condition checking. For example, Figure 4 shows a `while` loop containing exception handling. The assignment statement at line #130 can throw an exception when the operation on the right hand side processes a null argument. As a result, the variable `key` is not updated and remains to be the default value which is `null`. So when the exception is triggered, the `if` statement (line #186) always returns true. Thus, all the statements in the `else` branch (line #188-206) are unreachable and any path consists of those statements are infeasible paths. In this example, DScope only generates the loop path as {120, 128, 129, 130, 139, 140, 141, 185, 186, 187, 207, 255, 256, 120}.

It can be computationally expensive to traverse the CFG of the loops containing multiple exceptions because every statement in the `try` block has two branches (i.e., triggering or not triggering the exception) resulting in a large CFG. DScope addresses the problem by grouping all the statements based on the data they process. Specifically, DScope identifies all the statements which involve function invocations in the `try` blocks and groups them based on the *arguments* of those function invocations. Since DScope aims

³DScope analyzes IR code directly to extract the execution path of different loops. For easy understanding, we illustrate the execution paths using source code in the rest of the paper.

```
// Soot IR
198 $i1 = r0.<InputStream: read()>(r2) // $i1 is an I/O
    // related variable
199 if $i1 == -1 goto line #203 // $i1 == -1 is the
    // exit condition
202 goto line #198
```

Figure 5: The example of the loop's exit condition directly depends on I/O operations. It is in the IOUtils class of Compress v1.0.

```
// Soot IR
3 if l8 >= 10 goto line #12 // l8 >= 10 is the
    // exit condition
5 $l2 = 10 - l8
6 $l4 = $r2.<InputStream: skip>($l2) // $l4 is an I/O
    // related variable
7 $b5 = $l4 cmp 0L // related variable
8 if $b5 == 0 goto line #12 // $b5 == 0 is the
    // exit condition
9 $l7 = $l8 + $l4 // exit condition
10 l8 = $l7
11 goto line #3
```

Figure 6: The example of the loop's exit condition indirectly depends on I/O operations. It is in the NonSyncDataInputBuffer class of Hive v2.3.2.

at detecting data corruption hang bugs, we can assume all the statements in the same group throw exceptions when their *arguments* get corrupted. Suppose there are m statements in the try blocks. DScope divides all m statements into n groups and runs the loop path discovery algorithm 2^n times. Thus, DScope can reduce the loop path search space from 2^m to 2^n ($n \ll m$), which reduces DScope's analysis time and resource requirements (e.g., avoiding analysis failures caused by OutOfMemoryException).

I/O dependent loop identification. To discover candidate data corruption hang bugs, DScope identifies those loops whose exit conditions depend on I/O operations, which are called *I/O dependent loops*. After extracting a loop path, DScope identifies all the loop exit instructions and derives the loop path's exit conditions by performing a union over the exit conditions of all the branch statements. We consider a loop path is I/O dependent if *any* of its exit conditions depend on I/O operations. The rationale is that a data corruption can cause the corresponding I/O operations to return unexpected values or throw exceptions, making the loop never exit and thus software hang. DScope considers the operations performed on *I/O classes* via virtual invocations or on the *I/O variables* via instance invocations as *I/O operations*. The I/O classes include all classes and interfaces in `java.io` and `java.nio` packages and their subclasses and implementation classes. The instances of the I/O classes are called *I/O variables*. Additionally, DScope allows users to easily add application I/O classes in the configuration files to maximize detection coverage by identifying more application I/O dependent loops.

DScope checks whether the loop exit conditions *directly* depend on I/O operations by identifying the appearance of I/O classes in the exit checking statements. Figure 5 shows an example where the loop exit condition directly depends on the I/O operations. In this example, the variable `$i1` in the exit condition checking statement (line #199) is directly derived from a Java I/O class called `InputStream`.

DScope checks whether the loop exit conditions *indirectly* depend on I/O operations by performing data dependency analysis

```
// Soot IR
10 $r13 = new java.util.HashMap

// Java source code
269 HashMap<OneBlockInfo, String[]> blockToNodes =
270 new HashMap<OneBlockInfo, String[]>();

// Java bytecode
Constant pool:
#219 = blockToNodes
#232 = Ljava/util/HashMap<
    Lorg/apache/hadoop/mapreduce/lib/
    input/CombineFileInputFormat$OneBlockInfo;
    [Ljava/lang/String; >;
LocalVariableTable:
name index signature index
#219 .....>#232
```

Figure 7: The `java.util.HashMap<K,V>` example in the `CombineFileInputFormat` class in Hadoop v0.23.0.

on all the statements of the corresponding application function. Specifically, DScope first identifies all the I/O related variables which are assigned with the return values of the I/O operations. Second, for each assignment statement of the application function that involves any I/O related variables on its right-hand-side, DScope iteratively labels the variable on the left-hand-side of the assignment statement as I/O related variables as well. After identifying all the I/O related variables, DScope checks whether the loop exit conditions are I/O dependent by identifying the appearance of I/O related variables in the exit checking statements. Figure 6 shows an example of indirectly I/O dependent loop exit condition. In this example, the loop exit checking involves `$l8` (line #3) and `b5` (line #8) whose value is derived from `l4` which is derived from a Java I/O operation `InputStream.skip()`.

To further check whether the loop exit conditions depend on I/O operations conducted on *complex I/O related variables* (i.e., variables with composite types), DScope performs an integrated analysis by linking variable information from IR code, Java source code, and Java bytecode⁴. DScope considers a variable with composite type as I/O related if any of the variable's elements is I/O related. Note that, by checking only the IR code, DScope might miss identifying some complex variables as I/O related. For example, in Figure 7, by checking only the IR code, DScope cannot identify the variable `$r13` as an I/O related variable, thus all the operations conducted on `$r13` will not be considered as I/O operations. This is because `$r13` is of type `HashMap` and `HashMap` is not an I/O class. To identify complex I/O related variables, DScope needs to retrieve the full type information (i.e., class path) in the Java bytecode for a target variable in the IR code. However, there is no direct mapping from IR code to Java bytecode. So, DScope has to leverage the source code to establish the mapping from. Specifically, DScope first retrieves the source code line number from Soot via `getLineNumber()` API for each variable `valIR` in the IR code. DScope then analyzes the corresponding source code and extracts `valIR`'s name in the source code, denoted as `valsrc`. In Figure 7, DScope extracts that the variable `$r13` is defined at line #269 in the source code with name `blockToNodes`. Next, DScope

⁴DScope mainly works on Soot IR code, except the integrated analysis in the I/O dependent loop identification module.

Table 1: The 60 commonly used Java classes and interfaces which contain APIs related to the loop index, stride and bound.

Prefix	Class	# of classes or interfaces
java.io	DataInput family	2
	File	1
	InputStream family	12
	Reader family	10
java.nio	Buffer family	8
	channels.Channel family	20
	Iterator, Enumeration	2
java.util	List, Queue, Set, Stack	4
	StringTokenizer	1

leverages Coffi [43], a Java bytecode parser, to extract the full type information for the target variable val_{src} . Specifically, the constant pool provides index lookup for each variable and LocalVariableTypeTable provides the mapping from the variable's index to the index of its signature which contains the full type information. In Figure 7, the constant pool indicates that the index of the variable `blockToNodes` is #219 and the LocalVariableTypeTable lookup tells its corresponding signature index is #232. DScope checks the constant pool using the index number #232 to derive the full type information of `blockToNodes`. Since `blockToNodes` consists of `CombineFileInputFormat$OneBlockInfo` class which is I/O related (i.e., an application I/O class), DScope infers `blockToNodes` is also I/O related. Thus the operations conducted on `blockToNodes` are I/O operations and the loops whose exit conditions depend on those I/O operations are I/O dependent loops.

2.3 Prune False Positives

Since our goal of candidate bug detection is to maximize coverage, false positives can be inevitably included in the candidate list. To improve DScope's bug detection precision, we further develop false positive pattern filtering schemes by identifying those loops which will always exit without causing any software hang. Our false positive filtering is achieved by analyzing the loop stride and loop bounds. DScope prunes false positive candidates by checking whether 1) the loop stride is *always* positive when the loop has an upper bound or the loop stride is always negative when the loop has a lower bound; 2) the loop bound value is unchanged in *every* loop iteration; and 3) the loop exit conditions contain bound checking. Intuitively, any loops satisfying all those conditions will always exit without causing software hang, which should be pruned from DScope's detection list.

DScope's loop stride and bound analysis schemes consider two cases: a) the loop index, stride, and bounds are denoted by numeric primitives (e.g., integer); and b) the loop index, stride, and bounds are denoted by APIs in 60 commonly used Java classes and interfaces, shown in Table 1. Note that those Java classes and interfaces are not necessarily the *I/O classes* but appear frequently in the *I/O dependent loops*. Moreover, they do not include all the Java classes and interfaces which contain the loop related APIs. We plan to further extend our analysis to cover other Java classes

```
// Soot IR
127 $b6 = i1 cmp i0 // i1 is the loop index
128 if $b6 >= 0 goto line #139 // i0 is the upper bound
129 ...
130 i7 = i1 + 1
131 i12 = i7 + 1
132 i17 = i12 + 1
133 i22 = i17 + 1
134 i27 = i22 + 1
135 i32 = i27 + 1
136 i37 = i32 + 1
137 i1 = i37 + 1 // all the 1's are the strides
138 goto line #127
```

Figure 8: The example of multiple strides. It is in the `OffHeapBitSet` class of `Cassandra v2.0.8`.

```
// Soot IR
530 $i5=r1.<ByteBuffer:limit> // $i5 is the upper bound
531 if i1 >= $i5 goto line #536 // i1 is the loop index
532 $i9 = <STEP_LENGTH>; // $i9 is the stride
533 i1 = i1 + $i9
...
535 goto line #531
```

Figure 9: The example of the stride is assigned outside of the function `total` where the loop resides. It is in the `CounterContext` class of `Cassandra v2.0.8`. The stride `STEP_LENGTH` is a static variable, which is assigned with 34 in the class initializer.

in our future work, which can further improve our false positive filtering efficacy.

When the loop index, stride, and bounds are denoted by numeric primitives, DScope first extracts the loop index variable from the loop exit conditions. The *loop index* is a variable that appears in the loop exit conditions and is updated by another variable in an assignment statement via arithmetic operations (e.g., addition and subtraction). After identifying the loop index, the variable to which it compares in the exit conditions is the *loop bound* and the variable which is added to or subtracted from the loop index is the *loop stride*. DScope further checks whether the numeric stride is positive when the loop has an upper bound or negative when the loop has a lower bound. For example, in the expression "*index = index op stride*", if the "*op*" is an addition operation, the loop stride is positive. In the expression "*index symbol bound*", if the "*symbol*" is \leq or $<$, the loop has an upper bound. Finally, DScope examines all the loop paths and checks whether the bound is unchanged within every loop path. Based on all the extracted information, DScope can make decision whether a discovered loop bug is a false positive.

When the loop index, stride, and bounds are denoted by *multiple* numeric primitives or the numeric primitives *outside* the current application function where the loop resides, DScope performs intra-procedure data flow analysis on all the statements of the corresponding application class to achieve accurate false positive filtering. For example, Figure 8 shows a loop with multiple strides in the `OffHeapBitSet` class in `Cassandra`. The variable `i1` is the loop index while the variable `i0` is the upper bound. The loop index `i1` is updated multiple times from line #130 to line #137. DScope recursively applies all the assignments from line #130 to line #137 to get `i1 = i1 + 8`, and extracts the aggregated stride (i.e., 8). Figure 9 shows an example where the stride variable `$i9` is updated by `STEP_LENGTH` in the `CounterContext` class in `Cassandra`.

Table 2: The APIs that are related to loop stride and bound update in 60 commonly used Java classes and interfaces. “*”: a set of APIs perform similar operations; and “-”: does not contain the corresponding type APIs.

Class	The type and name of APIs				
	Forward index	Reverse index	Reset index	Check bounds	Update bounds
File	create*	get*	new	is*, can* exists, get*	-
InputStream & Reader family	read*	reset	new	read*	new
DataInput family	read*	-	new	read*	new
Buffer family	position get* put*	position reset clear	duplicate allocate new	has* remaining	flip limit clear new
Channel family	read write	-	-	read write	-
List & Set	-	remove	new	is*	add clear new
Queue	-	poll remove	-	poll remove	add offer new
Stack	-	pop	-	empty pop	push new
Iterator & Enumeration	next	-	iterator elements new	has*	new
StringTokenizer	next	-	new	has*	new

The loop resides in the function `total()` while `STEP_LENGTH` is a static variable defined in the class initializer. DScope performs data flow analysis to extract the value of `STEP_LENGTH` from the class initializer and then checks whether the stride is positive because the loop index has an upper bound.

We now describe how to perform false positive filtering when the loop index, stride and bounds are denoted by the APIs in 60 commonly used Java classes and interfaces, listed in Table 1.

We classify those APIs into five categories, shown by Table 2: 1) the APIs which move the index forward when the Java class/interface has an upper bound; 2) the APIs which move the index backward when the Java class/interface has a lower bound; 3) the APIs which reset the index; 4) the APIs which check bounds; and 5) the APIs which update bounds. The APIs’ names ending with “*” denote those APIs which perform similar operations and share the same prefix in their names. For example, in the `InputStream` family, there are `read()` and `readLine()` functions which both perform read operations on the corresponding `InputStream`.

DScope first extracts all the invoked APIs for each of the 60 Java classes and interfaces in the loop paths, and then prunes false positive candidates by checking whether 1) the “forward index” APIs or the “reverse index” APIs are invoked; 2) the “reset index” APIs and “update bounds” APIs are not invoked; and 3) the “check bounds” APIs are invoked in the exit conditions. Note that, DScope

```
//DFSOutputStream.java #HDFS-5438(v0.23.0)
1665 private void completeFile(ExtendedBlock last) ... {
    ...
1667 boolean fileComplete = false;
1668 while (!fileComplete) {
1669     fileComplete = dfsClient namenode.complete(src,
        dfsClient.clientName, last);
    ...
1689 }}
```

Figure 10: The code snippet of the HDFS-5438 Bug. When the `ExtendedBlock last` is corrupted, the `fileComplete` variable is never set to be true, causing an infinite loop in `DFSOutputStream`.

cannot prune the case where both the “forward index” APIs and the “reverse index” APIs are invoked in the loop paths because the loop stride cannot be guaranteed to be *always* positive or negative.

The APIs in the five categories do not necessarily change the loop index or bounds. Those APIs’ arguments should also be considered when DScope performs the false positive filtering. For example, `ByteBuffer` contains overloading methods which have an attribute called *relative* or *absolute*. The `ByteBuffer.get()` is a relative method while the `ByteBuffer.get(int)` is an absolute method. Invoking a relative method can change the loop index (i.e., `ByteBuffer.position`) while invoking absolute methods cannot. Another example is the `InputStream` class. Invoking the `InputStream.read(byte[], int, int)` with a zero size byte array or with 0 as the third parameter cannot change the loop index.

To achieve accurate pruning, DScope first annotates all the commonly used Java APIs with the attribute *change-positive*, *change-negative* or *change-possible*. The positive APIs can change the loop index (or bounds). The negative APIs cannot change either one. The possible APIs can possibly change the loop index (or bounds). For positive APIs, DScope’s pruning steps are the same. For negative APIs in the type of “forward index”, “reverse index”, “reset index” or “update bounds”, DScope ignores them when performing the pruning. For possible APIs, DScope performs intra-procedural data flow analysis on their parameters to decide whether these APIs change loop index/bounds or not.

DScope’s false positive filtering only considers the commonly used Java APIs. If the loop index, stride or bounds are *only* related to specific application functions, which means the loop paths do not invoke any Java APIs in Table 2, DScope skips analyzing the loop and simply considers it as a false positive — this design decision may introduce false negatives, but greatly help the efficiency and accuracy of DScope.

One false negative example is the HDFS-5438 bug, shown by Figure 10. This hang bug is caused by a corrupted block, i.e. `last`. `DFSOutputStream` keeps polling `NameNode` to check the completeness of the committing block operation (line #1669). When the `last` block is corrupted, `NameNode` fails to commit it to the disk but returns `false` instead. This results in an infinite loop (line #1668-1689) causing a software hang in `DFSOutputStream`. DScope prunes this case because the loop paths do not invoke any Java APIs in Table 2. In fact, the loop path only invokes a specific application function, i.e., `complete()`. DScope should be able to detect this bug after adding inter-procedural analysis, which is however beyond the scope of this work.


```
// IOUtils.java #Hadoop-8614(v0.23.0)
183 public static void skipFully(InputStream in, long len
    ) throws IOException {
184     while (len > 0) {
185         long ret = in.skip(len); /* in is corrupted */
186         if (ret < 0) { /* ret = 0 */
187             throw new IOException (...);
188         }
189         len -= ret;
    }
}
```

Figure 11: The example when error code returned by I/O operations directly impacts the loop stride. Data corruption causes the I/O function, `InputStream.skip` returns 0, and 0 is used as the stride.

3 DATA CORRUPTION HANG BUG TYPES

This section summarizes common types of corruption-hang bugs based on the detection results of DScope (the details of all the bugs detected by DScope will be presented in §4). Although DScope design was **not** affected by these types, this categorization can help future work on avoiding, detecting, and fixing corruption-hang bugs, and help developers better understand the impact of data corruption and corruption-hang bugs.

Our categorization is along two dimensions:

- What is the cause — is it specific error code returned by data operations (Type 1), or specific corrupted data content (Type 2), or specific exception thrown by data operations (Type 3, Type 4)?
- How did the cause lead to an infinite loop — is it through a direct data assignment (Type 1) or control-flow change (Type 3), or indirect data and control flow (Type 2, Type 4)?

Type 1: Error codes returned by I/O operations directly affect loop strides. For this type, the loop stride is directly assigned with a return value of an I/O operation. An infinite loop occurs when an unexpected error code is returned due to underlying data corruption. For example, as shown by Figure 11, when the log file (`InputStream in` at line #183) is corrupted due to bad encoding (Yarn-2724) or corruption propagation (Yarn-7179), the `InputStream in` becomes null. The `skip()` function returns 0 instead of the EOF indicator -1 (line #185). The return value `ret` is then used as the stride at line #189, which makes `len` never get updated but always stay larger than the lower bound (`len > 0`). As a result, the `skipFully()` function causes the system to hang by spinning in the loop forever. Variations of this hang bug type include the cases where the stride is always negative when the loop exit condition contains an upper bound or the stride is always positive when the loop exit condition contains a lower bound.

Type 2: Corrupted data content indirectly affects loop strides. This type of bugs occur when a specific piece of data is corrupted to certain unexpected values. Those values will then affect loop strides through data and/or control flow propagation and lead to infinite loops. For example, as shown by Figure 12, when a configuration file (`conf` at line #190) is corrupted, the variable `BUFFER_SIZE` read from `conf` becomes 0. Calling `read()` function on a zero-size byte array at line #87 causes the loop stride to be zero and zero is then returned, which makes the loop's exit condition (`size < 0`) never be satisfied.

Figure 13 shows an example when the loop stride is indirectly impacted by the corrupted data content which involves *multiple*

```
// BenchmarkThroughput.java #HDFS-13514(v2.5.0)
172 public int run(...) throws IOException {
190     Configuration conf = getConf(); /* conf is corrupted */
    ...
194     BUFFER_SIZE = conf.getInt(...); /* BUFFER_SIZE = 0 */
    ...
229 }
```

```
78 private void readLocalFile(Path path, ... throws
    IOException {
    ...
83     InputStream in = new FileInputStream(...);
84     byte[] data = new byte[BUFFER_SIZE];
85     long size = 0;
86     while (size >= 0) { /* size = 0 */
87         size = in.read(data);
    }
}
```

Figure 12: The example when corrupted data content indirectly impacts the loop stride. The corrupted configuration file causes “`BUFFER_SIZE = 0`”, which in turn makes the `InputStream in` perform read operation on a zero-size byte array and return 0. The loop's exit condition become infeasible because “`size < 0`” is never satisfied.

```
// CombineFileInputFormat.java #Mapreduce-2185(v0.23)
477 private static class OneFileInfo {
    ...
544     for (OneBlockInfo oneblock : blocks) {
545         blockToNodes.put(oneblock, oneblock.hosts);
        ... /* corrupted block's racks.length is 0 */
549         for (int j = 0; j < oneblock.racks.length; j++) {
550             String rack = oneblock.racks[j];
            ...
554             rackToBlocks.put(rack, blklist);
            ...
        }
    }
}
```

```
255 private void getMoreSplits(...) throws ... {
    ...
348     while (blockToNodes.size() > 0) {
        ...
359         for (Iterator<...> iter = rackToBlocks.
360             entrySet().iterator(); iter.hasNext(); ) {
361             Map.Entry<...> one = iter.next();
            ...
363             List<OneBlockInfo> blocks = one.getValue();
            ...
369             for (OneBlockInfo oneblock : blocks) {
370                 if (blockToNodes.containsKey(oneblock)) {
371                     blockToNodes.remove(oneblock);
                    ...
                }
            }
        }
    }
}
```

Figure 13: The example when corrupted data content indirectly impacts the loop stride. Data corruption causes `blockToNodes` and `rackToBlocks` to be different on the dimension of the blocks' number. This difference makes the corrupted block never been removed from the `blockToNodes` (i.e., zero-stride), causing the loop's exit condition to be infeasible. This is because “`blockToNodes.size() <= 0`” is never satisfied.

I/O related variables. The `blockToNodes` and `rackToBlock` are two maps which store different metadata information about every data block. If everything works correctly, these two maps should contain information about exactly the same set of blocks (i.e., every record in the blocks on line #544). However, if a block (`oneblock` at line #549) is corrupted and its `racks.length` becomes 0, this block will still be inserted into `blockToNodes` at line #545, but not

```

//TestProcsBasedProcessTree.java      #Yarn-6991(v0.23.0)
//Thread #1
62 private class RogueTaskThread extends Thread {
63     public void run() {
64         try {
65             ...
66             args.add(" echo $$ > " + pidFile + ".");
67             shexec = new ShellCommandExecutor(args...);
68             shexec.execute();
69         } catch (IOException ioe) {
70             LOG.info("Error executing cmd");
71         }
72     }
73 }
74 }
75 }
76 }
77 }
78 }
79 }
80 }
81 }
82 }
83 }
84 }
85 }
86 }
87 }
88 }
89 }
90 }
91 }
92 }
93 }
94 }
95 }
96 }
97 }
98 }
99 }
100 }
101 }
102 }
103 }
104 }
105 }
106 }
107 }
108 }
109 }
110 }
111 }
112 }
113 }
114 }
115 }
116 }
117 }
118 }
119 }
120 }
121 }
122 }
123 }
124 }
125 }
126 }
127 }
128 }
129 }
130 }
131 }
132 }
133 }
134 }
135 }
136 }
137 }
138 }
139 }
140 }
141 }
142 }
143 }
144 }
145 }
146 }
147 }
148 }
149 }
150 }
151 }
152 }
153 }
154 }
155 }
156 }
157 }
158 }
159 }
160 }
161 }
162 }
163 }
164 }
165 }
166 }
167 }
168 }
169 }
170 }
171 }
172 }
173 }
174 }
175 }
176 }
177 }
178 }
179 }
180 }
181 }
182 }
183 }
184 }
185 }
186 }
187 }
188 }
189 }
190 }
191 }
192 }
193 }
194 }
195 }
196 }
197 }
198 }
199 }
200 }
201 }
202 }
203 }
204 }
205 }
206 }
207 }
208 }
209 }
210 }
211 }
212 }
213 }
214 }
215 }
216 }
217 }
218 }
219 }
220 }
221 }
222 }
223 }
224 }
225 }
226 }
227 }
228 }
229 }
230 }
231 }
232 }
233 }
234 }
235 }
236 }
237 }
238 }
239 }
240 }
241 }
242 }
243 }
244 }
245 }
246 }
247 }
248 }
249 }
250 }
251 }
252 }
253 }
254 }
255 }
256 }
257 }
258 }
259 }
260 }
261 }
262 }
263 }
264 }
265 }
266 }
267 }
268 }
269 }
270 }
271 }
272 }
273 }
274 }
275 }
276 }
277 }
278 }
279 }
280 }
281 }
282 }
283 }
284 }
285 }
286 }
287 }
288 }
289 }
290 }
291 }
292 }
293 }
294 }
295 }
296 }
297 }
298 }
299 }
300 }
301 }
302 }
303 }
304 }
305 }
306 }
307 }
308 }
309 }
310 }
311 }
312 }
313 }
314 }
315 }
316 }
317 }
318 }
319 }
320 }
321 }
322 }
323 }
324 }
325 }
326 }
327 }
328 }
329 }
330 }
331 }
332 }
333 }
334 }
335 }
336 }
337 }
338 }
339 }
340 }
341 }
342 }
343 }
344 }
345 }
346 }
347 }
348 }
349 }
350 }
351 }
352 }
353 }
354 }
355 }
356 }
357 }
358 }
359 }
360 }
361 }
362 }
363 }
364 }
365 }
366 }
367 }
368 }
369 }
370 }
371 }
372 }
373 }
374 }
375 }
376 }
377 }
378 }
379 }
380 }
381 }
382 }
383 }
384 }
385 }
386 }
387 }
388 }
389 }
390 }
391 }
392 }
393 }
394 }
395 }
396 }
397 }
398 }
399 }
400 }
401 }
402 }
403 }
404 }
405 }
406 }
407 }
408 }
409 }
410 }
411 }
412 }
413 }
414 }
415 }
416 }
417 }
418 }
419 }
420 }
421 }
422 }
423 }
424 }
425 }
426 }
427 }
428 }
429 }
430 }
431 }
432 }
433 }
434 }
435 }
436 }
437 }
438 }
439 }
440 }
441 }
442 }
443 }
444 }
445 }
446 }
447 }
448 }
449 }
450 }
451 }
452 }
453 }
454 }
455 }
456 }
457 }
458 }
459 }
460 }
461 }
462 }
463 }
464 }
465 }
466 }
467 }
468 }
469 }
470 }
471 }
472 }
473 }
474 }
475 }
476 }
477 }
478 }
479 }
480 }
481 }
482 }
483 }
484 }
485 }
486 }
487 }
488 }
489 }
490 }
491 }
492 }
493 }
494 }
495 }
496 }
497 }
498 }
499 }
500 }
501 }
502 }
503 }
504 }
505 }
506 }
507 }
508 }
509 }
510 }
511 }
512 }
513 }
514 }
515 }
516 }
517 }
518 }
519 }
520 }
521 }
522 }
523 }
524 }
525 }
526 }
527 }
528 }
529 }
530 }
531 }
532 }
533 }
534 }
535 }
536 }
537 }
538 }
539 }
540 }
541 }
542 }
543 }
544 }
545 }
546 }
547 }
548 }
549 }
550 }
551 }
552 }
553 }
554 }
555 }
556 }
557 }
558 }
559 }
560 }
561 }
562 }
563 }
564 }
565 }
566 }
567 }
568 }
569 }
570 }
571 }
572 }
573 }
574 }
575 }
576 }
577 }
578 }
579 }
580 }
581 }
582 }
583 }
584 }
585 }
586 }
587 }
588 }
589 }
590 }
591 }
592 }
593 }
594 }
595 }
596 }
597 }
598 }
599 }
600 }
601 }
602 }
603 }
604 }
605 }
606 }
607 }
608 }
609 }
610 }
611 }
612 }
613 }
614 }
615 }
616 }
617 }
618 }
619 }
620 }
621 }
622 }
623 }
624 }
625 }
626 }
627 }
628 }
629 }
630 }
631 }
632 }
633 }
634 }
635 }
636 }
637 }
638 }
639 }
640 }
641 }
642 }
643 }
644 }
645 }
646 }
647 }
648 }
649 }
650 }
651 }
652 }
653 }
654 }
655 }
656 }
657 }
658 }
659 }
660 }
661 }
662 }
663 }
664 }
665 }
666 }
667 }
668 }
669 }
670 }
671 }
672 }
673 }
674 }
675 }
676 }
677 }
678 }
679 }
680 }
681 }
682 }
683 }
684 }
685 }
686 }
687 }
688 }
689 }
690 }
691 }
692 }
693 }
694 }
695 }
696 }
697 }
698 }
699 }
700 }
701 }
702 }
703 }
704 }
705 }
706 }
707 }
708 }
709 }
710 }
711 }
712 }
713 }
714 }
715 }
716 }
717 }
718 }
719 }
720 }
721 }
722 }
723 }
724 }
725 }
726 }
727 }
728 }
729 }
730 }
731 }
732 }
733 }
734 }
735 }
736 }
737 }
738 }
739 }
740 }
741 }
742 }
743 }
744 }
745 }
746 }
747 }
748 }
749 }
750 }
751 }
752 }
753 }
754 }
755 }
756 }
757 }
758 }
759 }
760 }
761 }
762 }
763 }
764 }
765 }
766 }
767 }
768 }
769 }
770 }
771 }
772 }
773 }
774 }
775 }
776 }
777 }
778 }
779 }
780 }
781 }
782 }
783 }
784 }
785 }
786 }
787 }
788 }
789 }
790 }
791 }
792 }
793 }
794 }
795 }
796 }
797 }
798 }
799 }
800 }
801 }
802 }
803 }
804 }
805 }
806 }
807 }
808 }
809 }
810 }
811 }
812 }
813 }
814 }
815 }
816 }
817 }
818 }
819 }
820 }
821 }
822 }
823 }
824 }
825 }
826 }
827 }
828 }
829 }
830 }
831 }
832 }
833 }
834 }
835 }
836 }
837 }
838 }
839 }
840 }
841 }
842 }
843 }
844 }
845 }
846 }
847 }
848 }
849 }
850 }
851 }
852 }
853 }
854 }
855 }
856 }
857 }
858 }
859 }
860 }
861 }
862 }
863 }
864 }
865 }
866 }
867 }
868 }
869 }
870 }
871 }
872 }
873 }
874 }
875 }
876 }
877 }
878 }
879 }
880 }
881 }
882 }
883 }
884 }
885 }
886 }
887 }
888 }
889 }
890 }
891 }
892 }
893 }
894 }
895 }
896 }
897 }
898 }
899 }
900 }
901 }
902 }
903 }
904 }
905 }
906 }
907 }
908 }
909 }
910 }
911 }
912 }
913 }
914 }
915 }
916 }
917 }
918 }
919 }
920 }
921 }
922 }
923 }
924 }
925 }
926 }
927 }
928 }
929 }
930 }
931 }
932 }
933 }
934 }
935 }
936 }
937 }
938 }
939 }
940 }
941 }
942 }
943 }
944 }
945 }
946 }
947 }
948 }
949 }
950 }
951 }
952 }
953 }
954 }
955 }
956 }
957 }
958 }
959 }
960 }
961 }
962 }
963 }
964 }
965 }
966 }
967 }
968 }
969 }
970 }
971 }
972 }
973 }
974 }
975 }
976 }
977 }
978 }
979 }
980 }
981 }
982 }
983 }
984 }
985 }
986 }
987 }
988 }
989 }
990 }
991 }
992 }
993 }
994 }
995 }
996 }
997 }
998 }
999 }
1000 }

```

Figure 14: The example when improper exception handling directly impacts the loop stride. `ShellCommandExecutor.execute()` causes `IOException`. The exception is simply logged, and the creation of the `pidFile` is silently failed (i.e., zero-stride), which makes `File.exists()` always be false.

```

//Scrubber.java      #Cassandra-9881(v2.0.8)
44 private final RandomAccessReader dataFile;
...
103 public void scrub() {
120     while (!dataFile.isEOF()) {
129         try {
130             key = sstable.partitioner.decorateKey( //key=null
131                 ByteBufferUtil.readWithShortLength(dataFile));
134             dataSize = dataFile.readLong(); //skipped
139         } catch (Throwable th) {
140             throwIfFatal(th); //ignore Exception
141         }
185         try {
186             if (key == null)
187                 throw new IOError(...);
207         } catch (Throwable th) {
208             throwIfFatal(th); //ignore IOError
209         }
210     }
211 }
212 }
213 }
214 }
215 }
216 }
217 }
218 }
219 }
220 }
221 }
222 }
223 }
224 }
225 }
226 }
227 }
228 }
229 }
230 }
231 }
232 }
233 }
234 }
235 }
236 }
237 }
238 }
239 }
240 }
241 }
242 }
243 }
244 }
245 }
246 }
247 }
248 }
249 }
250 }
251 }
252 }
253 }
254 }
255 }
256 }
257 }
258 }
259 }
260 }
261 }
262 }
263 }
264 }
265 }
266 }
267 }
268 }
269 }
270 }
271 }
272 }
273 }
274 }
275 }
276 }
277 }
278 }
279 }
280 }
281 }
282 }
283 }
284 }
285 }
286 }
287 }
288 }
289 }
290 }
291 }
292 }
293 }
294 }
295 }
296 }
297 }
298 }
299 }
300 }
301 }
302 }
303 }
304 }
305 }
306 }
307 }
308 }
309 }
310 }
311 }
312 }
313 }
314 }
315 }
316 }
317 }
318 }
319 }
320 }
321 }
322 }
323 }
324 }
325 }
326 }
327 }
328 }
329 }
330 }
331 }
332 }
333 }
334 }
335 }
336 }
337 }
338 }
339 }
340 }
341 }
342 }
343 }
344 }
345 }
346 }
347 }
348 }
349 }
350 }
351 }
352 }
353 }
354 }
355 }
356 }
357 }
358 }
359 }
360 }
361 }
362 }
363 }
364 }
365 }
366 }
367 }
368 }
369 }
370 }
371 }
372 }
373 }
374 }
375 }
376 }
377 }
378 }
379 }
380 }
381 }
382 }
383 }
384 }
385 }
386 }
387 }
388 }
389 }
390 }
391 }
392 }
393 }
394 }
395 }
396 }
397 }
398 }
399 }
400 }
401 }
402 }
403 }
404 }
405 }
406 }
407 }
408 }
409 }
410 }
411 }
412 }
413 }
414 }
415 }
416 }
417 }
418 }
419 }
420 }
421 }
422 }
423 }
424 }
425 }
426 }
427 }
428 }
429 }
430 }
431 }
432 }
433 }
434 }
435 }
436 }
437 }
438 }
439 }
440 }
441 }
442 }
443 }
444 }
445 }
446 }
447 }
448 }
449 }
450 }
451 }
452 }
453 }
454 }
455 }
456 }
457 }
458 }
459 }
460 }
461 }
462 }
463 }
464 }
465 }
466 }
467 }
468 }
469 }
470 }
471 }
472 }
473 }
474 }
475 }
476 }
477 }
478 }
479 }
480 }
481 }
482 }
483 }
484 }
485 }
486 }
487 }
488 }
489 }
490 }
491 }
492 }
493 }
494 }
495 }
496 }
497 }
498 }
499 }
500 }
501 }
502 }
503 }
504 }
505 }
506 }
507 }
508 }
509 }
510 }
511 }
512 }
513 }
514 }
515 }
516 }
517 }
518 }
519 }
520 }
521 }
522 }
523 }
524 }
525 }
526 }
527 }
528 }
529 }
530 }
531 }
532 }
533 }
534 }
535 }
536 }
537 }
538 }
539 }
540 }
541 }
542 }
543 }
544 }
545 }
546 }
547 }
548 }
549 }
550 }
551 }
552 }
553 }
554 }
555 }
556 }
557 }
558 }
559 }
560 }
561 }
562 }
563 }
564 }
565 }
566 }
567 }
568 }
569 }
570 }
571 }
572 }
573 }
574 }
575 }
576 }
577 }
578 }
579 }
580 }
581 }
582 }
583 }
584 }
585 }
586 }
587 }
588 }
589 }
590 }
591 }
592 }
593 }
594 }
595 }
596 }
597 }
598 }
599 }
600 }
601 }
602 }
603 }
604 }
605 }
606 }
607 }
608 }
609 }
610 }
611 }
612 }
613 }
614 }
615 }
616 }
617 }
618 }
619 }
620 }
621 }
622 }
623 }
624 }
625 }
626 }
627 }
628 }
629 }
630 }
631 }
632 }
633 }
634 }
635 }
636 }
637 }
638 }
639 }
640 }
641 }
642 }
643 }
644 }
645 }
646 }
647 }
648 }
649 }
650 }
651 }
652 }
653 }
654 }
655 }
656 }
657 }
658 }
659 }
660 }
661 }
662 }
663 }
664 }
665 }
666 }
667 }
668 }
669 }
670 }
671 }
672 }
673 }
674 }
675 }
676 }
677 }
678 }
679 }
680 }
681 }
682 }
683 }
684 }
685 }
686 }
687 }
688 }
689 }
690 }
691 }
692 }
693 }
694 }
695 }
696 }
697 }
698 }
699 }
700 }
701 }
702 }
703 }
704 }
705 }
706 }
707 }
708 }
709 }
710 }
711 }
712 }
713 }
714 }
715 }
716 }
717 }
718 }
719 }
720 }
721 }
722 }
723 }
724 }
725 }
726 }
727 }
728 }
729 }
730 }
731 }
732 }
733 }
734 }
735 }
736 }
737 }
738 }
739 }
740 }
741 }
742 }
743 }
744 }
745 }
746 }
747 }
748 }
749 }
750 }
751 }
752 }
753 }
754 }
755 }
756 }
757 }
758 }
759 }
760 }
761 }
762 }
763 }
764 }
765 }
766 }
767 }
768 }
769 }
770 }
771 }
772 }
773 }
774 }
775 }
776 }
777 }
778 }
779 }
780 }
781 }
782 }
783 }
784 }
785 }
786 }
787 }
788 }
789 }
790 }
791 }
792 }
793 }
794 }
795 }
796 }
797 }
798 }
799 }
800 }
801 }
802 }
803 }
804 }
805 }
806 }
807 }
808 }
809 }
810 }
811 }
812 }
813 }
814 }
815 }
816 }
817 }
818 }
819 }
820 }
821 }
822 }
823 }
824 }
825 }
826 }
827 }
828 }
829 }
830 }
831 }
832 }
833 }
834 }
835 }
836 }
837 }
838 }
839 }
840 }
841 }
842 }
843 }
844 }
845 }
846 }
847 }
848 }
849 }
850 }
851 }
852 }
853 }
854 }
855 }
856 }
857 }
858 }
859 }
860 }
861 }
862 }
863 }
864 }
865 }
866 }
867 }
868 }
869 }
870 }
871 }
872 }
873 }
874 }
875 }
876 }
877 }
878 }
879 }
880 }
881 }
882 }
883 }
884 }
885 }
886 }
887 }
888 }
889 }
890 }
891 }
892 }
893 }
894 }
895 }
896 }
897 }
898 }
899 }
900 }
901 }
902 }
903 }
904 }
905 }
906 }
907 }
908 }
909 }
910 }
911 }
912 }
913 }
914 }
915 }
916 }
917 }
918 }
919 }
920 }
921 }
922 }
923 }
924 }
925 }
926 }
927 }
928 }
929 }
930 }
931 }
932 }
933 }
934 }
935 }
936 }
937 }
938 }
939 }
940 }
941 }
942 }
943 }
944 }
945 }
946 }
947 }
948 }
949 }
950 }
951 }
952 }
953 }
954 }
955 }
956 }
957 }
958 }
959 }
960 }
961 }
962 }
963 }
964 }
965 }
966 }
967 }
968 }
969 }
970 }
971 }
972 }
973 }
974 }
975 }
976 }
977 }
978 }
979 }
980 }
981 }
982 }
983 }
984 }
985 }
986 }
987 }
988 }
989 }
990 }
991 }
992 }
993 }
994 }
995 }
996 }
997 }
998 }
999 }
1000 }

```

Figure 15: The example when improper exception handling indirectly impacts the loop stride. Data corruption causes the I/O function `decorateKey()` to throw exception at line #130-131, which makes the loop skip the index updating statement (i.e., zero-stride) at line #134.

be put into the `rackToBlocks` map with line #554 skipped. This would eventually cause the while loop on line #348 to hang. The reason is that this while loop keeps iterating until every block in `blockToNodes` is removed. Unfortunately, since only blocks that also exist in `rackToBlocks` map can be removed (line #369 – #371), the corrupted block will never be removed from `blockToNodes` and cause an infinite loop.

Type 3: Improper exception handling directly affects loop strides. Sometimes, a data-related operation itself is expected to

Table 3: The cloud server systems used in our evaluation and the number of detected data corruption hang bugs in each system.

System	Description	# of bugs
Cassandra	Distributed database management system	2
Compress	Libraries for I/O ops on compressed file	2
HD Common	Hadoop utilities and libraries	10
Mapreduce	Hadoop big data processing framework	5
HDFS	Hadoop distributed file system	4
Yarn	Hadoop resource management platform	4
Hive	Data warehouse	12
Kafka	Distributed streaming platform	1
Lucene	Indexing and search server	2
Total		42

update the loop stride. When this operation throws an exception, an improper exception handling may give up the operation, together with the associated stride updates, causing infinite loops. For example, the Yarn-6991 bug belongs to this type, shown by Figure 14. The `ShellCommandExecutor.execute()` function is expected to create a `pidFile`, whose existence will help a while loop (line #89) to exit. When the disk is full, `ShellCommandExecutor.execute()` throws an exception at line #74. This exception is simply logged. Consequently, without the creation of `pidFile`, the while loop at line #89 never exits.

Type 4: Improper exception handling indirectly affects loop strides. For this type of bugs, the stride-update operation itself did not raise any exceptions. However, an exception handling of another operation, a data-related operation, changes the control flow and causes the stride update to be skipped. For example, the Cassandra-9881 bug matches this type, shown by Figure 15. When the `dataFile` (`RandomAccessReader` at line #131) is corrupted, the `decorateKey()` function cannot recognize it, thus throws an exception without assigning `key` at line #130 (i.e., `key == null`), or executing `dataFile.readLong()` at line #134. But this exception is simply ignored because it's not fatal at line #140. When the `key` is null, the `scrub()` function throws an `IOError` (line #187), catches it (line #207), and ignores it because it's not a fatal error (line #208). Without moving the index (i.e., zero-stride) by calling `dataFile.readLong()` at line #134, the `scrub()` function keeps reading from the same place, looping forever.

Discussion Theoretically, other types of corruption-hang bugs could exist, like corruption affecting loop bounds, instead of loop strides, or corrupted data content directly, instead of indirectly, affects loop strides. DScope bug detection algorithm *can* detect those types of bugs too. However, we did not observe them in the real-world bugs that we have encountered.

4 EVALUATION

In this section, we present our experimental evaluations on DScope. We first describe our evaluation methodology and then discuss our evaluation results in detail.

4.1 Evaluation Methodology

DScope is implemented on top of Soot v2.5.0 [8], a Java bytecode analysis infrastructure, with the latest Coffi library [43], written in Java language with about 18,000 lines of code. Our experimental evaluation covers a wide range of popular cloud server systems listed in Table 3: Cassandra is a distributed key-value store; Compress provides libraries for I/O operations on compressed files; Hadoop common provides utilities and libraries for all Hadoop projects; Hadoop MapReduce is a big data processing platform; HDFS is a distributed file system; Hadoop Yarn is a distributed resource management service; Hive is a data warehouse; Kafka is a distributed streaming system; and Lucene is a data indexing and searching server. We try to cover as many cloud server systems as possible to show that data corruption hang bugs are widespread in the real world.

All the experiments were conducted in our lab machine with an Intel® Xeon® E5-1630 Octa-core 3.7GHz CPU, 16GB memory, running 64-bit Ubuntu 16.04 with kernel v4.13.0. Our evaluation considers both coverage (i.e., true positives) and precision (i.e., false positives) of data corruption hang bug detection. We also compare DScope with two state-of-the-art static bug detection tools, Findbugs(v3.0.1) [7] and Infer(v0.9.2) [6].

For all the hang bugs reported by DScope, we first manually validate them by checking whether we can reproduce the software hang symptom after injecting data corruption into the corresponding data. We first check DScope’s analysis results to identify which faulty I/O operations affect the loop strides. We then inject the faults (e.g., corrupted data content, corrupted configuration files, disk exhaustion) into the corresponding I/O operations. If the software hang does happen, we mark the bug as a true positive. Otherwise, we consider it as a false positive. For all the true positives, we then search the bug repository (i.e., JIRA [5]) to see whether they are already reported. If they are, we mark them as the existing data corruption hang bugs. Otherwise, we report them in the bug repository and mark them as newly discovered data corruption hang bugs.

We then use the true positives detected by DScope as the benchmark to evaluate the detection efficacy of Findbugs and Infer. For these two tools, if they report at least one line of the code related to a data corruption hang bug (e.g. a line of the data corruption loop body or a line contains a variable which is then used in the loop), we consider the reported issue as a true positive. We omit the false positives of Findbugs and Infer in our evaluation because these two generic bug detection tools can report hundreds or thousands of suspicious issues. For example, Findbugs and Infer reports 5,434 and 13,993 issues in Hive v2.3.2, respectively. It is extremely time-consuming to validate all of their detection results manually. It is also wrong to label all the issues identified by Findbugs or Infer but not DScope as false positives since some of those issues are true bugs although they are not related to data corruption hang bugs.

4.2 Bug Detection and Precision Results

Table 4 and 5 show the detection results achieved by different schemes. DScope reports 79 data corruption hang bugs, with 42 of them being true bugs, and 29 out of the 42 bugs are newly discovered

Table 4: The detection comparison of DScope with Findbugs and Infer on all the 9 systems. “TP”: the number of true positive bugs by each scheme; “FP”: the number of false positive bugs reported by DScope; “-”: runtime execution errors (Infer).

System	Release date	DScope		Findbugs		Infer
		TP	FP	TP	TP	TP
Cassandra	v2.0.8	2014/05/29	2	1	0	1
Compress	v1.0	2009/05/21	2	2	0	-
HD Common	v0.23.0	2011/11/11	4	6	0	0
	v2.5.0	2014/08/11	6	6	0	0
Mapreduce	v0.23.0	2011/11/11	3	0	0	0
	v2.5.0	2014/08/11	2	0	0	0
HDFS	v0.23.0	2011/11/11	1	1	0	0
	v2.5.0	2014/08/11	3	5	1	-
Yarn	v0.23.0	2011/11/11	2	2	1	0
	v2.5.0	2014/08/11	2	5	0	0
Hive	v1.0.0	2015/05/20	7	6	0	-
	v2.3.2	2017/11/18	5	1	0	0
Kafka	v0.10.0.0	2016/05/22	1	1	0	0
Lucene	v2.1.0	2007/02/17	2	1	0	0
Total			42	37	2	1

bugs. Note that, we ran DScope on the target cloud server systems and identified those 42 bugs. But it does not mean that those 42 bugs include *all* the data corruption bugs in those systems. There are some other types of data corruption hang bugs that we cannot identify. For example, data corruption causes the recursive functions never end, making system hang. However, it is out of the scope of this paper, which is part of our future work.

In contrast, existing generic bug detection tools cannot detect most of those 42 data corruption hang bugs. Findbugs only identifies the HDFS-5892 and Yarn-163 bugs while Infer only identifies the Cassandra-9881 bug. Those results are expected because no previous static analysis tools, including Findbugs and Infer, have targeted data corruption hang bugs. Findbugs targets bugs that follow specific anti-patterns in Java programs, such as “private method is never called”, “method concatenates strings using + in a loop”, and “unchecked type in generic call”, none of which are related to data corruption hang bugs detected by DScope. Note that, Findbugs does contain one specific anti-pattern called “an apparent infinite loop” which is related to data corruption hang bugs. However, Findbugs only reports two suspicious issues on the target cloud server systems and both issues involve a `while(true)` type loop. After further inspection, these two loops can exit eventually due to timeouts. Infer mostly focuses on memory and resource leak bugs, and hence cannot detect most corruption-hang bugs shown in Table 5.

Findbugs identifies the HDFS-5892 bug, as it discovers `getFinalizedDir()` can be “null” in the loop body. This bug happens when corrupted data content indirectly affects the loop stride (i.e. the `getFinalizedDir().length` becomes 0). Indeed, the `getFinalizedDir()` function is called during the loop’s execution, but it is not the root cause of this data corruption hang bug. Findbugs identifies the Yarn-163 bug, as it discovers that encoding the `InputStreamReader` reader to a `FileReader` can corrupt the reader, which

Table 5: The detection comparison of DScope with Findbugs and Infer on all the 42 data corruption hang bugs.

#	Bug name	System version	Bug type	Known or new	Detected		
					DScope	Findbugs	Infer
1	Cassandra-7330	v2.0.8	#1	known	✓	✗	✗
2	Cassandra-9881	v2.0.8	#3	known	✓	✗	✓
3	Compress-87	v1.0	#1	known	✓	✗	✗
4	Compress-451	v1.0	#2	new	✓	✗	✗
5	Hadoop-8614	v0.23.0	#1	known	✓	✗	✗
6	Hadoop-15088	v2.5.0	#1	new	✓	✗	✗
7	Hadoop-15415	v0.23.0	#2	new	✓	✗	✗
8		v2.5.0	#2	new	✓	✗	✗
9	Hadoop-15417	v0.23.0	#2	new	✓	✗	✗
10		v2.5.0	#2	new	✓	✗	✗
11	Hadoop-15424	v2.5.0	#1	new	✓	✗	✗
12	Hadoop-15425	v2.5.0	#1	new	✓	✗	✗
13	Hadoop-15429	v0.23.0	#2	new	✓	✗	✗
14		v2.5.0	#2	new	✓	✗	✗
15	HDFS-4882	v0.23.0	#3	known	✓	✗	✗
16	HDFS-5892	v2.5.0	#2	known	✓	✓	✗
17	HDFS-13513	v2.5.0	#2	new	✓	✗	✗
18	HDFS-13514	v2.5.0	#2	new	✓	✗	✗
19	Mapreduce-2185	v0.23.0	#2	known	✓	✗	✗
20	Mapreduce-2862	v0.23.0	#2	known	✓	✗	✗
21	Mapreduce-6990	v0.23.0	#1	new	✓	✗	✗
24	Mapreduce-7088	v2.5.0	#1	new	✓	✗	✗
25	Mapreduce-7089	v2.5.0	#1	new	✓	✗	✗
26	Yarn-163	v0.23.0	#1	known	✓	✓	✗
27	Yarn-2905	v2.5.0	#1	known	✓	✗	✗
22	Yarn-6991	v0.23.0	#4	new	✓	✗	✗
23		v2.5.0	#4	new	✓	✗	✗
28	Hive-5235	v1.0.0	#1	known	✓	✗	✗
29	Hive-13397	v1.0.0	#2	known	✓	✗	✗
30	Hive-18142	v1.0.0	#2	new	✓	✗	✗
31	Hive-18216	v2.3.2	#1	new	✓	✗	✗
32	Hive-18217	v2.3.2	#1	new	✓	✗	✗
33	Hive-18219	v1.0.0	#2	new	✓	✗	✗
34		v2.3.2	#2	new	✓	✗	✗
35	Hive-19391	v1.0.0	#2	new	✓	✗	✗
36	Hive-19392	v1.0.0	#2	new	✓	✗	✗
37		v2.3.2	#2	new	✓	✗	✗
38	Hive-19395	v1.0.0	#1	new	✓	✗	✗
39	Hive-19406	v2.3.2	#2	new	✓	✗	✗
40	Kafka-6271	v0.10.0	#1	new	✓	✗	✗
41	Lucene-772	v2.1.0	#2	known	✓	✗	✗
42	Lucene-8294	v2.1.0	#2	new	✓	✗	✗
Total #					42	2	1

is related to the data corruption hang bugs — performing skip operations on a corrupted FileReader can cause the skip function to return error code (i.e., 0).

Infer identifies the Cassandra-9881 bug, as it discovers that the scrub() function in the Scrubber class could invoke a throwIfCommutative() function at line #248 with null parameter, shown by Figure 16. As we discussed in §3, when data corruption happens, key fails to be assigned to new values and sticks with the default

```
//cassandra-2.0.8: Scrubber.java
41 private boolean isCommutative = false;
...
103 public void scrub() {
...
120 while (!dataFile.isEOF()) {
...
127 DecoratedKey key = null;
...
248 throwIfCommutative(key, th); //Infer: null parameter
...
}}
```

```
327 private void throwIfCommutative(DecoratedKey key,
328                                 Throwable th) {
329     if ((isCommutative && !skipCorrupted) {
331         outputHandler.warn(String.format("...", key));
...
}}
```

Figure 16: Infer identifies a null parameter problem in the throwIfCommutative() function at line #248. The Cassandra-9881 bug happens at line #103-256.**Table 6: The types of false positives pruned by DScope.**

System		Pruned FP	
		Numeric primitives	Java APIs
Cassandra	v2.0.8	386	71
Compress	v1.0	147	20
HD Common	v0.23.0	1023	378
	v2.5.0	1650	790
Mapreduce	v0.23.0	377	363
	v2.5.0	938	641
HDFS	v0.23.0	312	323
	v2.5.0	1723	1073
Yarn	v0.23.0	151	214
	v2.5.0	451	665
Hive	v1.0.0	4268	3003
	v2.3.2	5269	3663
Kafka	v0.10.0.0	186	441
Lucene	v2.1.0	287	44
Total		17168	11689

value, “null”. This makes the scrub() function skip updating the index, causing an infinite loop. Indeed, the throwIfCommutative() function is called during the loop’s execution, but it is not the root cause of this data corruption hang bug. In fact, it does not break the loop to prevent scrub() from hanging. This is because the isCommutative variable is false, which makes the if branch at line #329 unreachable. Thus, even with a null parameter, the throwIfCommutative() can still execute successfully at line #248.

Table 5 also shows the types of the detected data corruption hang bugs. As we can see, “Type 1” and “Type 2” cover most of the detected bugs — 16 and 22 bugs respectively. This indicates that most of the data corruption hang bugs happen when the data corruption causes the error code returned by I/O operations to directly impact the loop stride or corrupted data content indirectly impacts the loop stride.

To understand how DScope does not prune all the false positives, we manually study those 37 false positives in Table 4. We find

most of cases require inter-procedural analysis to identify. We will discuss it in §5.

As shown in Table 6, DScope prunes 28,857 false positives in total, including 17,168 cases where the loop index, stride and bounds are denoted by numeric primitives, and 11,689 cases where the loop index, stride and bounds are denoted by commonly used Java APIs.

We should note that, we do not intend to claim that DScope can replace those generic bug detection tools such as Findbugs and Infer. We believe our bug detection schemes are complementary to those existing tools and could be used in combination by the software developer.

5 DISCUSSION

We observe that in most of the 37 false positive cases, the forwarding-index/reversing-index Java APIs and the checking-bounds Java APIs are located in different application functions. These APIs are indirectly invoked in the application functions which are invoked in the loop paths. To further reduce false positives, we plan to conduct inter-procedural analysis on all the bug candidates to generate the loop paths where the loop index, stride, and bounds are denoted by either numeric primitives or Java APIs. We then adopt DScope's false positive pruning principles to prune the false positives without missing true positives.

DScope focuses on detecting data-corruption hang bugs. We plan to explore auto-fix schemes to correct those bugs based on their types, as described in §3. For example, when the error code returned by the I/O operations directly affects the loop strides, one possible fix is to add extra error code checking statements in the loop exit conditions to avoid the software hang. If the corrupted data content indirectly affects loop strides, one possible fix could be adding additional check over the data content before it is used in the loop body. When the improper exception handling causes the loop stride update to be skipped, one possible fix is to add additional exception handling to properly update the stride when data corruption occurs.

6 RELATED WORK

Data corruption study and detection: Previous work has been extensively studied the data corruption problems in storage systems. Hwang et al. [27] and Schroeder et al. [39] studied the data corruptions in memory devices. They found that DRAM failures occur more frequently than expected. Bairavasundaram et al. [11, 12] and Oleksenko et al. [36] detected the data pointer corruptions on disks. They showed that disk failures are prevalent for data corruptions. Previous works have also been done to detect data corruptions in file systems. ZFS [14] detected file system corruption caused by storage hardware, e.g., latent sector errors. Fryer et al. [22, 41] implemented runtime data corruption detectors for the Ext3 and Btrfs file systems.

The above work provides motivations for us to study data corruption induced performance problems. Our work focuses on detecting data corruption hang bugs in software-level rather than detecting the data corruption itself (hardware-level).

Performance bug detection and diagnosis: Much work has been done to detect and diagnose performance bugs in large scale

systems. X-ray [9] uses symbolic execution to automatically identify and suggest fixes to performance bugs caused by configuration or input-based problems. Xu et al. [47] presented a clustering-based scheme to detect system anomalies. Jin et al. [29] employed rule-based methods to detect performance bugs that violate known efficiency rules. Caramel [33] statically detects inefficient loops that can be fixed by adding conditional-breaks. LDoctor [40] provides statistical diagnosis for inefficient loops. Jolt [16] dynamically detects infinite loops by checking each loop iteration's run-time state. Tools also exist to detect inefficient nested loops [34] and workload-dependent loops [46].

In comparison, our work focuses on detecting data corruption induced software hang problems before they are triggered in the production system. We adopt a pattern-driven approach instead of rule-based or anomaly detection approaches to achieving both high coverage and precision for our targeted data corruption hang bugs.

Previous work Carburizer [30] statically analyzes device driver code and identifies infinite driver-polling problems. That is, a driver may wait for a device to enter a given state by polling a device register. Once the register data is corrupted, a buggy driver may be stuck forever. DScope and Carburizer both statically analyze loops and loop-exit conditions. However, they face different design challenges due to the different types of bugs they target. DScope targets cloud systems written in Java, instead of low-level device drivers, and hence needs to handle a much broader set of I/O functions and I/O related data (e.g., not only data retrieved by I/O operations but also status returned by I/O operations), and more complicated control flows caused by Java exceptions. Carburizer false-positive pruning only involves identifying loop timeouts. However, DScope has to conduct sophisticated loop stride and bound analysis in its false-positive pruning. Finally, as indicated in §3, the type of corruption-hang bugs identified by DScope in cloud systems go much beyond simple I/O-state infinite polling problems, where the device register content often directly updates the loop stride.

Fault injection: Previous work [13, 24, 42] used fault injection techniques to analyze the failure behaviors (e.g., hang, crash) of both software and hardware systems. For example, HSFI [42] injected faults in the source code. Fault injection is also widely used to check whether file systems can handle certain type of data corruptions [21, 23, 38, 48]. For instance, Bairavasundaram et al. [10] used context aware fault injections to find disk errors in virtual memory systems. Zhang et al. [48] conducted a comprehensive reliability case study of local file systems to analyze both on-disk and in-memory data integrity in Sun's ZFS. Their results show that file systems are robust to disk corruption but less resilient to memory corruptions. Cords [23] exposed data losses, block corruptions, and unavailability problems commonly exist in distributed file systems. Cords also indicated that modern distributed file systems are not equipped to effectively use redundancy across replicas to recover from local file system faults. In contrast, our work focuses on detecting potential data corruption hang bugs before they are triggered by the data corruption faults. We only rely on static code analysis, which can be easily applied to different cloud server systems. We believe our work is complementary to the fault injection

based approaches which can be used to validate our candidate bugs and further reduce false positives.

Functional bug detection: Apart from performance bugs, recent works have also been done to detect functional bugs. pbSE [45] conducted concolic execution to detect functional bugs and generate test cases for those bugs. Kollenda et al. [32] detected the crash bugs by identifying the crash-resistant primitives via system calls on Linux, Windows API functions, and exception handlers. In contrast, our work focuses on detecting performance bugs, which requires the bug detection system to focus on different aspects of the program such as loop exit checking.

Software testing: DeepXplore [37] is a whitebox framework to test deep learning systems. DeepXplore takes unlabeled test inputs as seeds in DNN systems. It uses gradient ascent to modify the input to maximize chance of finding rare corner cases. Fex [35] is a software system evaluator, which collects a set of reused scripts to develop a matured evaluation framework. Fex addressed the limitation of rigid, simplistic and inconsistent in large system testing. Elia et al. [20] designed an interoperability certification model, which facilitates testing interoperability among different web applications. Our work is complementary to those software testing tools. Our tool can identify potential buggy functions with infinite loops, which can guide the test case generation to further test our detection results.

7 CONCLUSION

In this paper, we have presented DScope, a new data corruption hang bug detection tool for cloud server systems. DScope combines candidate bug discovery and false positive pattern filtering to detect software hang bugs that are related to data corruptions. DScope is fully automatic without requiring any user input or pre-defined rules. We have implemented a prototype of DScope and evaluated it over 9 commonly used cloud server systems. DScope successfully detects 42 true corruption hang bugs (29 of them are new bugs) while existing bug detection tools can only detect very few of them (2 by Findbugs and 1 by Infer).

ACKNOWLEDGMENTS

We would like to thank Ennan Zhai, our shepherd, and the anonymous reviewers for their insightful feedback and valuable comments. This work was sponsored in part by NSF CNS1513942 grant and NSF CNS1149445 grant. Any opinions expressed in this paper are those of the authors and do not necessarily reflect the views of NSF or U.S. Government.

REFERENCES

- [1] 2013. HDFS-4882. <https://issues.apache.org/jira/browse/HDFS-4882>.
- [2] 2017. What Lessons can be Learned from BA's Systems Outage? <http://www.extraordinarymanagementservices.com/news/what-lessons-can-be-learned-from-bas-systems-outage/>.
- [3] 2018. Apache Cassandra. <http://cassandra.apache.org/>.
- [4] 2018. Apache Hadoop. <http://hadoop.apache.org/>.
- [5] 2018. Apache JIRA. <https://issues.apache.org/jira>.
- [6] 2018. Facebook Infer. <http://fbinfer.com/>.
- [7] 2018. Findbugs. <http://findbugs.sourceforge.net/>.
- [8] 2018. Soot: A Framework for Analyzing and Transforming Java and Android Applications. <https://sable.github.io/soot/>.
- [9] Mona Attariyan, Michael Chow, and Jason Flinn. 2012. X-ray: Automating Root-Cause Diagnosis of Performance Anomalies in Production Software. In *OSDI*.
- [10] Lakshmi N. Bairavasundaram, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. 2006. Dependability Analysis of Virtual Memory Systems. In *DSN*.
- [11] Lakshmi N. Bairavasundaram, Andrea C. Arpaci-Dusseau, Remzi H. Arpaci-Dusseau, Garth R. Goodson, and Bianca Schroeder. 2008. An Analysis of Data Corruption in the Storage Stack. *TOS* 4, 3 (nov 2008), 8:1–8:28.
- [12] Lakshmi N. Bairavasundaram, Meenali Rungta, Nitin Agrawal, Andrea C. Arpaci-Dusseau, Remzi H. Arpaci-Dusseau, and Michael M. Swift. 2008. Analyzing the Effects of Disk-Pointer Corruption. In *DSN*.
- [13] James H. Barton, Edward W. Czeck, Zary Z. Segall, and Daniel P. Siewiorek. 1990. Fault Injection Experiments Using FIAT. *TC* 39, 4 (apr 1990).
- [14] Jeff Bonwick and Bill Moore. 2007. *ZFS—The Last Word In File Systems*. https://wiki.illumos.org/download/attachments/1146951/zfs_last.pdf.
- [15] Nedyalko Borisov, Shivnath Babu, Nagapramod Mandagere, and Sandeep Utamchandani. 2011. Dealing Proactively with Data Corruption: Challenges and Opportunities. In *SMDB*.
- [16] Michael Carbin, Sasa Misailovic, Michael Kling, and Martin C. Rinard. 2011. Detecting and Escaping Infinite Loops with Jolt. In *ECOOP*.
- [17] Ting Dai, Daniel Dean, Peipei Wang, Xiaohui Gu, and Shan Lu. 2018. Hytrace: A Hybrid Approach to Performance Bug Diagnosis in Production Cloud Infrastructures. *IEEE Transactions on Parallel and Distributed Systems* (2018).
- [18] Daniel J. Dean, Hiep Nguyen, Xiaohui Gu, Hui Zhang, Junghwan Rhee, Nipun Arora, and Geoff Jiang. 2014. PerfScope: Practical Online Server Performance Bug Inference in Production Cloud Computing Infrastructures. In *SOCC*.
- [19] Daniel J. Dean, Peipei Wang, Xiaohui Gu, William Enck, and Guoliang Jin. 2015. Automatic Server Hang Bug Diagnosis: Feasible Reality or Pipe Dream?. In *ICAC*.
- [20] Ivano Alessandro Elia, Nuno Laranjeiro, and Marco Vieira. 2015. Test-Based Interoperability Certification for Web Services. In *DSN*.
- [21] David Fiala, Frank Mueller, Christian Engelmann, Rolf Riesen, Kurt Ferreira, and Ron Brightwell. 2012. Detection and Correction of Silent Data Corruption for Large-scale High-performance Computing. In *SC*.
- [22] Daniel Fryer, Mike Qin, Jack Sun, Kah Wai Lee, Angela Demke Brown, and Ashvin Goel. 2014. Checking the Integrity of Transactional Mechanisms. *TOS* 10, 4 (oct 2014), 17:1–17:23.
- [23] Aishwarya Ganesan, Rammatthan Alagappan, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. 2017. Redundancy Does Not Imply Fault Tolerance: Analysis of Distributed Storage Reactions to Single Errors and Corruptions. In *FAST*.
- [24] Weining Gu, Zbigniew Kalbarczyk, Ravishankar K. Iyer, and Zhen-Yu Yang. 2003. Characterization of Linux Kernel Behavior Under Errors. In *DSN*.
- [25] Haryadi S. Gunawi, Mingzhe Hao, Tanakorn Leesatapornwongsa, Tirat Patanana-ake, Thanh Do, Jeffry Adityatama, Kurnia J. Eliazar, Agung Laksono, Jeffrey F. Lukman, Vincentius Martin, et al. 2014. What Bugs Live in the Cloud?: A Study of 3000+ Issues in Cloud Systems. In *SOCC*.
- [26] Jian Huang, Xuechen Zhang, and Karsten Schwan. 2015. Understanding Issue Correlations: A Case Study of the Hadoop System. In *SOCC*.
- [27] Andy A. Hwang, Ioan A. Stefanovici, and Bianca Schroeder. 2012. Cosmic Rays Don't Strike Twice: Understanding the Nature of DRAM Errors and the Implications for System Design. In *ASPLOS*.
- [28] Weihang Jiang, Chongfeng Hu, Arkady Kanevsky, and Yuanyuan Zhou. 2008. Is Disk the Dominant Contributor for Storage Subsystem Failures? A Comprehensive Study of Failure Characteristics. In *FAST*.
- [29] Guoliang Jin, Linhai Song, Xiaoming Shi, Joel Scherpelz, and Shan Lu. 2012. Understanding and Detecting Real-World Performance Bugs. In *PLDI*.
- [30] Asim Kadav, Matthew J. Renzelmann, and Michael M. Swift. 2009. Tolerating Hardware Device Failures in Software. In *SOSP*.
- [31] Jonathan Kaldor, Jonathan Mace, Michal Bejda, Edison Gao, Wiktor Kuropatwa, Joe O'Neill, Kian Win Ong, Bill Schaller, Pingjia Shan, Brendan Viscomi, Vinod Venkataraman, Kaushik Veeraraghavan, and Yee Jiun Song. 2017. Canopy: An End-to-End Performance Tracing And Analysis System. In *SOSP*.
- [32] Benjamin Kollenda, Enes Göktas, Tim Blazytko, Philipp Koppe, Robert Gawlik, RK Konoth, Cristiano Giuffrida, Herbert Bos, and Thorsten Holz. 2017. Towards Automated Discovery of Crash-Resistant Primitives in Binary Executables. In *DSN*.
- [33] Adrian Nistor, Po-Chun Chang, Cosmin Radoi, and Shan Lu. 2015. Caramel: Detecting and Fixing Performance Problems That Have Non-intrusive Fixes. In *ICSE*.
- [34] Adrian Nistor, Linhai Song, Darko Marinov, and Shan Lu. 2013. Toddler: Detecting Performance Problems via Similar Memory-Access Patterns. In *ICSE*.
- [35] Oleksii Oleksenko, Dmitrii Kuvaitskii, Pramod Bhatotia, and Christof Fetzer. 2017. Fex: A Software Systems Evaluator. In *DSN*.
- [36] Oleksii Oleksenko, Dmitrii Kuvaitskii, Pramod Bhatotia, Christof Fetzer, and Pascal Felber. 2016. Efficient Fault Tolerance using Intel MPX and TSX. In *DSN*.
- [37] Kexin Pei, Yinzi Cao, Junfeng Yang, and Suman Jana. 2017. DeepXplore: Automated Whitebox Testing of Deep Learning Systems. In *SOSP*.
- [38] Vijayan Prabhakaran, Lakshmi N. Bairavasundaram, Nitin Agrawal, Haryadi S. Gunawi, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. 2005. IROn File Systems. In *SOSP*.

- [39] Bianca Schroeder, Eduardo Pinheiro, and Wolf-Dietrich Weber. 2009. DRAM Errors in the Wild: A Large-Scale Field Study. In *SIGMETRICS*.
- [40] Linhai Song and Shan Lu. 2017. Performance Diagnosis for Inefficient Loops. In *ICSE*.
- [41] Kuei Sun, Daniel Fryer, Dai Qin, Angela Demke Brown, and Ashvin Goel. 2014. Robust Consistency Checking for Modern Filesystems. In *RV*.
- [42] Erik van der Kouwe and Andrew S Tanenbaum. 2016. HSFI: Accurate Fault Injection Scalable to Large Code Bases. In *DSN*.
- [43] Clark Verbrugge. 1996. *Using Coffi*. <http://www.sable.mcgill.ca/~clump/Coffi/Coffi.ps>.
- [44] Peipei Wang, Daniel J. Dean, and Xiaohui Gu. 2015. Understanding Real World Data Corruptions in Cloud Systems. In *IC2E*.
- [45] Qixue Xiao, Yu Chen, Chengang Wu, Kang Li, Junjie Mao, Shize Guo, and Yuanchun Shi. 2017. pbSE: Phase-Based Symbolic Execution. In *DSN*.
- [46] Xusheng Xiao, Shi Han, Dongmei Zhang, and Tao Xie. 2013. Context-Sensitive Delta Inference for Identifying Workload-Dependent Performance Bottlenecks. In *ISSTA*.
- [47] Kui Xu, Ke Tian, Danfeng Yao, and Barbara G. Ryder. 2016. A Sharper Sense of Self: Probabilistic Reasoning of Program Behaviors for Anomaly Detection with Context Sensitivity. In *DSN*.
- [48] Yupu Zhang, Abhishek Rajimwale, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. 2010. End-to-End Data Integrity for File Systems: A ZFS Case Study. In *FAST*.