

A Study on Container Vulnerability Exploit Detection

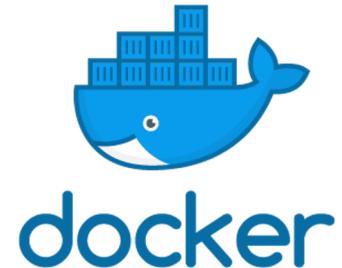
Olufogorehan Tunde-Onadele, Jingzhu He, Ting Dai, Xiaohui Gu

NC State University

IC2E '19

Motivation

- Great deployment benefits
 - Portability
 - Consistency
 - Isolation
- Recent platform still facing security issues



Real-World Examples

Juniper Networks: Cryptomining Exploit Targeting Docker Containers



by Michael Vizard on November 15, 2018

Attackers Continue to Target Docker Vulnerabilities

Brief, Security

Oct 29, 2018 By [Eric Carter](#), ProgrammableWeb Staff

Doomsday Docker security hole uncovered

A security vulnerability has been disclosed for a flaw in runc, Docker and Kubernetes' container runtime, which can be used to attack any host system running containers.



By [Steven J. Vaughan-Nichols](#) for [Linux and Open Source](#) | February 11, 2019 -- 18:53 GMT (10:53 PST) | Topic: Security



Detection Approaches

- 2 main categories of vulnerability detection
 - Static detection on container images
 - Dynamic detection on container instances

Static Detection Problems

Clair

- Container image scanning
- Depends on
 - Packages and versions
 - Vulnerability records
- Limitation
 - High false positive rate
 - Low detection rate



Dynamic Detection Challenges

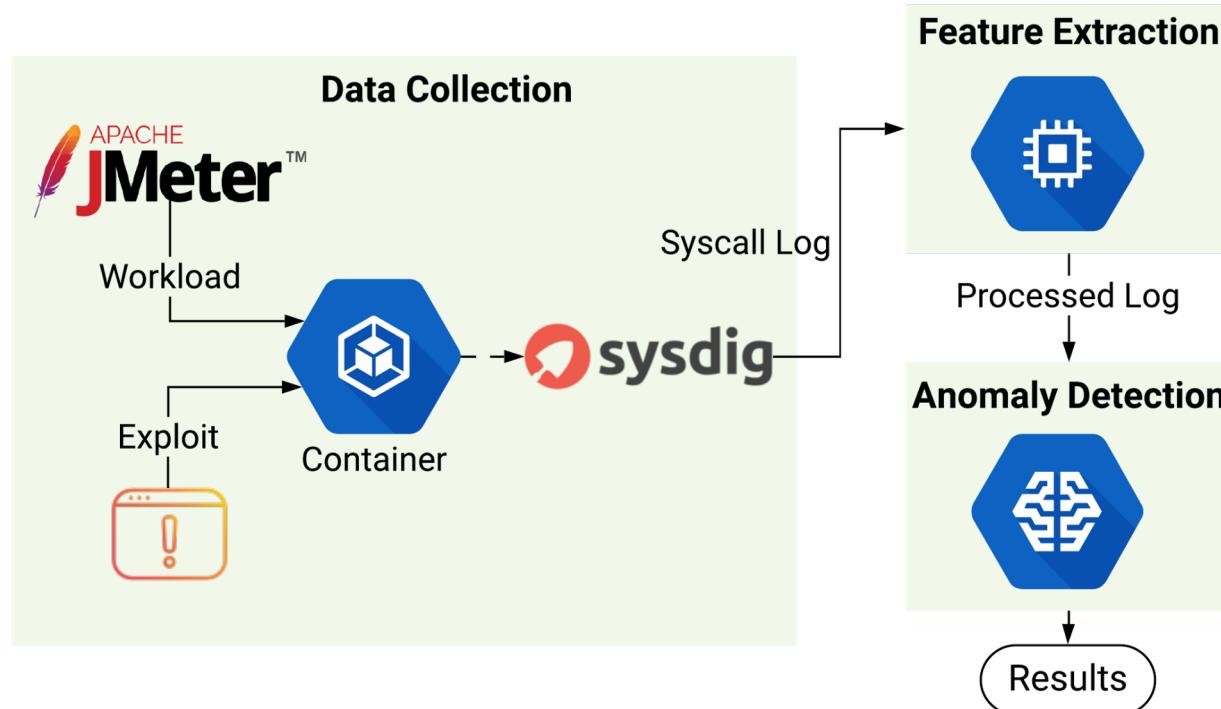
1. Containers are often short-lived
2. Containers have dynamic available resources & workloads
3. Containers are lightweight

Dynamic Detection

- Anomaly detection on system calls
- Studied Algorithms
 - k-Nearest Neighbors (k-NN)
 - Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)
 - k-Means
 - Self Organizing Map (SOM)

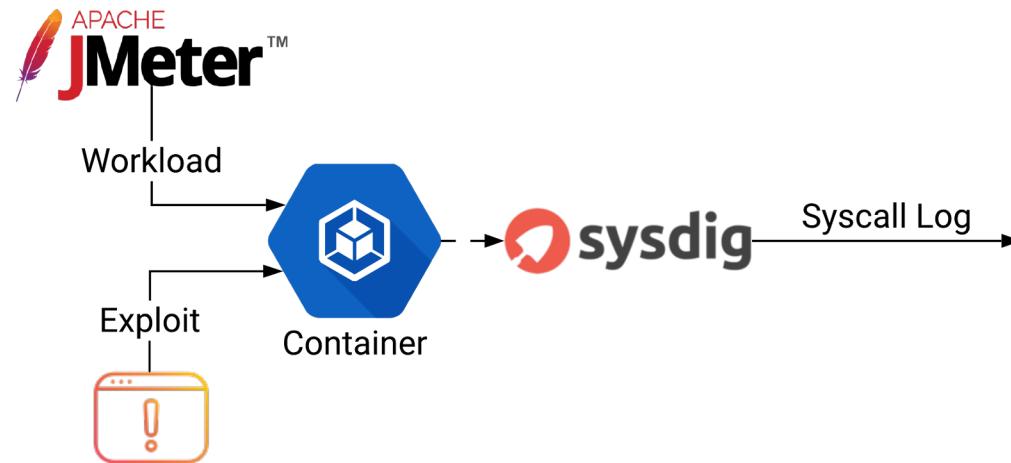
Experiment

- Setup



Data Collection

- Deliver workload for a particular application
- Trigger the exploit at appropriate time
- Collect a detailed log of the system call information



Data Processing

- Frequency vector
 - System call occurrences per sample

Timestamp	System Calls			
	write	read	futex	epoll_wait
1516544689186	100	256	430	78
1516544689286	300	759	726	356

- Time vector
 - System call runtime per sample

Case Study: ActiveMQ (CVE-2016-3088)

- File write vulnerability
 - Allows upload and execution of arbitrary files
 - Achieved using an HTTP PUT followed by an HTTP MOVE request
- Exploit requires
 - Knowledge of web directory
 - ActiveMQ run as root

Case Study: ActiveMQ (CVE-2016-3088)

- Exploit (Vulhub, 2017)
 - Set up a waiting shell
 - HTTP PUT payload with crontab commands
 - HTTP MOVE to a crontab location
 - Shell returned
1. **PUT /fileserver/1.txt HTTP/1.1**
2. **Host:** localhost:port#
...
7. **Content-Length:** 247
8. { *crontab command to initiate socket connection to shell*}

1. **MOVE /fileserver/1.txt HTTP/1.1**
2. **Destination:** file:///etc/cron.d/root
3. **Host:** localhost:port#
...
8. **Content-Length:** 0

Case Study: ActiveMQ (CVE-2016-3088)

- Sysdig log snippet

```
10:14:04.999140525 3 java (10306) > switch next=0 pgft maj=0 pgft_min=33...
10:14:05.049191227 3 java (10306) < futex res=-110(ETIMEDOUT)
10:14:05.049194706 3 java (10306) > futex addr=7F20F0236928 op=129... val=1
10:14:05.049195721 3 java (10306) < futex res=0
10:14:05.049202973 3 java (10306) > futex addr=7F20F0236954 op=137... val=1
...
10:14:05.089969920 3 java (10340) > getsockname
10:14:05.089971658 3 java (10340) < getsockname
10:14:05.089976977 3 java (10340) > getsockname
10:14:05.089978207 3 java (10340) < getsockname
...
10:14:05.099234302 3 java (10306) < futex res=-110(ETIMEDOUT)
10:14:05.099237402 3 java (10306) > futex addr=7F20F0236928 op=129...
10:14:05.099238268 3 java (10306) < futex res=0
10:14:05.099244128 3 java (10306) > futex addr=7F20F0236954 op=137...
10:14:05.099250022 3 java (10306) > switch next=0 pgft_maj=0 pgft_min=33...
10:14:05.128901808 2 java (10346) > write fd=137(<p>pipe:[1141873]) size=1
```

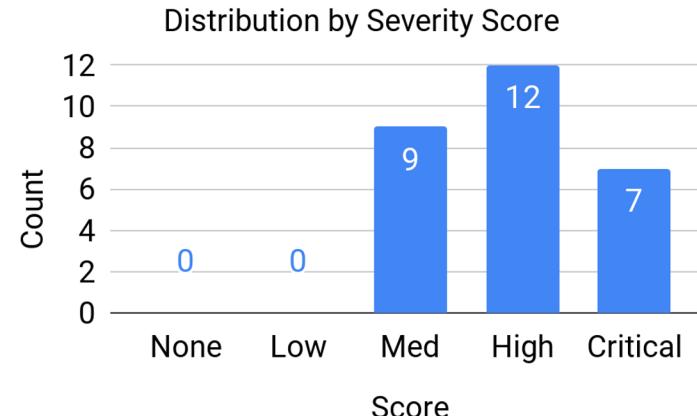
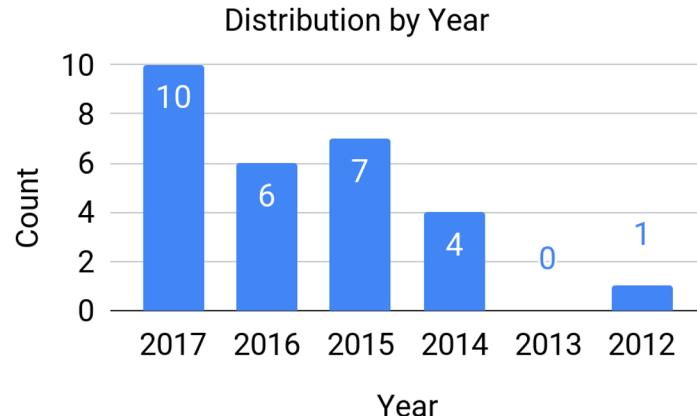
Case Study: ActiveMQ (CVE-2016-3088)

- Processed log snippet

timestamp	read	futex	accept	fcntl	getsockname
1528884844803	0	4	0	0	0
1528884844903	0	4	0	0	0
1528884845003	0	4	1	3	2
1528884845103	258	599	0	200	0
1528884845203	531	1542	0	436	0

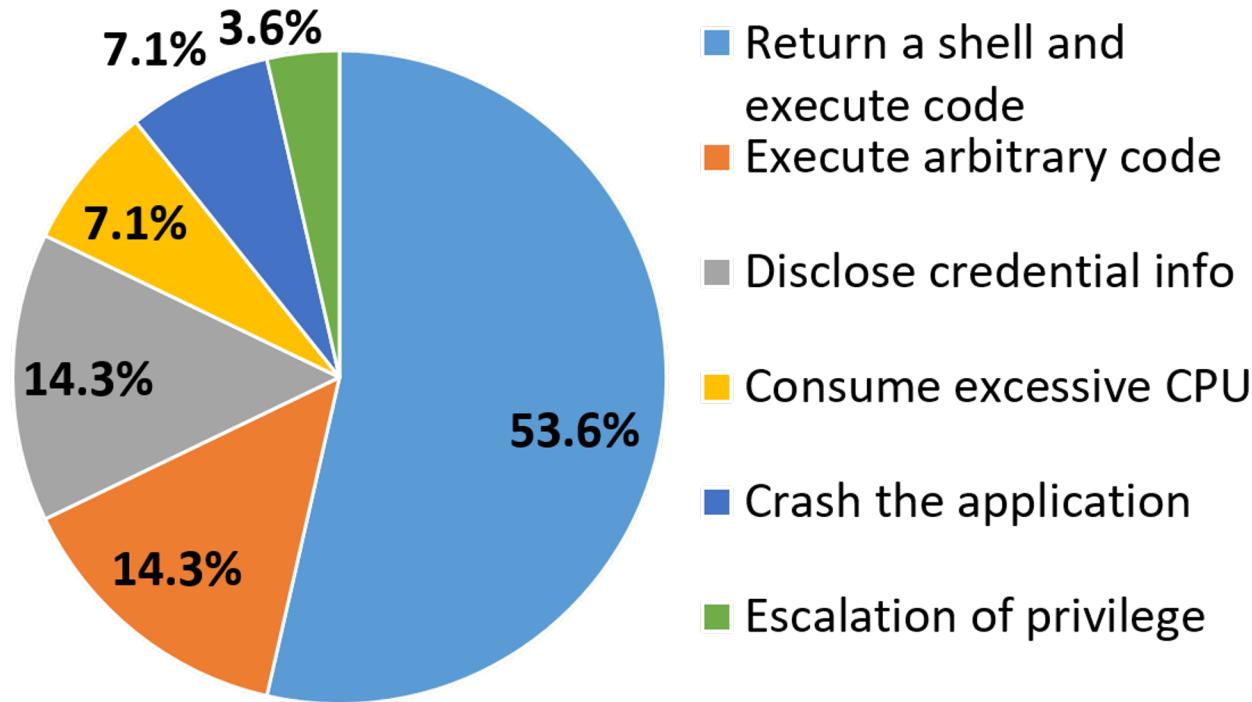
Studied Vulnerabilities

- 28 recent vulnerabilities of moderate to high severity
- Variety of applications
 - Web, file services



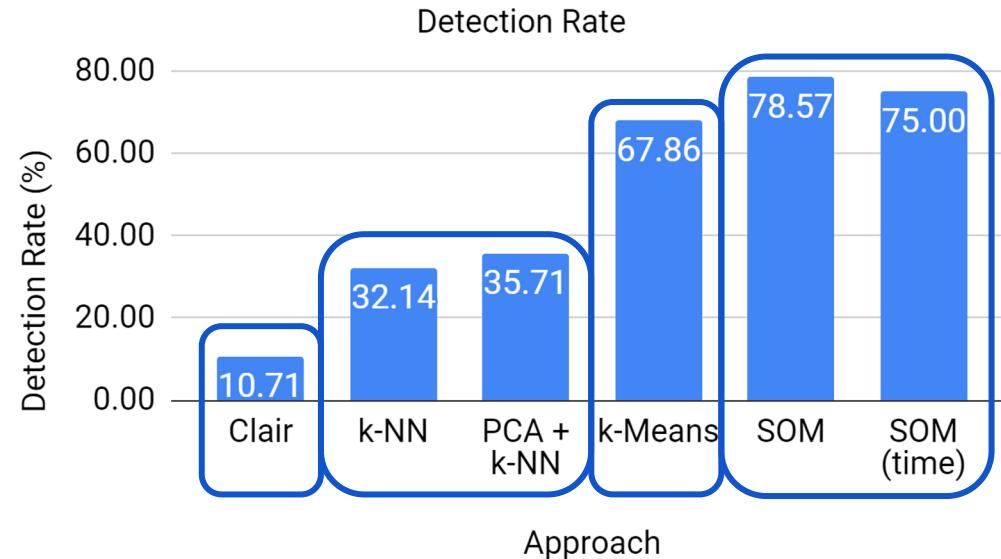
Studied Vulnerabilities

- Vulnerability categories



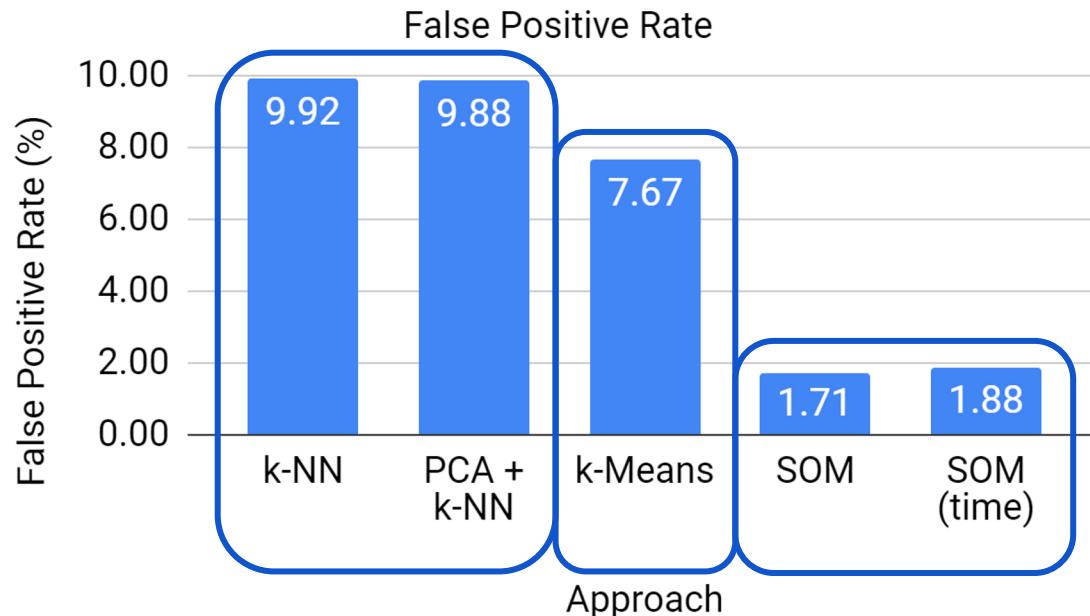
Results - Detection Rate

- Percentage of attacks in which an alarm is raised during exploitation



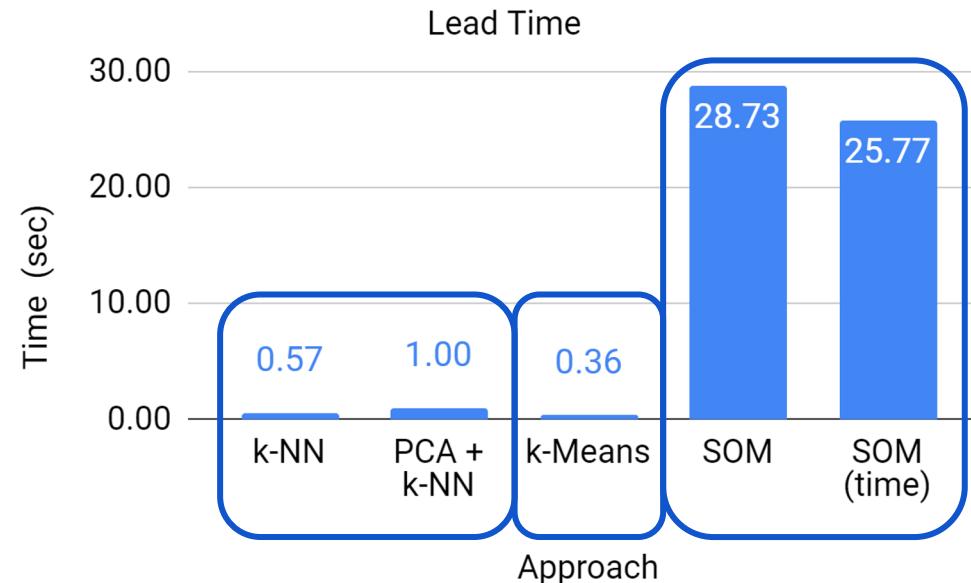
Results - False Positive Rate (FPR)

- False alarm rate



Results - Lead Time

- Time between the alarm notice and attack completion
- For the category of exploits that return a shell



Results - Summary

- Self Organizing Map (SOM) shows the most promising results
 - Detection rate
 - False positive rate
 - Lead time
- Detection over frequency vectors yields improved results over that of time vectors

Future Work

- Investigate more vulnerability case studies
- Improve detection accuracy of the studied schemes

Conclusion

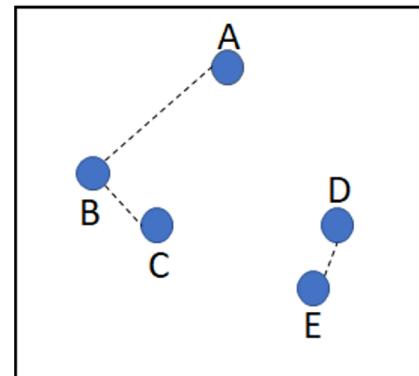
- Studied 28 real world vulnerabilities in 24 common containerized applications
- 24 of 28 vulnerability exploits detected (**85.7%**)
 - Static alone detects 3 of 28 exploits (**10.7%**)
 - Dynamic alone detects 22 of 28 exploits (**78.6%**)

Thank you!

Backup slides

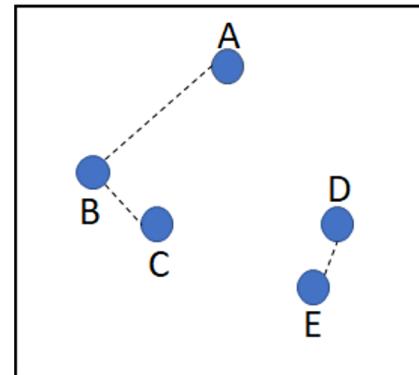
k-Nearest Neighbors

- *Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)*
- Example
 - 1-nearest neighbor
 - Largest 20% of nearest neighbor distance



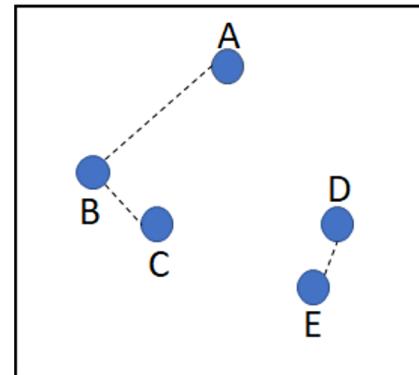
k-Nearest Neighbors

- *Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)*
- Example
 - 1-nearest neighbor
 - Largest 20% of nearest neighbor distance



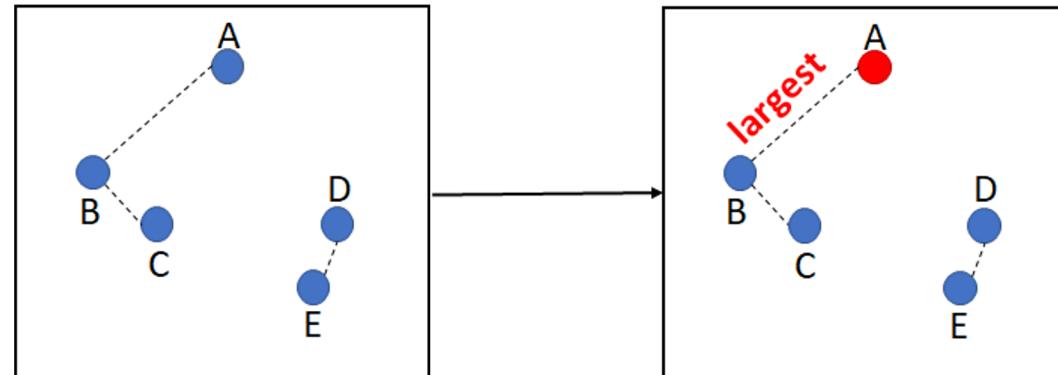
k-Nearest Neighbors

- *Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)*
- Example
 - 1-nearest neighbor
 - Largest 20% of nearest neighbor distance



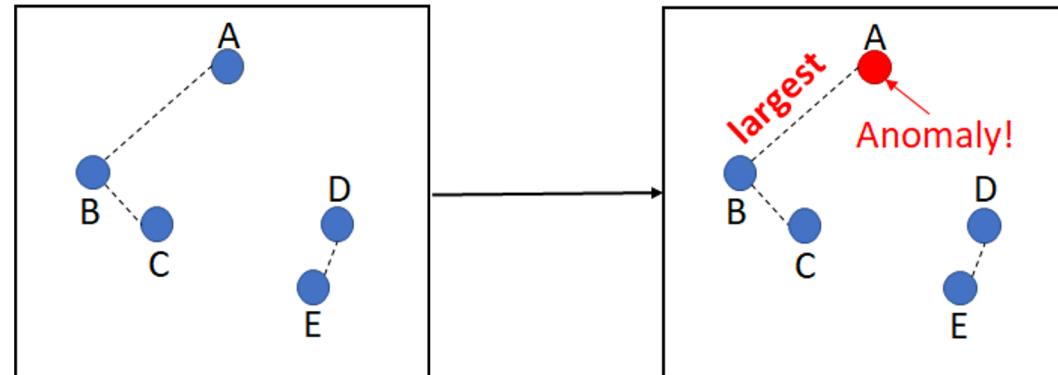
k-Nearest Neighbors

- *Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)*
- Example
 - 1-nearest neighbor
 - Largest 20% of nearest neighbor distance



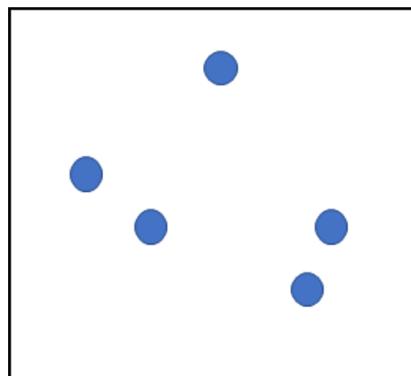
k-Nearest Neighbors

- *Principal Component Analysis (PCA) + k-Nearest Neighbors (k-NN)*
- Example
 - 1-nearest neighbor
 - Largest 20% of nearest neighbor distance



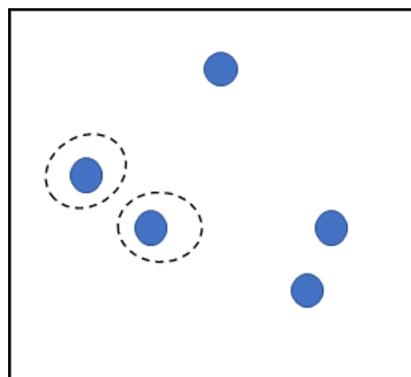
k-Means

- *k-Means*
- Example
 - 2 clusters



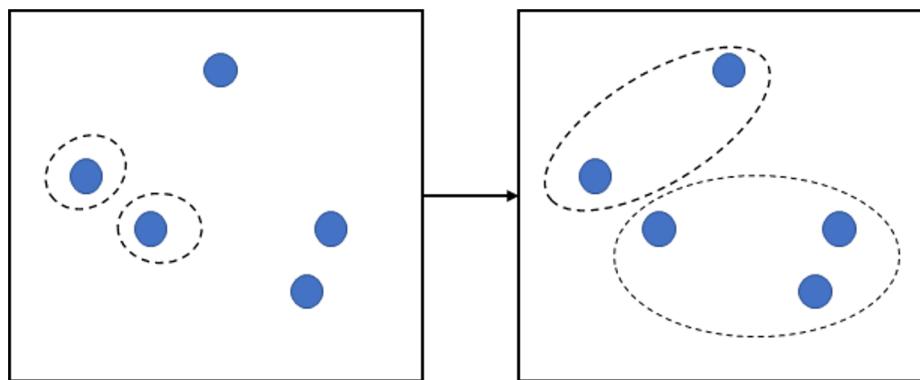
k-Means

- *k-Means*
- Example
 - 2 clusters



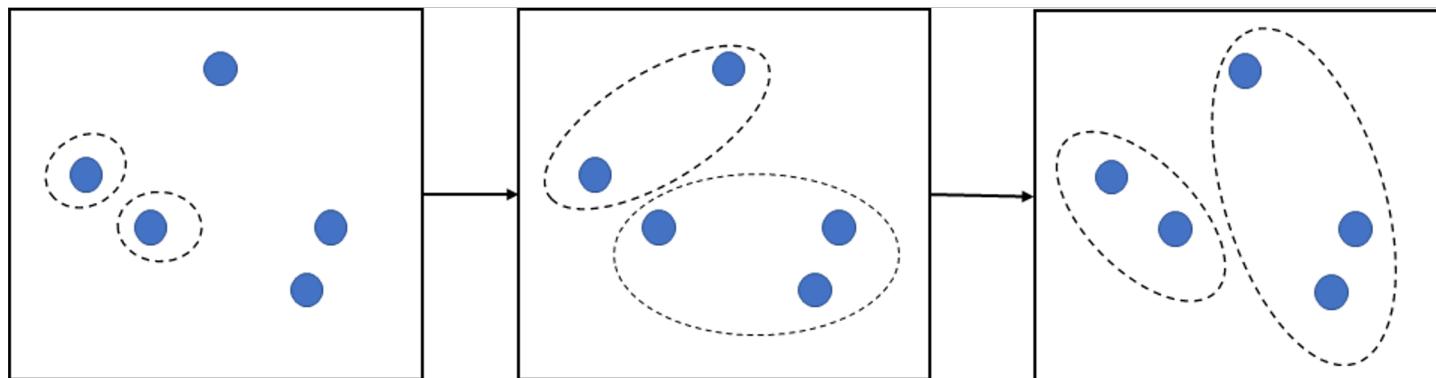
k-Means

- *k-Means*
- Example
 - 2 clusters



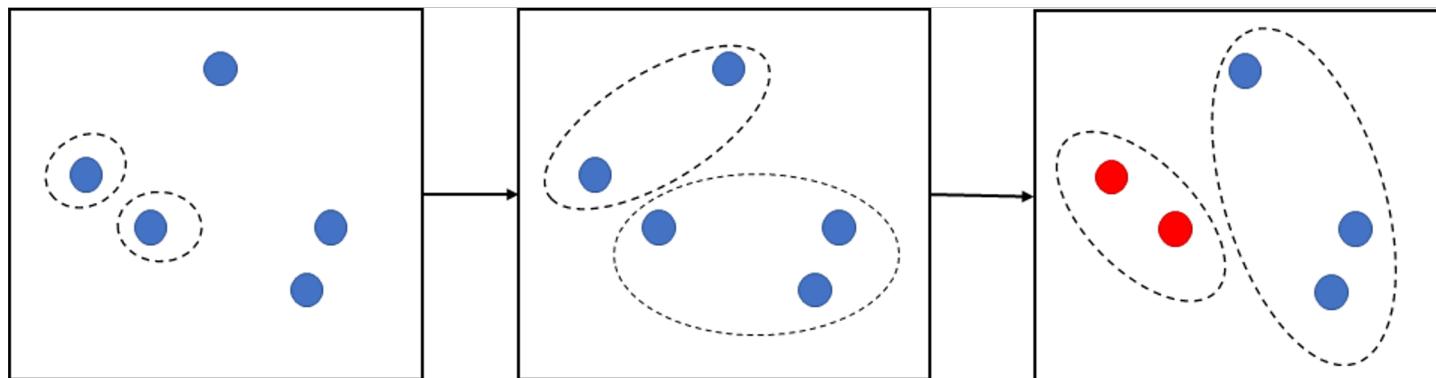
k-Means

- *k-Means*
- Example
 - 2 clusters
 - Cluster size threshold of 2



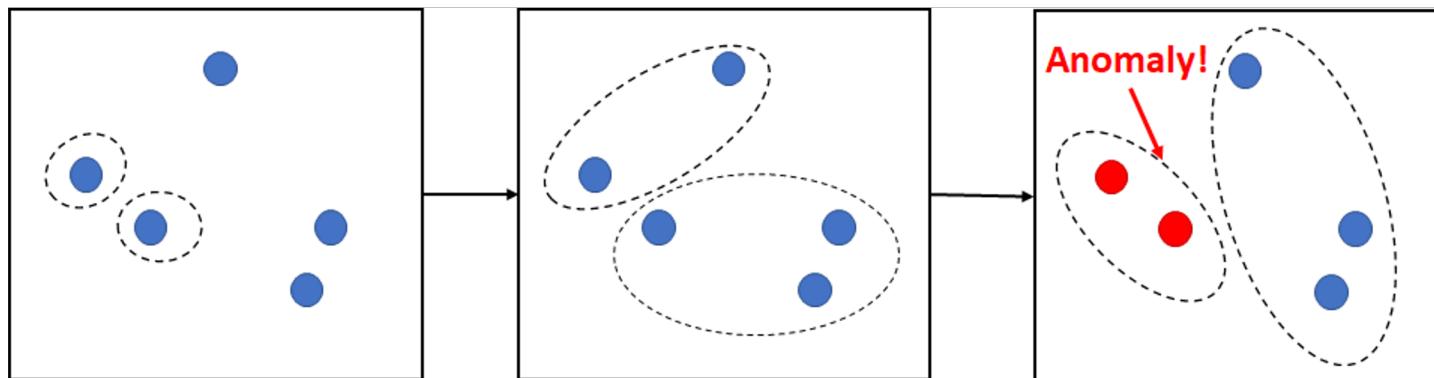
k-Means

- *k-Means*
- Example
 - 2 clusters
 - Cluster size threshold of 2



k-Means

- *k*-Means
- Example
 - 2 clusters
 - Cluster size threshold of 2



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + N(t)L(t)(V(t) - W(t))$
 - $W(t + 1) = W(t) + (1)(0.5)(V(t) - W(t))$



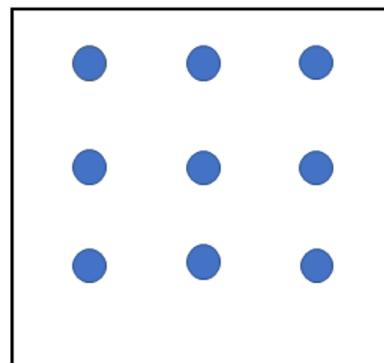
Winning Neuron



Neighbor Neurons



Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + (1)(0.5)(V(t) - W(t))$
 - $W(t + 1) = W(t) + (0.5)(V(t) - W(t))$



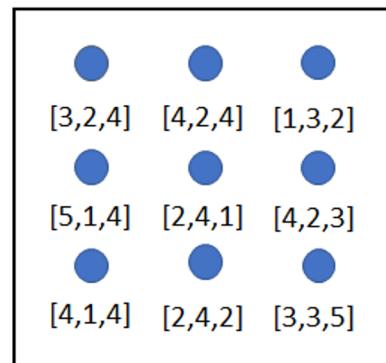
Winning Neuron



Neighbor Neurons



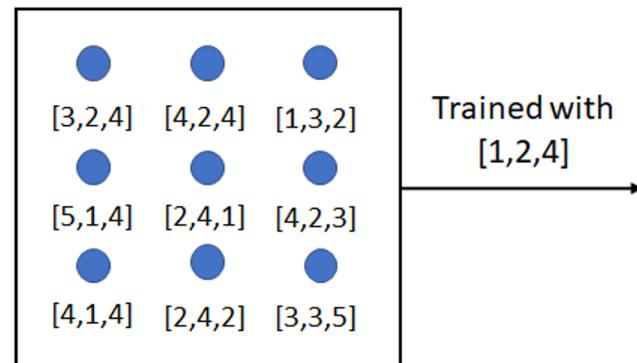
Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + (0.5)(V(t) - W(t))$
 - [1, 2, 4] closest to [3, 2, 4]

● Winning Neuron ● Neighbor Neurons ● Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + (0.5)(V(t) - W(t))$
 - $W(t + 1) = [3, 2, 4] + (0.5)[-2, 0, 0] = [2, 2, 4]$



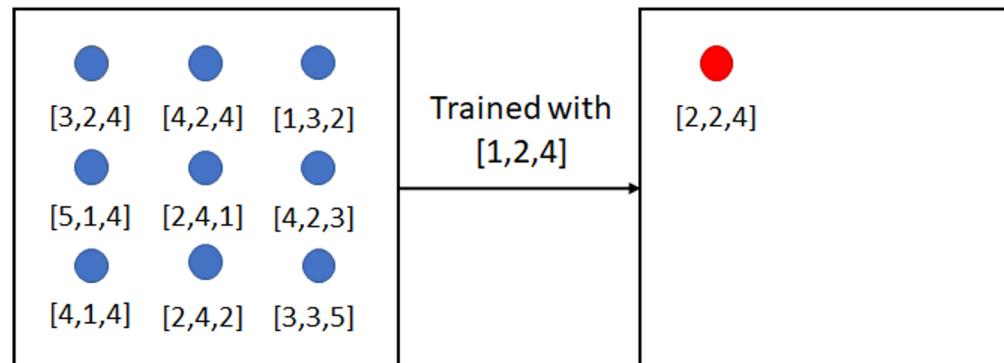
Winning Neuron



Neighbor Neurons



Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + N(t)L(t)(V(t) - W(t))$
 - $W(t + 1) = W(t) + (1)(0.5)(V(t) - W(t))$



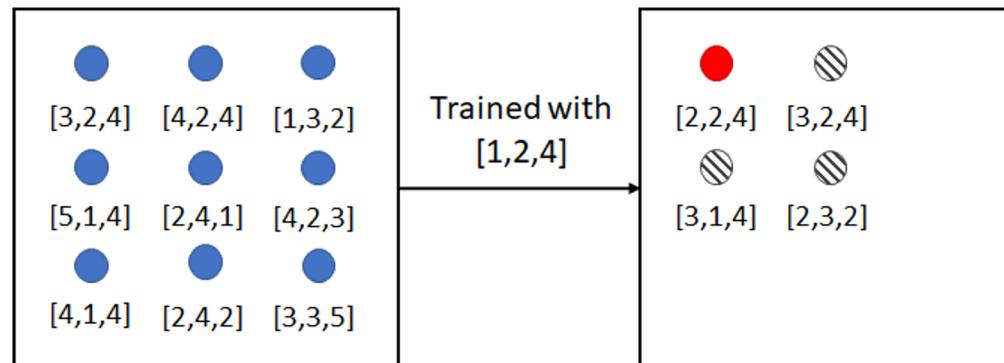
Winning Neuron



Neighbor Neurons



Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) learning phase*
- Example
 - $W(t + 1) = W(t) + N(t)L(t)(V(t) - W(t))$
 - $W(t + 1) = W(t) + (1)(0.5)(V(t) - W(t))$



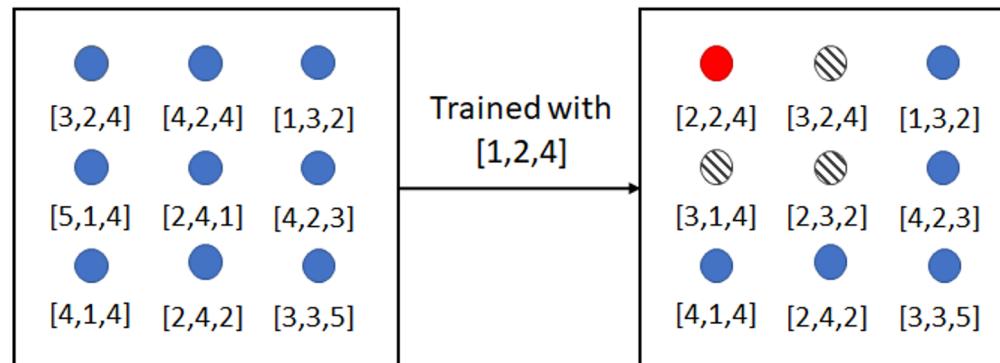
Winning Neuron



Neighbor Neurons



Other Neurons



Self Organizing Map

- *Self Organizing Map (SOM) sample mapping*
- For each test sample
 - Sum distances between the winning neuron and its neighbors (neighborhood area size)
 - Anomalies are in the largest percentile of sizes

● Winning Neuron ● Neighbor Neurons ● Other Neurons

